

Advanced Operating Systems  
MS degree in Computer Engineering  
University of Rome Tor Vergata   
Lecturer: Francesco Quaglia

## **Trap/interrupt architecture**

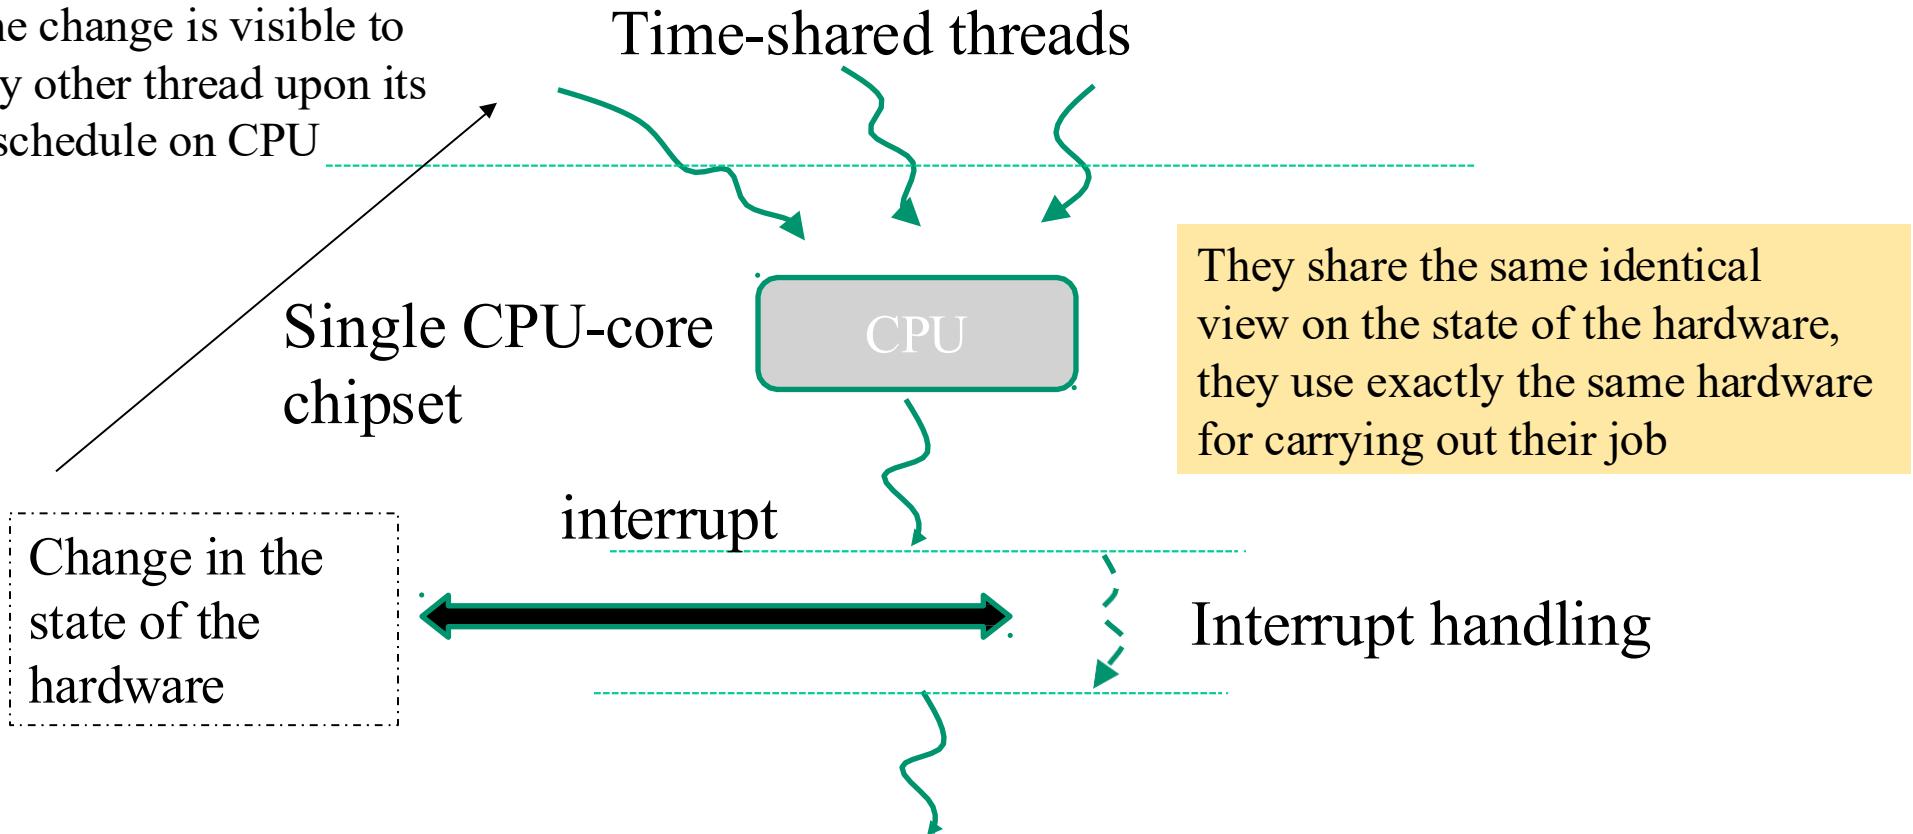
1. Architectural hints
2. Relations with software and its layering
3. Binding to the Linux kernel internals

# Single-core traditional concepts

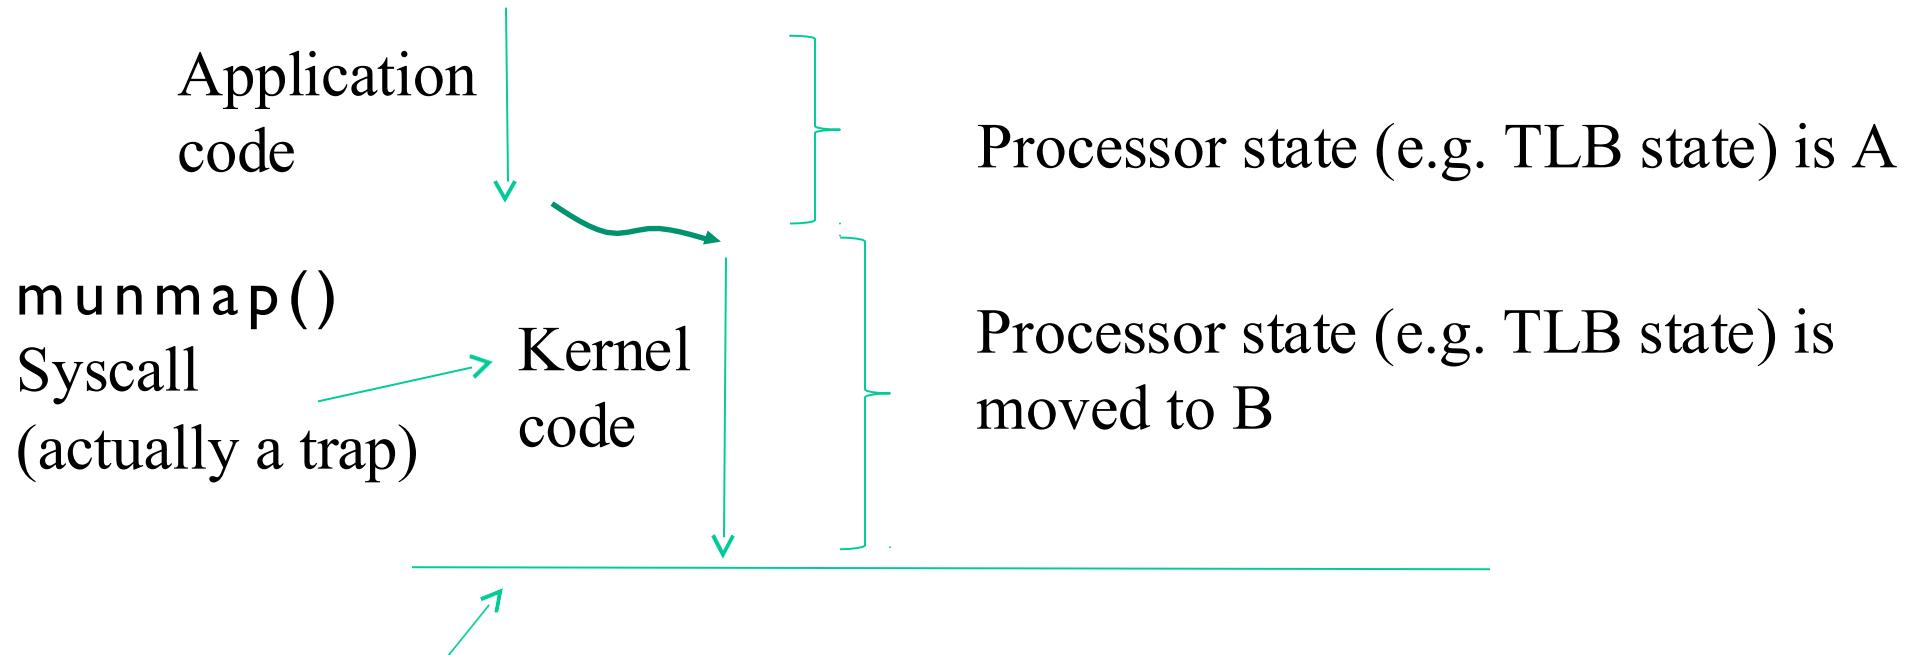
- Traditional single-core machines only relied on
  - Traps (synchronous events wrt software execution)
  - Interrupts from external devices (asynchronous events)
- The classical way of handling the event has been based on running operating system code on the **unique CPU** in the system (single CPU systems) upon event acceptance
- This has been enough (in terms of consistency) even for concurrent (multi-thread) applications given that the state of the hardware was time-shared across threads

# Some more insights

The change is visible to any other thread upon its reschedule on CPU

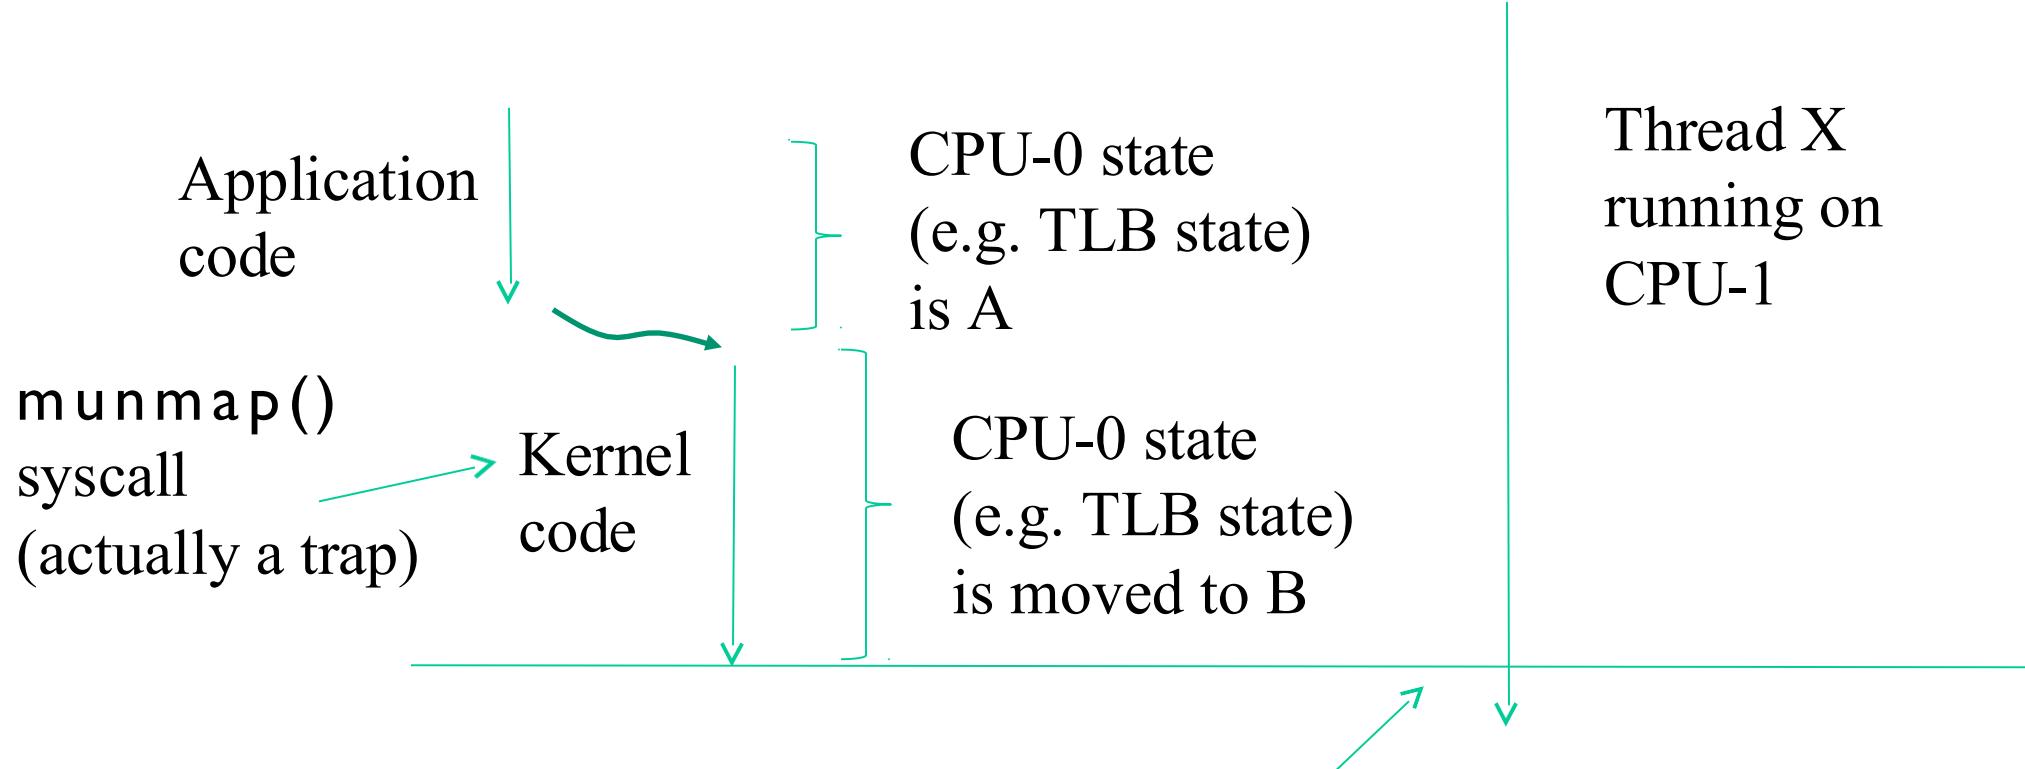


# An example with traps (e.g. syscalls)



from this point any time-shared thread sees the correct final state as determined by trap handling

# Moving to multiple CPU systems



This thread does not see state B – what if the TLB on CPU-1 caches the same page table (the same state portion) as the one of CPU-0??

# Core issues

- If the system state is distributed/replicated within the hardware architecture  
we need mechanisms for allowing state changes by traps/interrupts to be propagated
- As an example, a trap on CPU-0 needs to be propagated to CPU-1 etc.
- In some cases this is addressed by pure firmware protocols (such as when the event **is bound to deterministic handling**)
- Otherwise we need mechanisms to propagate and handle the event at the operating system (software) level

# The IPI (Inter Processor Interrupt) support

- IPI is a third type of event (beyond traps and classical interrupts) that may trigger the execution of specific operating system software on any CPU
- An IPI is a synchronous event at the sender CPU and an asynchronous one at the recipient CPU
- On the other hand, IPI is typically used to put in place cross CPU activities (e.g. request/reply protocols) allowing, e.g., a specific CPU to trigger a change in the state of another one
- Or to trigger a change on the hardware portion only observable by the other CPU

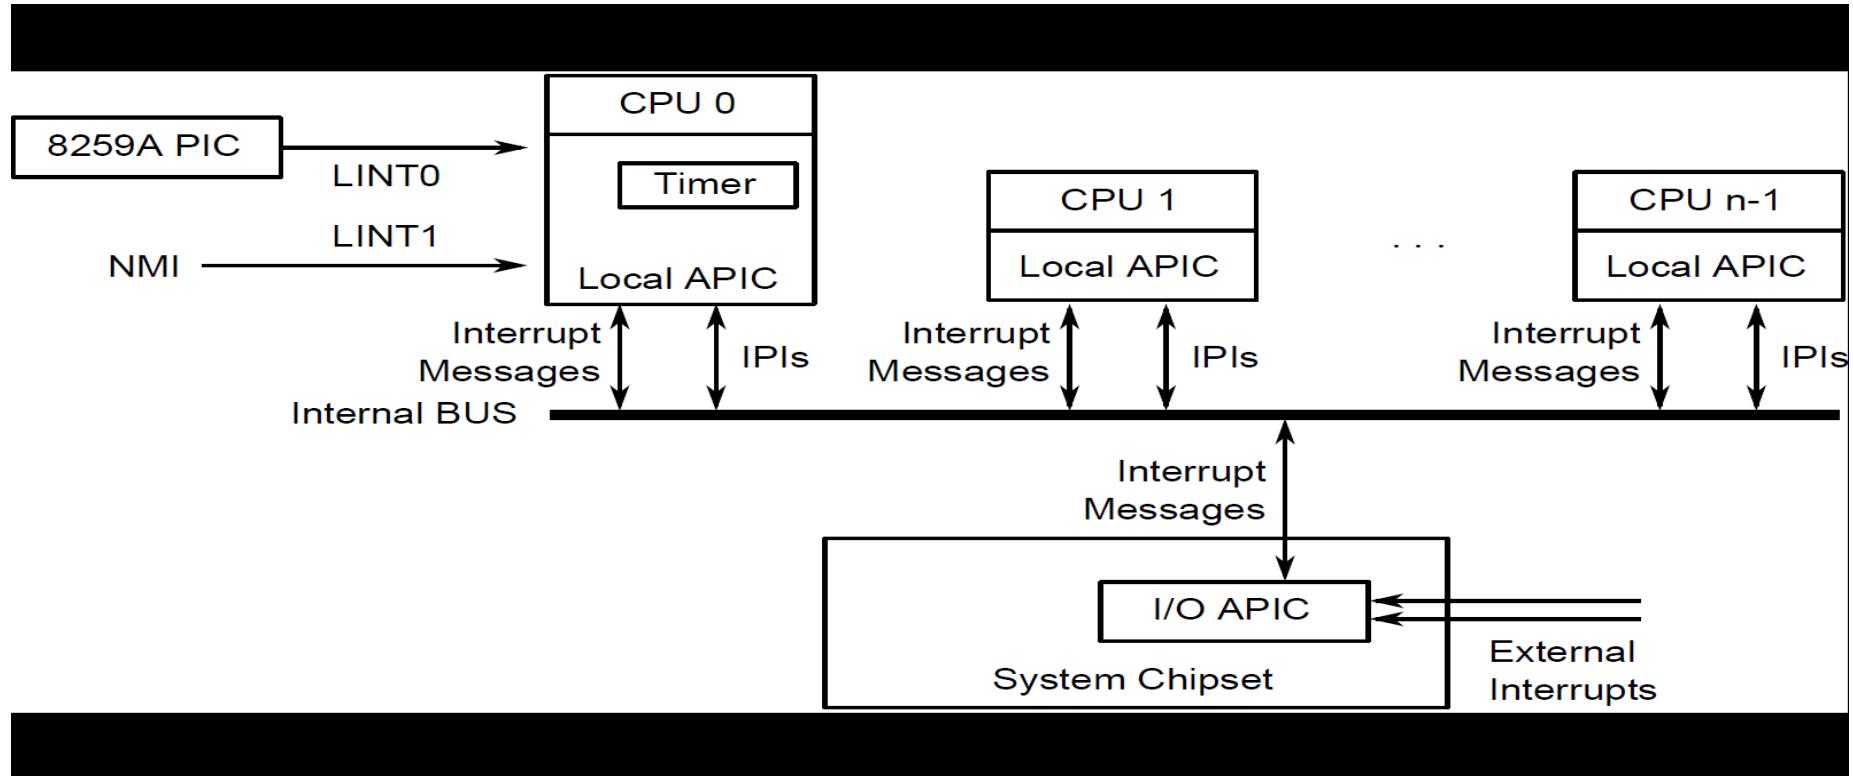
# Priorities

- IPIs are generated via firmware support, but are finally processed at software level (it becomes an OS matter)
- Classically, at least two priority levels are admitted
  - ✓ High
  - ✓ Low
- High priority leads to immediate processing of the IPI at the recipient (a single IPI is accepted and stands out at any point in time)
- Low priority generally leads to queue the requests and process them via sequentialization

# Actual support in x86 machines

- In x86 processors, the basic firmware support for interrupts is the so called APIC (Advanced Programmable Interrupt Controller)
- This offers a local instance to any CPU (called LAPIC – Local APIC)
- As an example, LAPIC offers a “CPU-local” programmable timer (for time tracking and time-sharing purposes) .... the LAPIC-T we already met
- It also offers pseudo-registers to be used for posting IPI requests in the system
- IPI requests travel along an ad-hoc APIC bus

# The architectural scheme



# Nomenclature

- IRQ is the actual code associated with the interrupt request (depending on hardware configuration carried out by the OS kernel)
- INT is the “interrupt line” as seen by the OS-kernel software
- In the essence  $\text{INT} = F(\text{IRQ})$  or  $\text{INT} = \text{kernel-decided-code}$
- The evaluation of the function F is typically hardware specific
- As it will be clear in a few slides, on x86 processors  $\text{INT} > \text{IRQ} + 32$
- Or the kernel-decided-code is larger than 32
- This means that the first 32 INT lines are reserved for something else – these are the predefined traps of the hardware architecture

# I/O APIC insights

- I/O APIC tracks how many CPUs are in the current chipset
- It can selectively direct interrupts to the different CPUs
- It uses so called local APIC-ID as an identifier of the CPU
- Fixed/physical operations
  - ✓ it sends interrupts from certain device to some single, predefined CPU
  - ✓ This occurs, e.g., when starting Linux with the **noapic pci=nomsi** command  
(all external devices go to the PIC - CPU0)
- Logical/low priority operations
  - ✓ it can deliver interrupts from certain device to multiple CPUs in a round robin fashion

# The Linux interface for APIC

- `/proc/interrupt` tells the actual accounting of the interrupt delivery to the different CPUs
- `/proc/irq/<IRQ#>/smp_affinity` tells what is the affinity of interrupts to CPUs in the logical/low priority operating mode
- The actual setup of the I/O APIC working mode is hard-coded into kernel boot rules and is generally observable via the `dmesg` buffer

# Linux core data structures - the IDT

- It is a table of entries that are used to describe the entry point (the GATE) for the handling of any interrupt
- x86 machines have IDTs formed by 256 entries
- the max amount of IRQ vectors we can generate with the I/O APIC architecture is 24 (standard case) or up to 48-64 (Multi-IOAPIC systems)
- The actual size and structure of the entries depends on the type of machine we are working with (say 32 vs 64 bit machines)
- Here is a high level view of the actual usage of the entries .....

# Linux IDT bindings

## Vector range

0-19 (0x0-0x13)

20-31 (0x14-0x1f)

32-127 (0x20-0x7f)

128 (0x80)

129-238 (0x81-0xee)

**239 (0xef)**

240-250 (0xf0-0xfa)

**251-255 (0xfb-0xff)**

The mixture changes with kernel releases (e.g. 255 is spurious)

## Use

Nonmaskable interrupts and exceptions

Intel-reserved

External interrupts (IRQs)

Programmed exception for system calls  
(segmented style)

External interrupts (IRQs)

**Local APIC timer interrupt**

Reserved by Linux for future use

**Inter-processor interrupts**

Back here in  
a while

# What we already saw - idtr

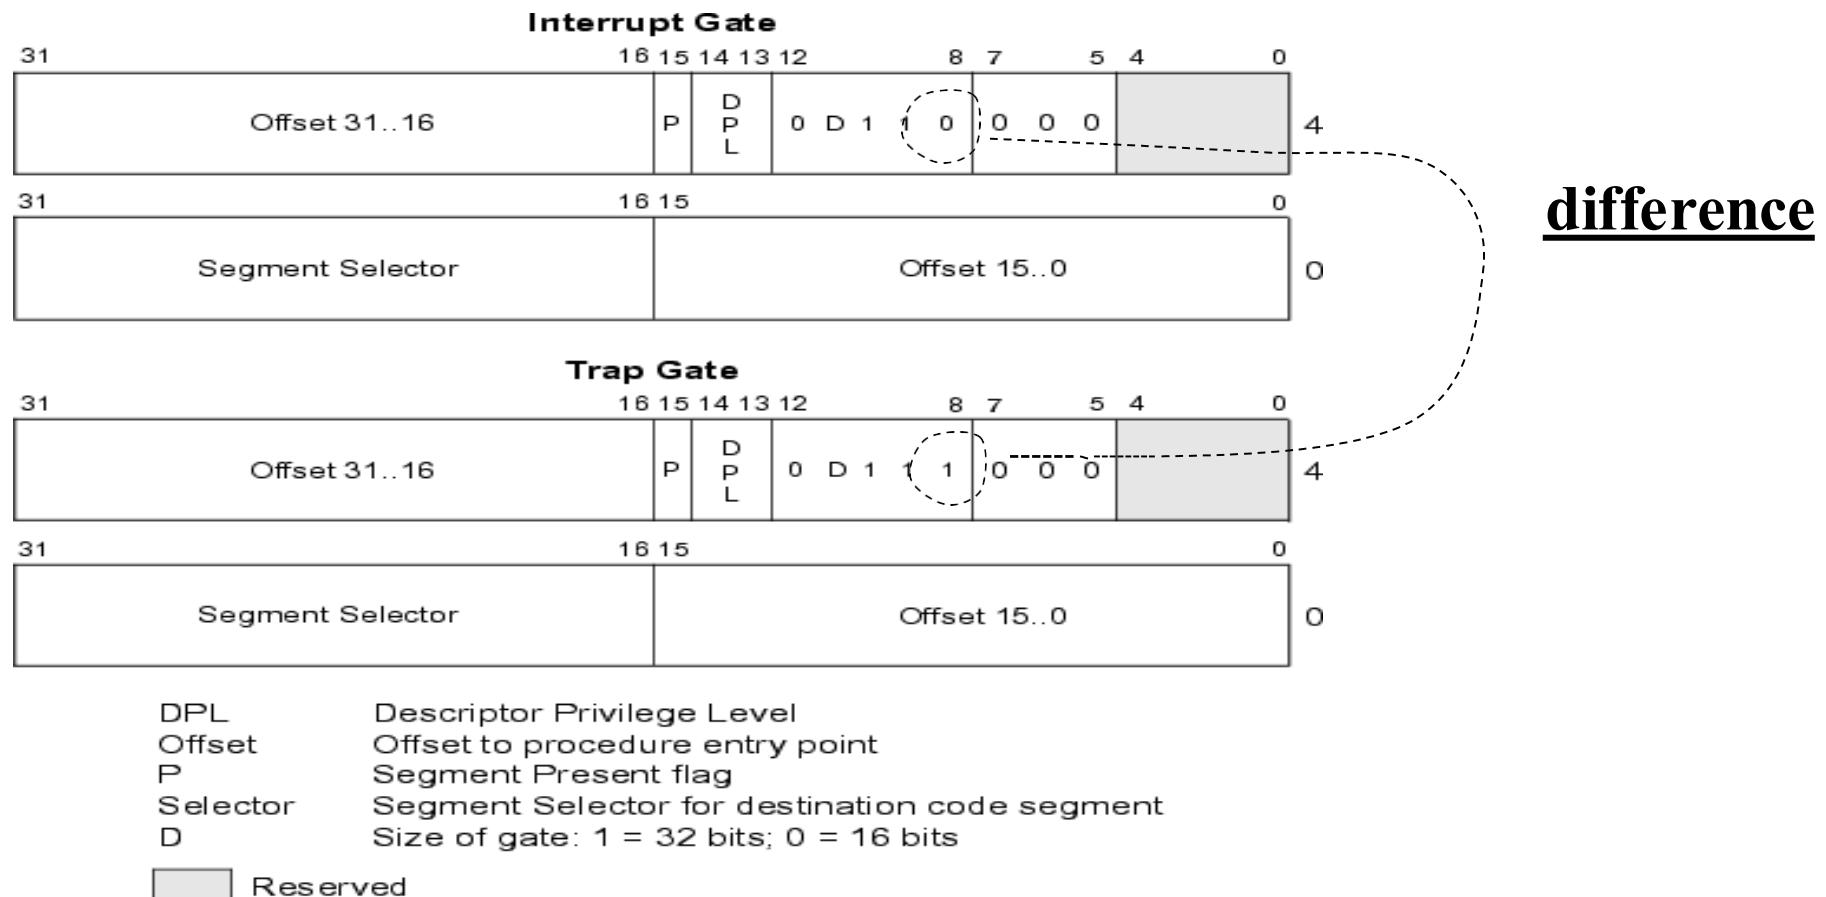
- The **idtr** register (interrupt descriptor table register) keeps on each CPU-core
  - ✓ the IDT virtual address (expressed as up to 6 bytes – 48bit – linear address)
  - ✓ The number of entries currently present in the IDT (expressed as 2 bytes – up to 256)
- This is a packed structure that we can manipulate with the LIDT (Load IDT) and SIDT (Store IDT) x86 machine instructions

# x86 protected mode

- The elements of the IDT are made up by 64-bit data structures
- In more detail, the data structure is of type `struct desc_struct`
- It is defined in `include/asm-i386/desc.h` as

```
struct desc_struct {  
    unsigned long a, b;  
}
```

# Structure of the x86 protected mode IDT entry



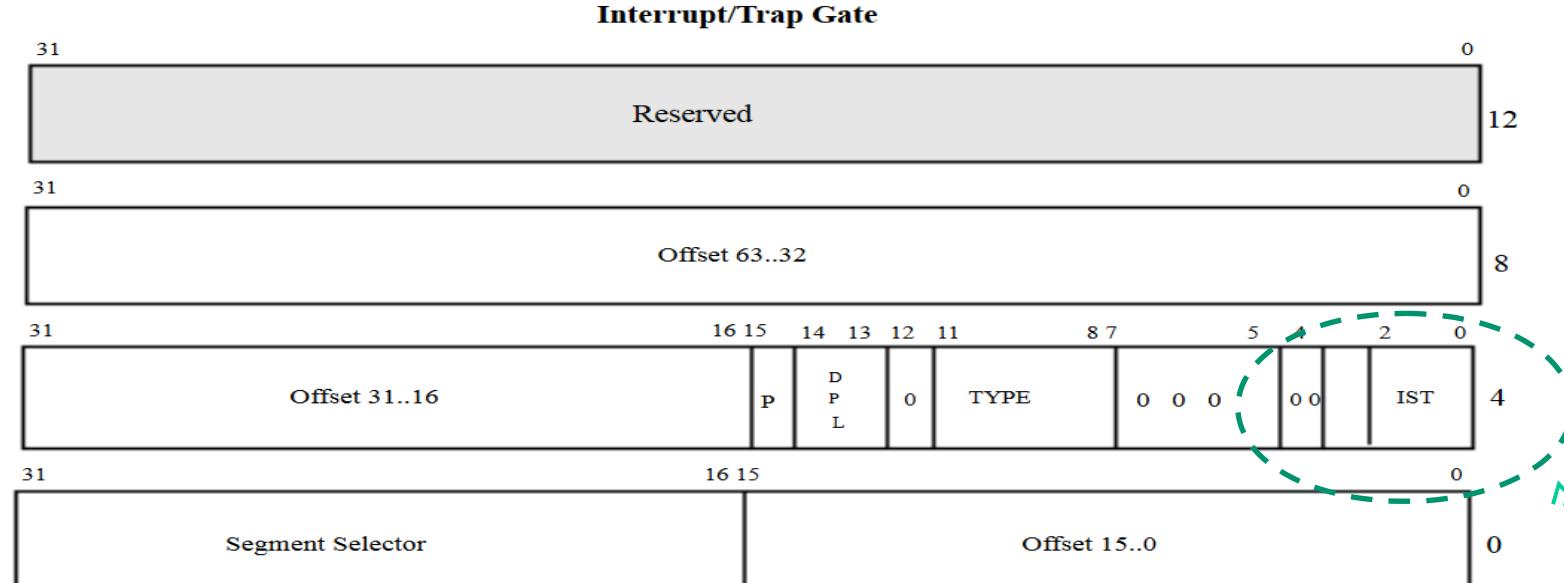
### IDT entry, Interrupt Gates

| Name            | Bit    | Full Name                  | Description  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
|-----------------|--------|----------------------------|--|--------|-----|---|--|--------|-----|---|---|--------|-----|---|--|--------|-----|----|---|--------|-----|----|--|
| <b>Offset</b>   | 48..63 | Offset 16..31              | Higher part of the offset.   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>P</b>        | 47     | Present                    | can be set to <b>0</b> for unused interrupts or for <a href="#">Paging</a> .   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>DPL</b>      | 45,46  | Descriptor Privilege Level | Gate call protection. Specifies which privilege Level the calling <a href="#">Descriptor</a> minimum should have. So hardware and CPU interrupts can be protected from being called out of userspace.  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>S</b>        | 44     | Storage Segment            | = <b>0</b> for interrupt gates.  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>Type</b>     | 40..43 | Gate Type 0..3             | <p>Possible IDT gate types :</p> <table border="1" style="margin-left: auto; margin-right: auto;"> <tr> <td>0b0101</td><td>0x5</td><td>5</td><td>80386 32 bit <a href="#">Task gate</a></td></tr> <tr> <td>0b0110</td><td>0x6</td><td>6</td><td>80286 16-bit <a href="#">interrupt gate</a></td></tr> <tr> <td>0b0111</td><td>0x7</td><td>7</td><td>80286 16-bit <a href="#">trap gate</a></td></tr> <tr> <td>0b1110</td><td>0xE</td><td>14</td><td>80386 32-bit <a href="#">interrupt gate</a></td></tr> <tr> <td>0b1111</td><td>0xF</td><td>15</td><td>80386 32-bit <a href="#">trap gate</a></td></tr> </table> | 0b0101 | 0x5 | 5 | 80386 32 bit <a href="#">Task gate</a> | 0b0110 | 0x6 | 6 | 80286 16-bit <a href="#">interrupt gate</a> | 0b0111 | 0x7 | 7 | 80286 16-bit <a href="#">trap gate</a> | 0b1110 | 0xE | 14 | 80386 32-bit <a href="#">interrupt gate</a> | 0b1111 | 0xF | 15 | 80386 32-bit <a href="#">trap gate</a> |
| 0b0101          | 0x5    | 5                          | 80386 32 bit <a href="#">Task gate</a>   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| 0b0110          | 0x6    | 6                          | 80286 16-bit <a href="#">interrupt gate</a>  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| 0b0111          | 0x7    | 7                          | 80286 16-bit <a href="#">trap gate</a>   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| 0b1110          | 0xE    | 14                         | 80386 32-bit <a href="#">interrupt gate</a>  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| 0b1111          | 0xF    | 15                         | 80386 32-bit <a href="#">trap gate</a>   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>0</b>        | 32..39 | Unused 0..7                | Have to be <b>0</b> .  |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>Selector</b> | 16..31 | Selector 0..15             | <a href="#">Selector</a> of the interrupt function (to make sense - the kernel's selector). The selector's descriptor's DPL field has to be <b>0</b> .   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |
| <b>Offset</b>   | 0..15  | Offset 0..15               | Lower part of the interrupt function's offset address (also known as pointer).   |        |     |   |  |        |     |   |   |        |     |   |  |        |     |    |   |        |     |    |  |

# Recap on relations with the GDT

- The segment identifier/selector allows accessing the entry of the GDT where we can find the base value for the target segment
- **NOTE**
  - As we already know, there are 4 valid data/code segments, all mapped to base 0x0
  - This is done in order to make Linux portable on architectures offering no segmentation support (i.e. only offering paging)
  - This is one reason why
    - ✓ Protection meta-data are also kept within page table entries
    - ✓ Setting up the offset for a GATE requires a displacement referring to 0x0, which can be denoted to the linker by the & operator

# Long mode IDT entry structure



|          |   |
|----------|---|
| DPL      | Descriptor Privilege Level                    |
| Offset   | Offset to procedure entry point               |
| P        | Segment Present flag                          |
| Selector | Segment Selector for destination code segment |
| IST      | Interrupt Stack Table                         |

Fully new

# Accessing the gate address - long mode

```
#define HML_TO_ADDR(h,m,l)      \
  ((unsigned long) (l) | ((unsigned long) (m) << 16) | \
  ((unsigned long) (h) << 32))

.....
gate_desc *gate_ptr;

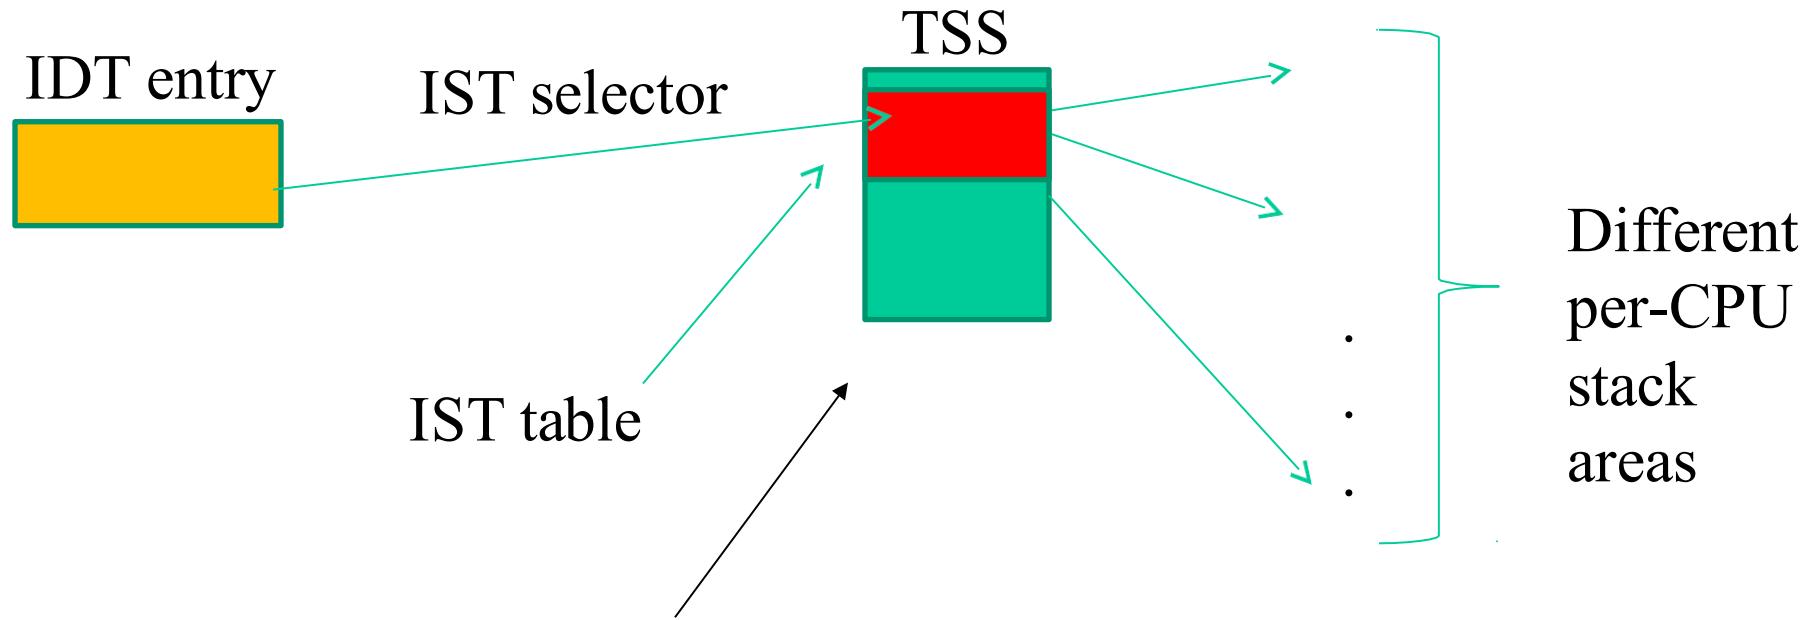
gate_ptr = ....;

HML_TO_ADDR(gate_ptr->offset_high, gate_ptr->offset_middle,
            gate_ptr->offset_low);
```

# x86 long mode fully new concepts - IST

- The **Interrupt Stack Table (IST)** is available as an alternative to handle stack switch upon traps/interrupts
- This mechanism unconditionally switches stacks when it is enabled on each individual interrupt-vector basis using a field in the IDT entry
- This means that some interrupt vectors can selectively use the IST mechanism
- IST provides a method for specific interrupts (such as NMI, double-fault, and machine-check) to always execute on a known good stack .... which is also visible when accessing the kernel
- The IST mechanism provides up to seven IST pointers in the TSS

# A scheme



These are typically the primary stacks (possibly of different size) for processing a given trap/interrupts

Software will then switch to the classical kernel level stack of the running task if nothing prevents it (e.g. a double fault)

# Macros for setting IDT entries - x86 protected mode

- within the `arch/i386/kernel/traps.c` file we can find the declaration of the following macros that can be used for setting up one entry of the IDT
  - `set_trap_gate(displacement, &symbol_name)`
  - `set_intr_gate(displacement, &symbol_name)`
  - `set_system_gate(displacement, &symbol_name)`
- displacement indicates the target entry of the IDT
- $\&\text{symbol\_name}$  identifies the segment displacement (starting from 0x0) which determines the address of the software module to be invoked for handling the trap or the interrupt

# Main differences among the modules

- The `set_trap_gate()` function initializes one IDT entry such in away to define the value 0 as the privilege level admitted for accessing the GATE via software
- Therefore we cannot rely on the INT assembly instruction unless we are already executing in kernel mode
- The `set_intr_gate()` function looks similar, however the handler activation relies on interrupt masking
- `set_system_gate()` is similar to `set_trap_gate()` however it defines the value 3 as the level of privilege admitted for accessing the GATE

# i386/kernel-2.4 examples

Handler managing division errors

```
set_trap_gate(0,&divide_error)
```

Handler for non-maskable interrupts

```
set_intr_gate(2,&nmi)
```

Handler used for dispatching system calls

```
set_system_gate(SYSCALL_VECTOR,&system_call)
```

# Variants for x86 long mode - kernel 3

## CODE SNIPPET FROM desc.h

```
409 /*
410 * This routine sets up an interrupt gate at directory privilege level 3.
411 */
412 static inline void set_system_intr_gate(unsigned int n, void *addr)
413 {
414     BUG_ON((unsigned)n > 0xFF);
415     _set_gate(n, GATE_INTERRUPT, addr, 0x3, 0, __KERNEL_CS);
416 }
417
418 static inline void set_system_trap_gate(unsigned int n, void *addr)
419 {
420     BUG_ON((unsigned)n > 0xFF);
421     _set_gate(n, GATE_TRAP, addr, 0x3, 0, __KERNEL_CS);
422 }
423
424 static inline void set_trap_gate(unsigned int n, void *addr)
425 {
426     BUG_ON((unsigned)n > 0xFF);
427     _set_gate(n, GATE_TRAP, addr, 0, 0, __KERNEL_CS);
428 }
```

# Variants for x86 long mode - kernel 4/5

```
#define write_ldt_entry(dt, entry, desc) \
    native_write_idt_entry(dt, entry, desc)

static inline void native_write_idt_entry(gate_desc *idt,
                                         int entry, const gate_desc *gate) {
    memcpy(&idt[entry], gate, sizeof(*gate));
}
```

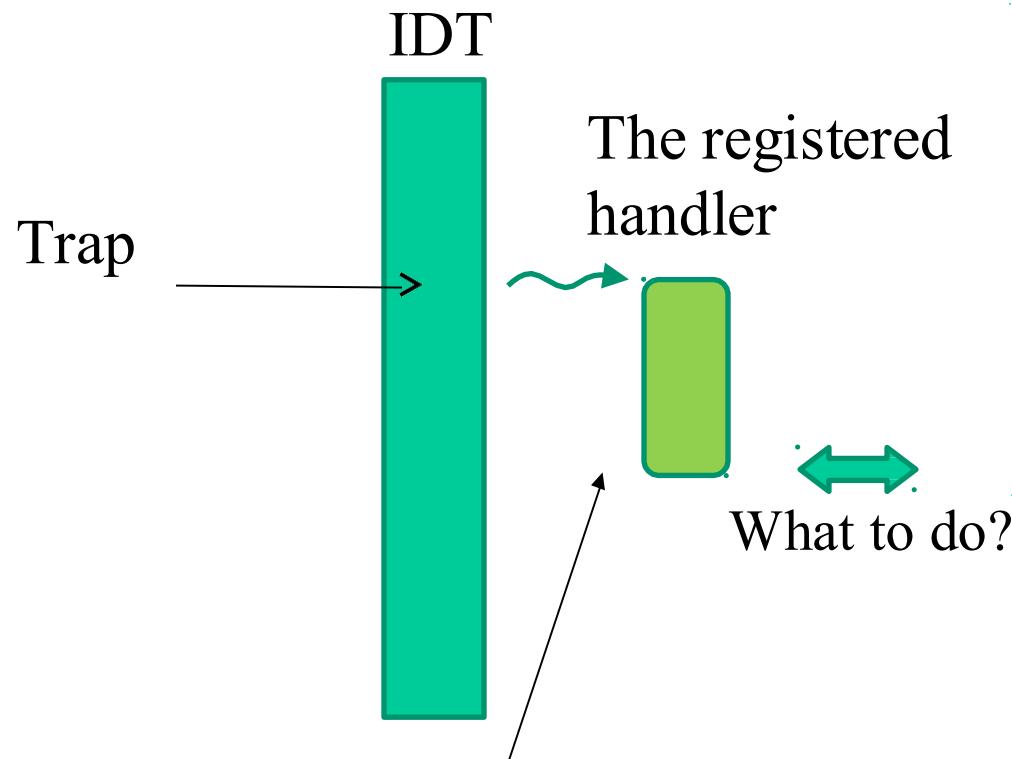
# Reserved vs available IDT entries (i)

- The entries from 0 to 31 are reserved for handlers that are used to manage specific (predefined) events/conditions (such as divide by 0 or page fault) or are already planned for future use ... these are mostly traps
- This is based on hardware design/requirements
- All the other entries are available for system programming purposes
- As an example, the entry at displacement 0x80 has been traditionally used for kernel level access via system calls
- We note that for some of the reserved entries, microcode tasks generate a **so called error-code to be passed to the handler .....**

# Reserved vs available IDT entries (ii)

- If needed, the handler needs to be structured such in a way to be aware of the production of the error-code
- Particularly, beyond exploiting the error-code value, it needs to remove it from, e.g., the stack right before returning from trap/interrupt (IRET)
- Non-reserved entries area managed by the microcode with no generation of any error-code value

# Recap on actions of trap/interrupt handlers



In modern kernels we also have the need for handling kernel isolation on page tables

CPU snapshot generation on the stack? YES

Management of the presence/absence of error code? YES

Additional stack change? YES/NO

Control passage to a second level handler? Typically YES (for sure if we are running a non top/bottom half task)

# Modular handler management

- The interrupt handlers are managed via an additional dispatcher
- Initially, each handler
  - Manages the actual access to kernel code (e.g. calling the PTI manager)
  - Logs a dummy-value into the target stack in case no error-code is generated in relation to the specific trap/interrupt – otherwise logs the generated one
  - Runs the CPU snapshot logger
- Then it passes control to the actual C-style interrupt handler

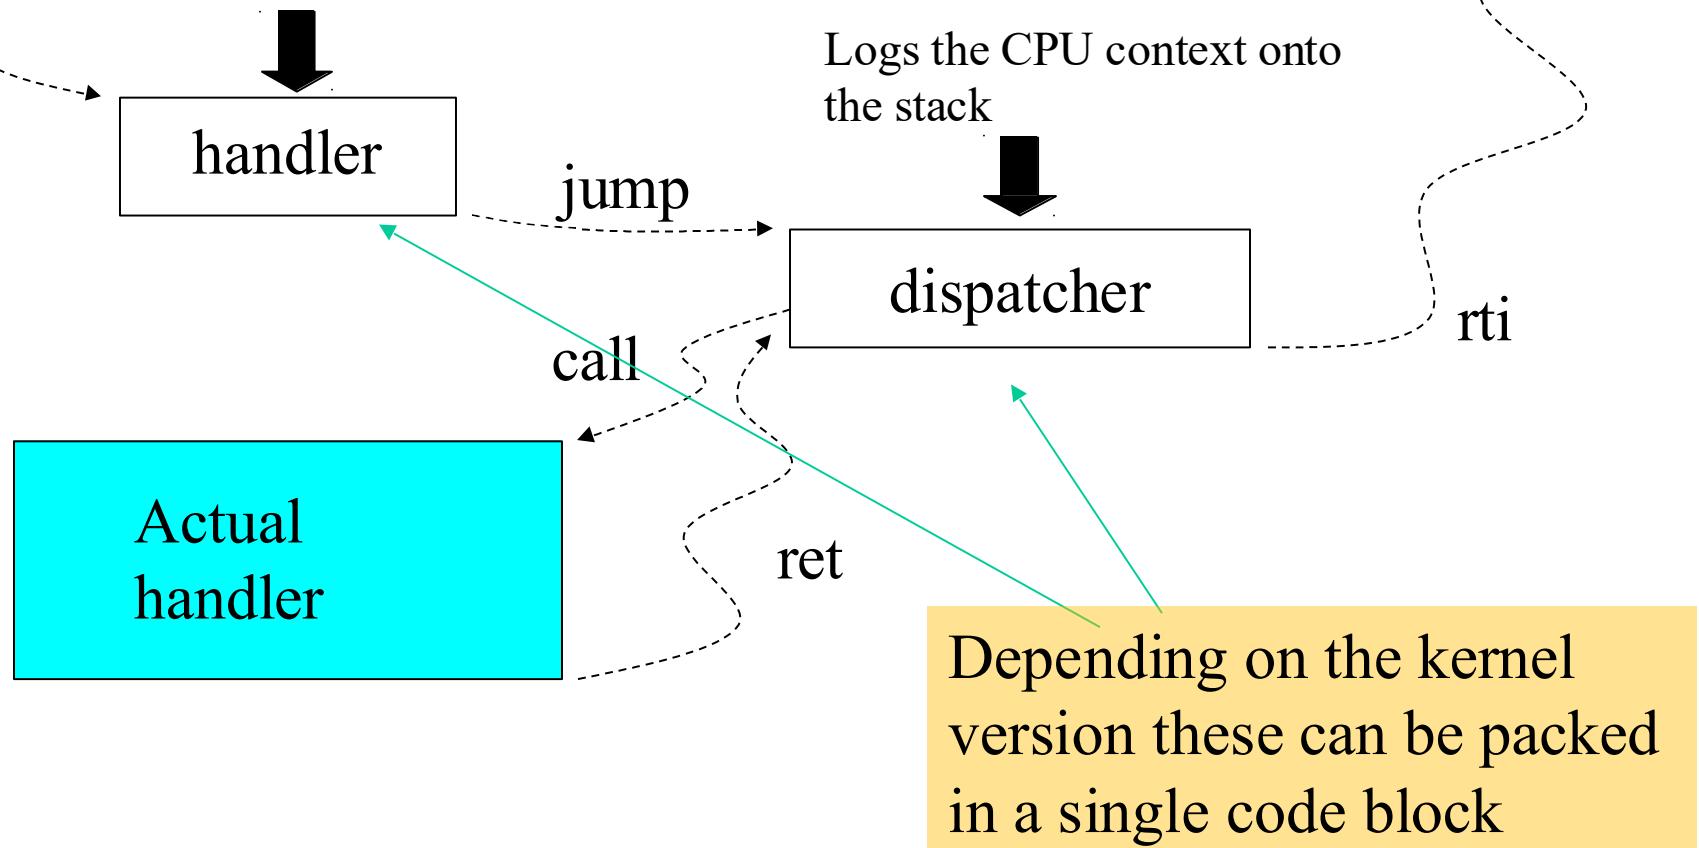
# The actual scheme

trap/interrupt

CALLs some entry level module (e.g. for PTI)

Logs the pointer/VECTOR\_INDEX for the handler

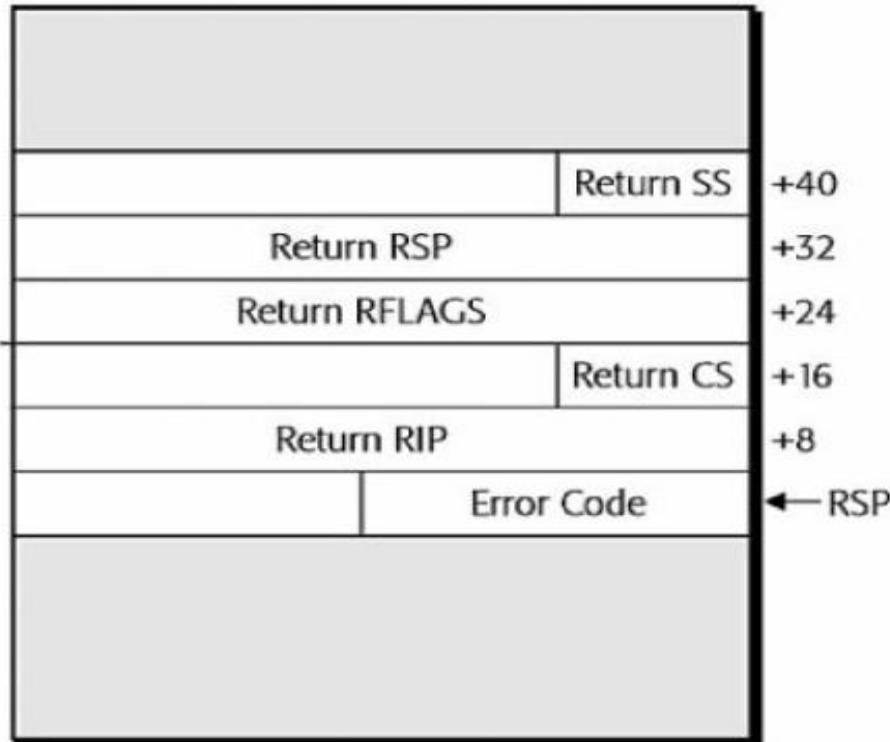
(and sometimes also the dummy value) onto final target stack



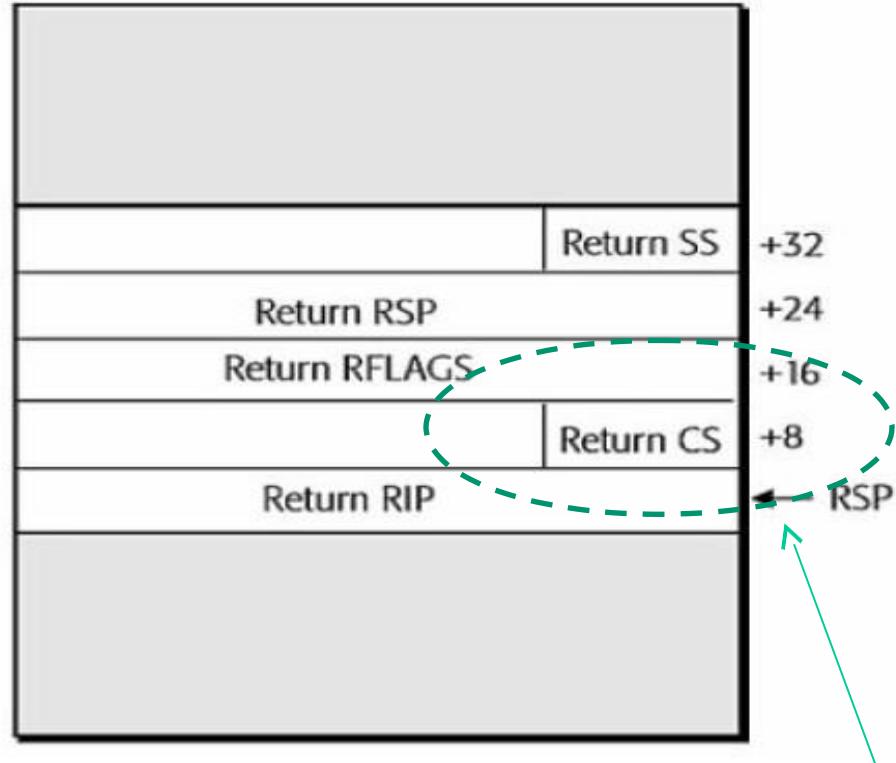
# x86-64 early trap/interrupt stack layout details

Interrupt-Handler Stack

With Error Code



With No Error Code



Coming from where?

# Examples (dated)

```
ENTRY(overflow)
```

```
    pushl $0  
    pushl $ SYMBOL_NAME(do_overflow)  
    jmp error_code
```

No error code by firmware

```
ENTRY(general_protection)
```

```
    pushl $ SYMBOL_NAME(do_general_protection)  
    jmp error_code
```

```
ENTRY(page_fault)
```

```
    pushl $ SYMBOL_NAME(do_page_fault)  
    jmp error_code
```

Error code already posted  
firmware

# The error\_code block - still i386 case

- The assembler code block called `error_code` is in charge of logging the CPU context into the stack
- This is done by aligning the stack content with the following data structure defined in `include/asm-i386/ptrace.h`

```
struct pt_regs {  
    long ebx;      long ecx;  
    long edx;      long esi;  
    long edi;      long ebp;  
    long eax;      int xds;  int xes;  
    long orig_eax; long eip;  int xcs;  
    long eflags;   long esp;  int xss;  
}
```

- The actual handler can take as input a `pt_regs*` pointer and, if needed, an `unsigned long` representing the error-code

# struct pt\_regs for x86 long mode

```
struct pt_regs {  
    unsigned long r15; ... unsigned long r12;  
    unsigned long bp;  
    unsigned long bx; /* arguments: non interrupts/non tracing syscalls only save up  
to here */  
    unsigned long r11; ... unsigned long r8;  
    unsigned long ax;  
    unsigned long cx;  
    unsigned long dx;  
    unsigned long si;  
    unsigned long di;  
    unsigned long orig_ax; /* end of arguments */ /* cpu exception frame or undefined  
*/  
    unsigned long ip;  
    unsigned long cs;  
    unsigned long flags;  
    unsigned long sp;  
    unsigned long ss; /* top of stack page */  
}
```

# The page fault handler - main features

```
asmlinkage void __kprobes do_page_fault(struct pt_regs *regs,  
                                         unsigned long write, unsigned long address);
```



Newer higher level interface compared to the below shown classical one

```
void do_page_fault(struct pt_regs *regs, unsigned long error_code);
```



bit 0 == 0 means no page found, 1 means protection fault

bit 1 == 0 means read, 1 means write

bit 2 == 0 means kernel, 1 means user-mode

# Back to IPI

- Immediate handling is allowed for the case in which there are no data structures that are shared across CPU-cores that need to be accessed for the handling (kind of stateless scenarios)
- An example is the system-halt (e.g. upon panic)
- Other classical usages of IPI are
  - ✓ Execution of a same function across all the CPU-cores (like the initialization of per-CPU variables)
  - ✓ Change of the state of hardware components across multiple CPU-cores in the system (e.g. the TLB state)
  - ✓ Ask some CPU to preempt the current thread

# Actual IPI usage in Linux - a few examples

## CALL\_FUNCTION\_VECTOR

Sent to all CPUs but the sender, forcing those CPUs to run a function passed by the sender. The corresponding interrupt handler is named `call_function_interrupt()`. Usually this interrupt is sent to all CPUs except the CPU executing the calling function by means of the `smp_call_function()` facility function.

## RESCHEDULE\_VECTOR

When a CPU receives this type of interrupt, the corresponding handler limits itself to acknowledge the interrupt.

## INVALIDATE\_TLB\_VECTOR

Sent to all CPUs but the sender, forcing them to invalidate their TLB.

# Actual IPI API in the APIC driver

`send_IPI_all( )`

Sends an IPI to all CPUs (including the sender)

`send_IPI_allbutself( )`

Sends an IPI to all CPUs except the sender

`send_IPI_self( )`

Sends an IPI to the sender CPU

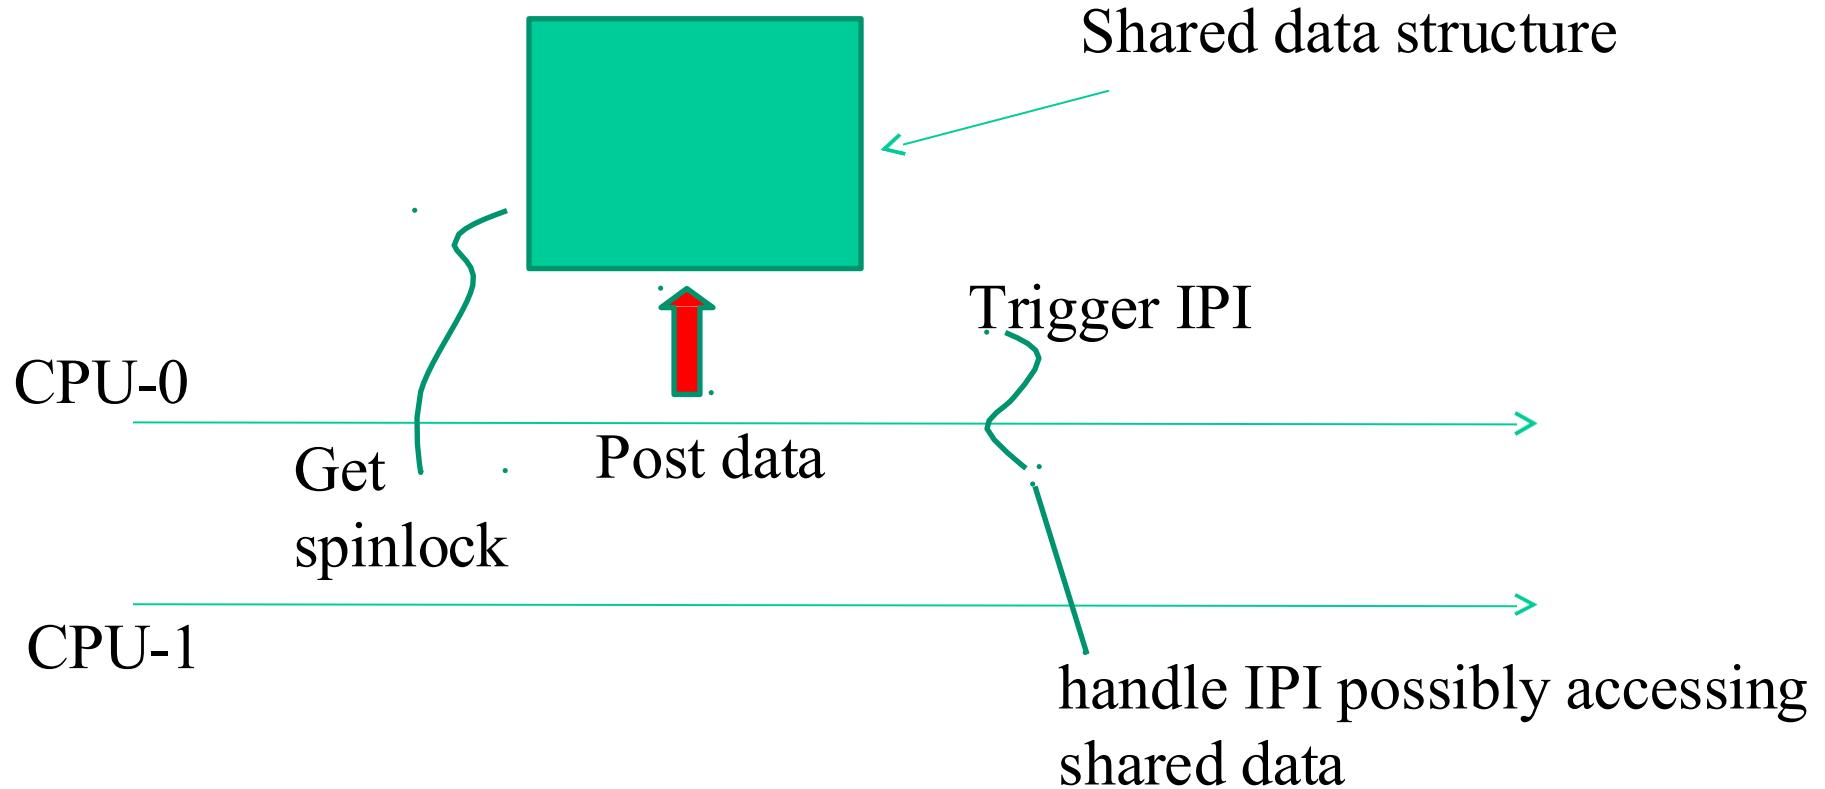
`send_IPI_mask( )`

Sends an IPI to a group of CPUs specified by a bit mask

# Sequentialization of IPI management

- The sequentializing approach is used in case the IPI requires managing a shared data structure across the threads
- This is the typical case of an IPI that requires specific parameters for correct management
- These parameters are in fact passed into predetermined memory locations accessible to all the CPU-cores, whose position in memory is predetermined
- The classical case is the one of smp-call-function, whose function pointer and parameter are both passed into a global table

# The scheme



```
207 int smp_call_function(void (*_func)(void *_info), void *_info, int wait)
208 {
209     .....
210
211     /* Can deadlock when called with interrupts disabled */
212     WARN_ON(irqs_disabled());
213
214     spin_lock_bh(&call_lock);
215
216     atomic_set(&scf_started, 0);
217     atomic_set(&scf_finished, 0);
218
219     func = _func;
220     info = _info;
221
222
223     for_each_online_cpu(i)
224         os_write_file(cpu_data[i].ipi_pipe[1], "C", 1);
225
226
227     while (atomic_read(&scf_started) != cpus)
228         barrier();
229
230     if (wait)
231         while (atomic_read(&scf_finished) != cpus)
232             barrier();
233
234     spin_unlock_bh(&call_lock);
235
236     return 0;
237 }
```



Beware this!!

# IPI additional effects

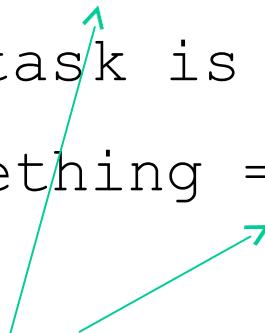
- As noted before, one IPI used by Linux is the **reschedule** one
- This may lead to preemption of the task running on the CPU targeted by the IPI
- This may have effects on both
  - ✓ Correctness/consistency
  - ✓ Performance

# Consistency aspects

- What about running a piece of code which is CPU-specific and preemption occurs??
- One example

```
struct _the_struct v[NR_CPUS];
v[smp_processor_id()] = some_value;
/* task is preempted here... */
```

```
something = v[smp_processor_id()];
```



We may be targeting different entries

# Performance aspects

- `smp_call_function()` typically runs with interrupts allowed ... just remember the deadlock issue!!
- But we cannot risk to have some `smp_call_function()` runner getting context switched off the CPU
- Otherwise the release of the `smp_call_function()` resources (e.g. the spinlock) might be delayed
- .... and we might even deadlock anyhow!!

# How to run with interrupts but no actual preemption

- We use per-thread atomic counters (we already saw)
- If the counter is not zero then no preemption will take place (although we can be targeted by interrupts)
- The check is clearly done via software upon attempting to process the preemption interrupt
- Beware managing the preemption counter explicitly if required!!

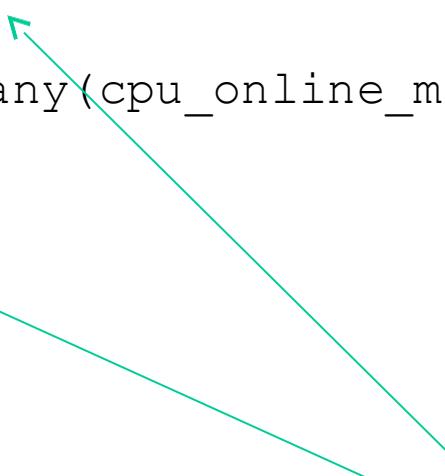
# Preemption enabling/disabling API recall

```
preempt_enable() // decrement the preempt counter  
preempt_disable() // increment the preempt counter  
preempt_enable_no_resched() decrement, but do not  
immediately preempt  
preempt_check_resched() // if needed, reschedule  
preempt_count() return the preempt counter  
put_cpu() /get_cpu() //decrease/increase the counter  
(enable/disable preemption)
```

Variants of each other

# Preemption vs SMP function calls

```
int smp_call_function(void (*func) (void *info), void *info, int wait)
{
    preempt_disable();
    smp_call_function_many(cpu_online_mask, func, info, wait );
    preempt_enable();
    return 0;
}
```



Internal structure with  
preemption awareness

# Be careful

- IPI is an extremely powerful technology
- However you need to consider scalability aspects
- This leads to conclude that IPI schemes involving large counts of CPU-cores need to be used only when mandatorily needed
- The classical example is when patching the kernel on line, e.g. upon mounting a module