Advanced Operating Systems
MS degree in Computer Engineering
University of Rome Tor Vergata
Lecturer: Francesco Quaglia

**Kernel programming basics**
- Addressing schemes and software protection models
- Hardware/software protection support
- Kernel access GATEs
- Per-CPU/per-thread memory
- System call dispatching
- Case study: LINUX (Kernels 2.4/2.6/3.xx/4.xx)

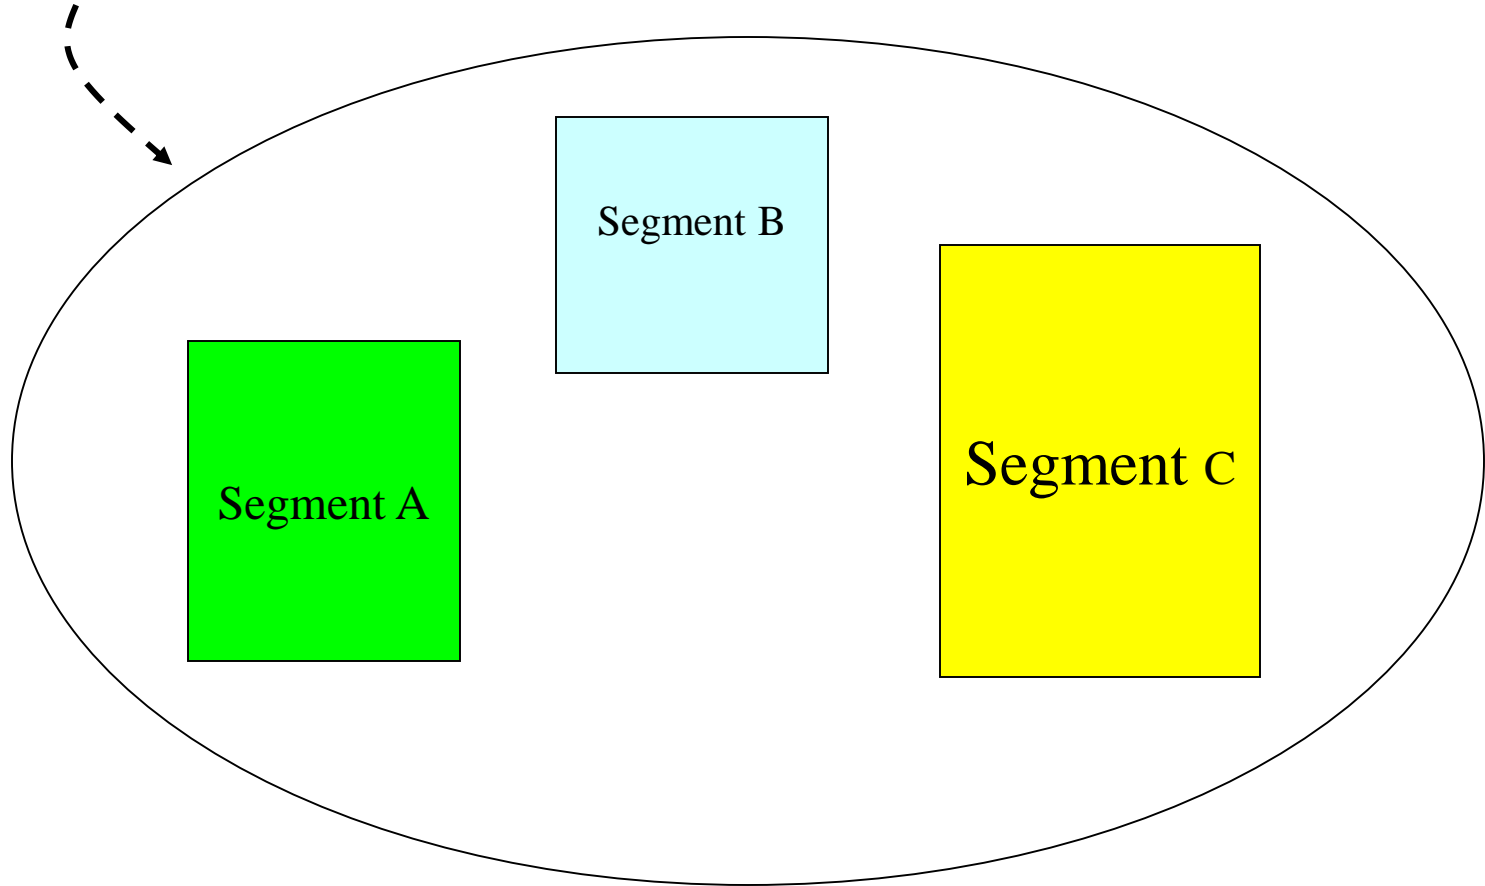# Linear addressing

Whatever memory slice available
for software execution (physical vs logical)
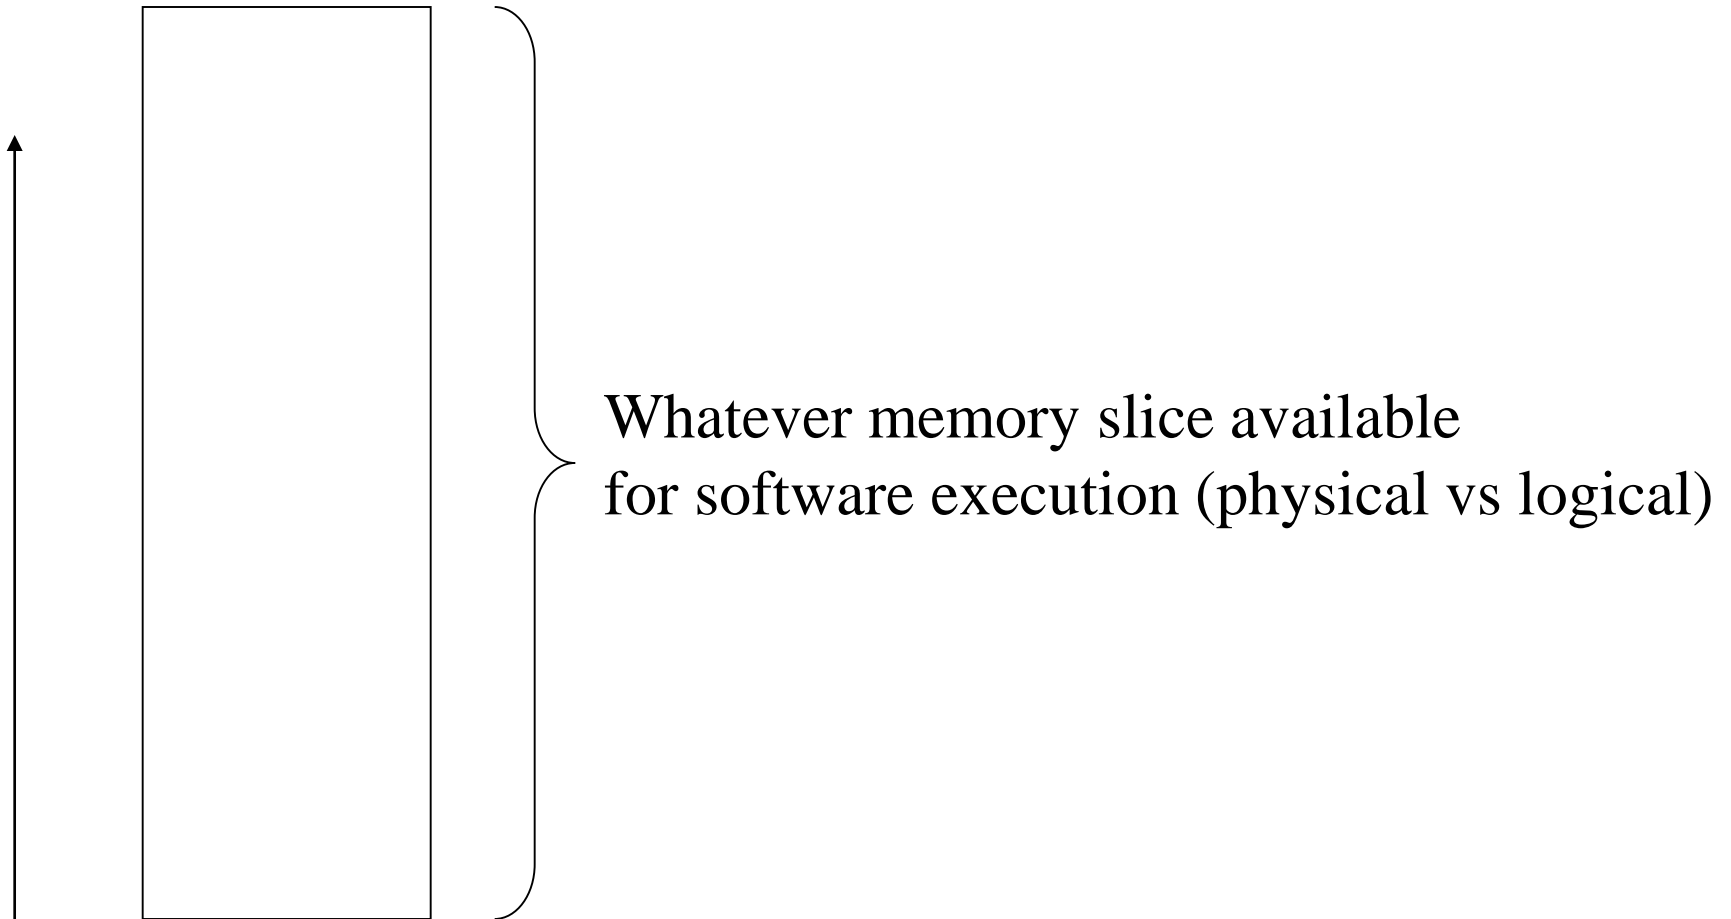
**Linear address (<offset>)**

# Segmentation

Address space (a linear one)

Segment B

Segment A

Segment C

**address = <seg.id,offset> (es. <A,0x10)**

# Combining segments in a linear address space

**Address specification = <seg.id,offset> (es. <B,offset>)**

Need to know where B is located in the linear address space (this is the "base" of B)

Then the linear address is **<base+offset>**

Segment C

Segment B

Segment A

# Virtual memory



Kernel mode (code + data/stack)

RAM

user mode (code + data/stack)
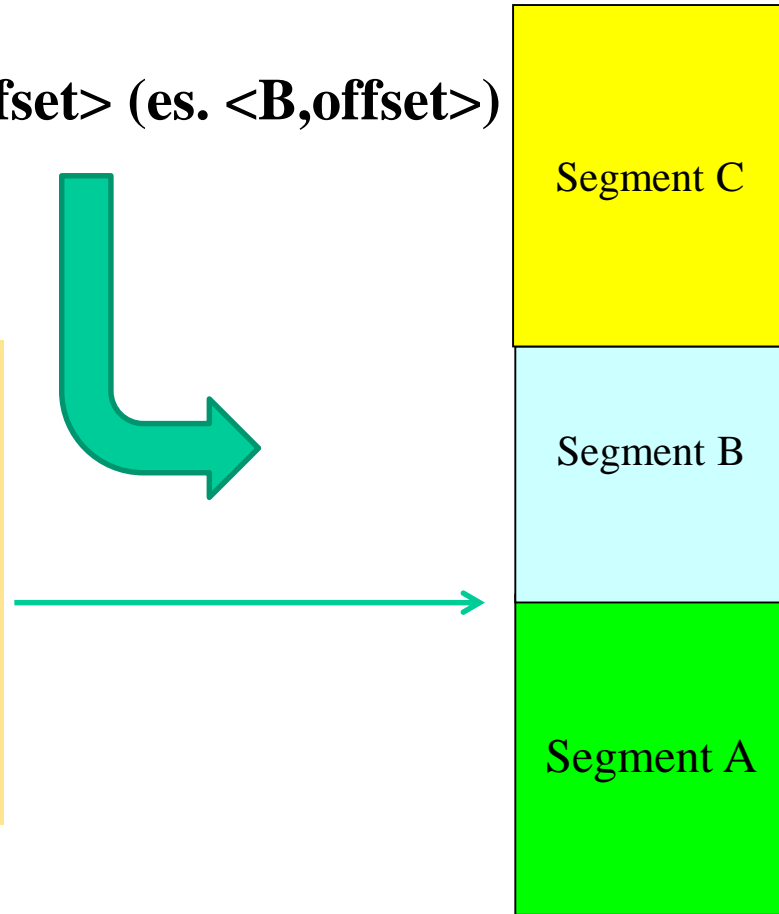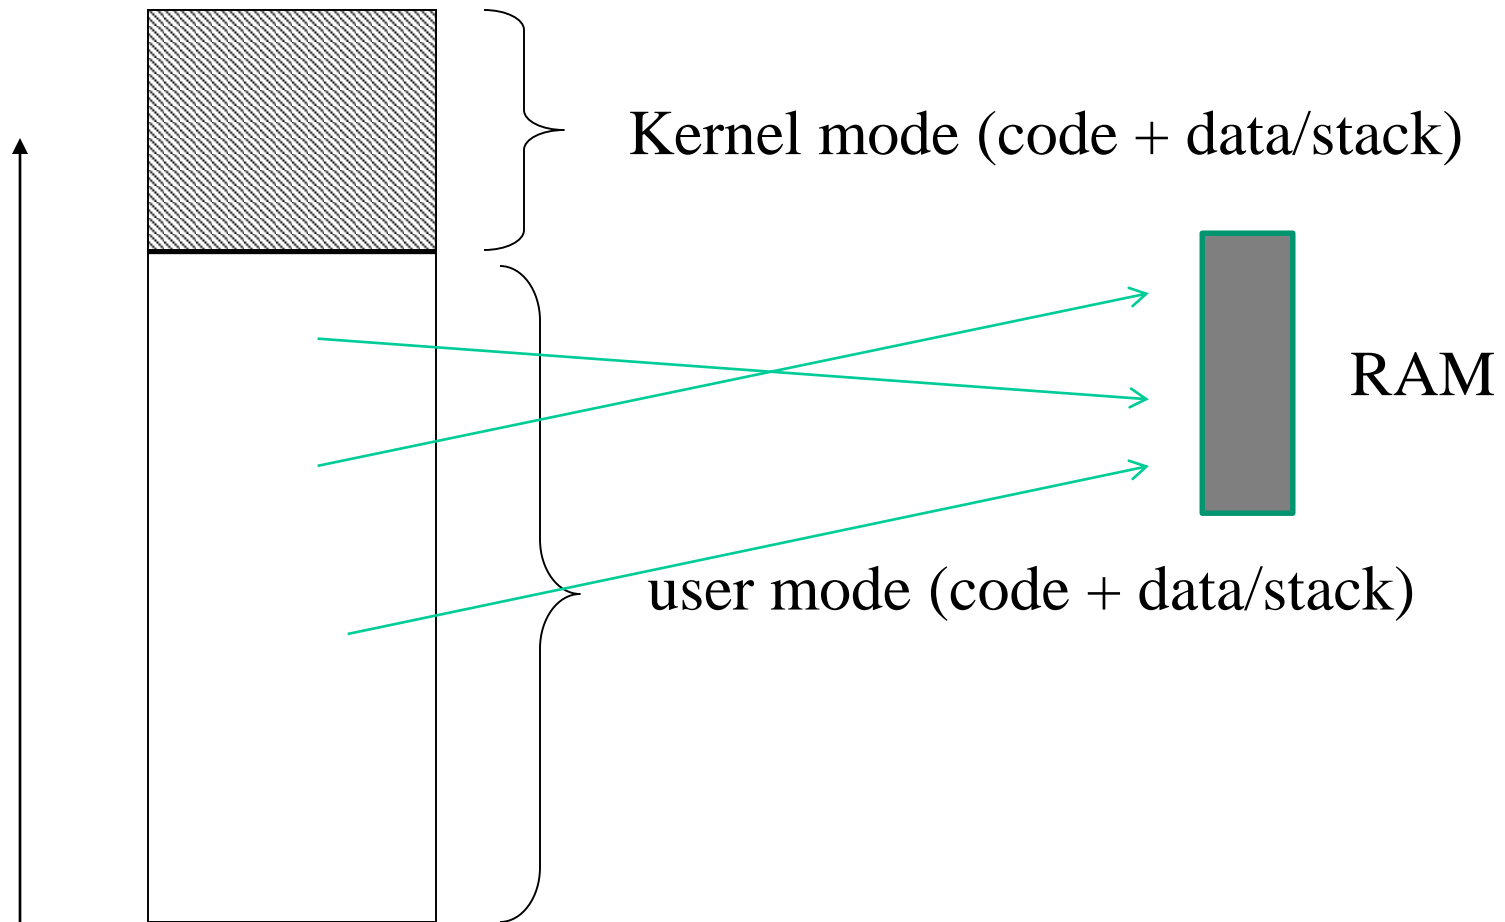
**Linear addressing + mapping to actual storage (if existing)**

# Segmentation based addresses

- Code relies on **addresses** formed by ****

- If segment numbers are not specified by the machine instruction, some <u>default segment</u> is used for each target datum (<u>instruction or operand</u>)

- Modern processors (system processors) are equipped such in a way to support segmentation efficiently, in combination with linear addressing and virtual memory (say paging)

- The whole architecture is therefore requested to handle a complex address mapping scheme such as

  **segmented addr $\Rightarrow$ linear addr $\Rightarrow$ paged addr $\Rightarrow$ physical addr**

# A very base x86 example

```
mov (%rax), %rbx
push %rbx
```

We might not imagine it but, when running this piece of code our x86 processor is implicitly using 3 different segments of memory!!

To have an exact idea of what is going on along program flow (in terms of reflection on the hardware usage) we need to know such segmentation related details

# "System" processors vs segmentation

- "system" processors (those oriented to host operating system software) rely on hardware components that allow **fast and transparent access to segmentation information (e.g. segment specific information**)

- These are
  - ➢CPU registers
  - ➢Main memory tables (directly pointed by registers)

# Segmentation with paging

| Segment number | offset |
|---|---|

- both are logical addresses
- the offset value may differ

HW supported translation

| PDE | PTE | page offset |
|---|---|---|
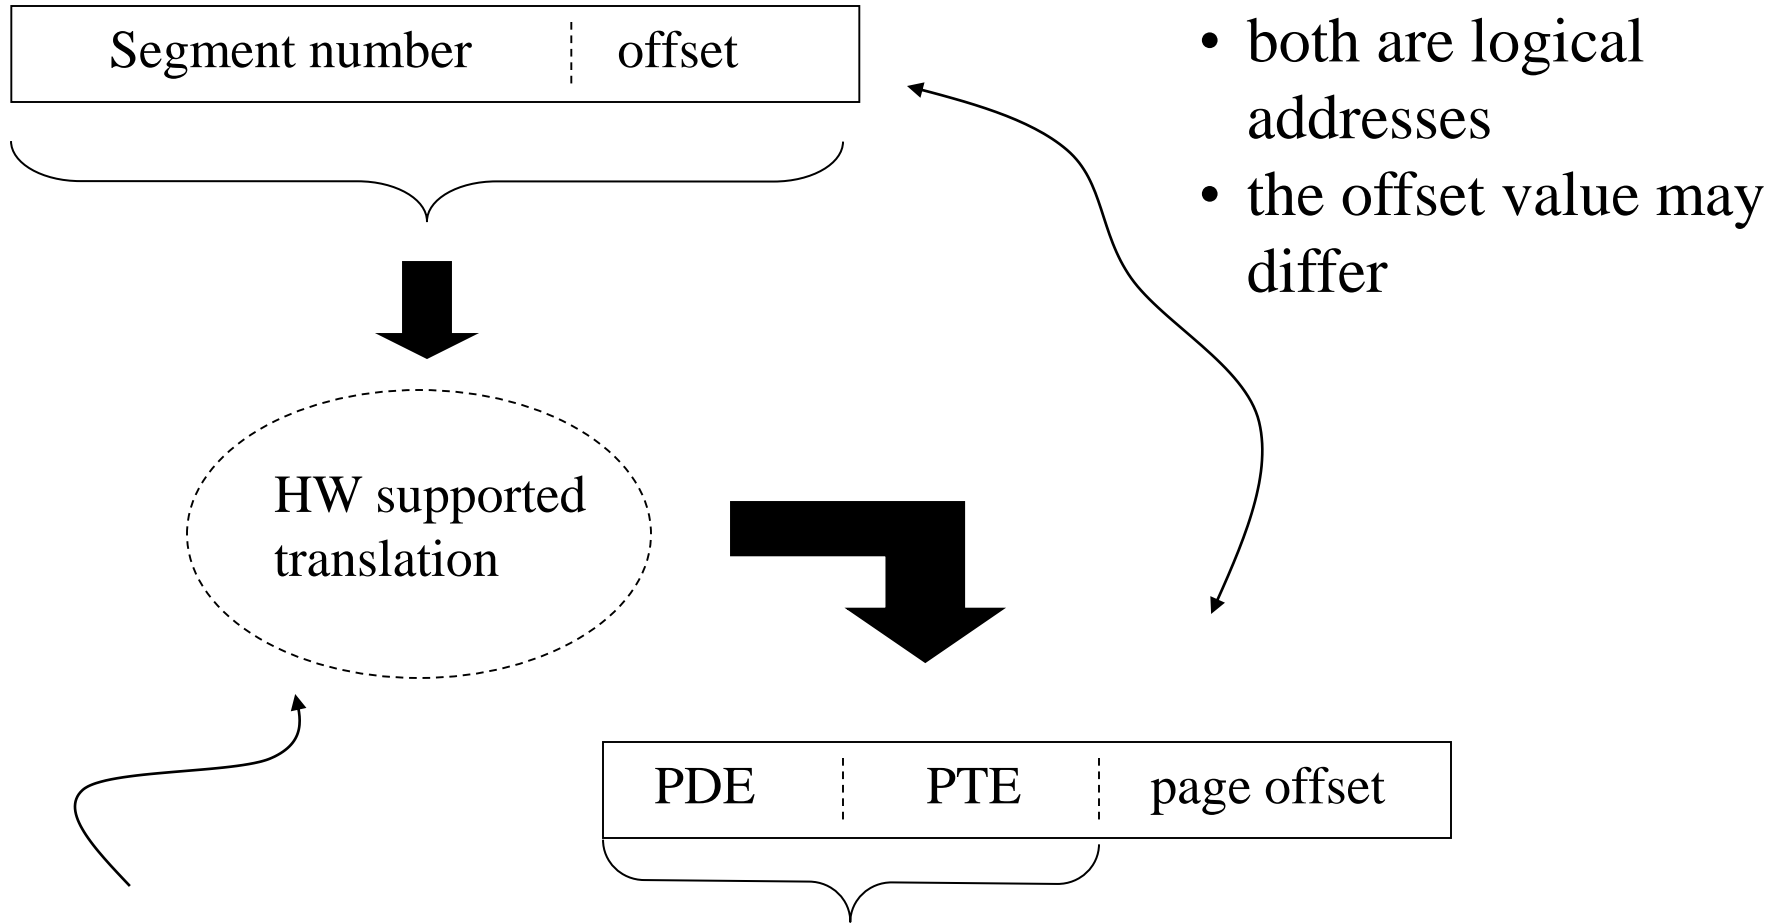
Determination of the linear address relying on <base,offset>

2-level paging example

# x86 memory access modes

- <u>Real mode</u>

  - ✓ Offers backward compatibility towards 286!!

  - ✓ a 16-bit segment register keeps the target segment ID

  - ✓ 16-bit (general) registers keep the segment offset

  - ✓ <u>Targeted addresses</u> are physical, and are computed as

    ```
    PhysicalAddress = Segment * 16 + Offset
    ```

  - ✓ Around 1MB ($2^{20}$B) of memory is allowed

  - ✓ Minimal support for separating chunks of memory in the addressing scheme

  - ✓ No segment specific protection information!!

  - ✓ Not suited for modern software systems!!!

# x86 memory access modes

- <u>80386 protected mode</u>

    - ✓ a 16-bit segment register keeps the target segment ID (using 13 bits)

    - ✓ 32-bit (general) registers keep the segment offset

    - ✓ The base of the segment in linear addressing is kept into a table in memory

    - ✓ <u>Targeted addresses</u> are linear and are computed as

        ```
        address = TABLE[segment].base + offset
        ```

    - ✓ Up to 4GB of linear (either physical or logical) memory is allowed

    - ✓ 3-bit for control (protection) are kept in the segment register …. much better for OS software!!!

# x86 memory access modes

- <u>Long mode (x86-64)</u>

    ✓ a 16-bit segment register keeps the target segment ID (using 13 bits)

    ✓ 64-bit (general) registers keep the segment offset (<u>limited to 48-bit global addressing in canonical form</u>)

    ✓ The base of the segment in linear addressing is kept into a table in memory

    ✓ <u>Targeted addresses</u> are linear and are computed as

    ```
    address = TABLE[segment].base + offset
    ```

    ✓ Up to 2^48 B (256 TB) of linear memory is allowed

    ✓ 3-bit for control (protection) are kept in the segment register

# x86 segment tables

- The are two table types keeping segments information: **Global Descriptor Table (GDT)** and **Local Descriptor Table (LDT)**

- Typically GDT and LDT are kept in main memory, and are directly accessible via pointers maintained by CPU registers

- GDT determines the mapping of linear addresses at least for kernel mode (namely kernel level segments) ... nowadays it is the unique used segment table in most operating systems

- LDT determines the mapping of linear addresses for user mode (namely user level segments), if not done via GDT

- These addresses are then used to access physical memory via page tables (if paging is activated)

# GDT organization

Segment protection and
usage rules

generic
entry

FLAGS

Segment base within
linear addressing

To be composed with
segment-offset upon access

# Segmentation vs paging

- Segmentation and paging typically have different targets

- Segmentation is a classical means for **protecting code and data**

- This protection mechanism is generally based on **coarse grain schemes** (in fact, segments may have very large sizes, covering up to the whole address space for the application)

- Paging (possibly coupled with virtual memory techniques) is generally employed as a means for **improving physical-memory management efficiency**

- Such "efficiency oriented" mechanism is based on a **fine-grain approach**, namely it relies on the size of the page frame for the specific hardware architecture (e.g. 4KB or 2/4MB for x86 architectures)

# Segmentation vs multi-cores/multi-threading

- … we know that paging schemes are still able to enforce protection of memory (via control bits in page-table entries)

- So we may think that segmentation is somehow useless in modern software systems

- This is a wrong concept, since as we will show <u>segmentation still plays a central role in multi-core architectures</u>

- <u>It also plays a central role in multi-thread programming</u>

- …… in 1985 paging was already there in the hardware but Intel further extended the segmentation support (e.g. in the 80386 processor)

- …. although the segmentation logic has been significantly revised in x86-64 processors

# The x86-64 revision

- Registers keeping track of segment IDs (also known as selectors) are not all managed the same way by firmware on board of the processor

- For some registers keeping segment IDs (hence for the corresponding segments in the GDT table) a fixed base of 0x0 is enforced for the segments

- Protection bits in the segment table entries associated with those segments IDs still work

- For a few registers keeping segment IDs the classical rule relying on arbitrary base values for the segments is adopted

# Segmentation based protection model (i)

- Each segment is associated with a given **protection level (or privilege level)**

- Each routine having protection level $h$ can invoke any other routine having protection level $h$, within any segment (**this can be achieved via intra-segment and cross-segment jumps**)

- Routines having protection level $h$ can invoke routines having protection level different from $h$ via **cross-segment jumps**

- **Cross-segment jumps always allow** jumping from protection level $h$ to protection level $h+i$

- Each segment having protection level $h$ is associated with a set of access points, called GATEs, each one identified as *<seg.id,offset>*

- Any GATE is associated with a maximum level *max=h+j* starting from which the GATE can be passed through

# Segmentation based protection model (ii)

- If *level(S)=h* and *max(GATE(S))=h+i* then segment *S* **entails a GATE** for accessing level *h* for modules associated with protection level up to *h+i*

- Cross-segment jumps <u>deny the access</u> to the destination if the source operates at protection level greater than the maximum one associated with the gate

- Overall, cross-segment jumps deny the access to the destination anytime we do not use a GATE as the destination *entry* for the jump

# Protection levels and jumps: the ring model

Level 2

Level 1

Level 0

———————▶ Always admitted

----------▶ Admitted depending on the *max* origin level associated with the target GATE

# An example

<S1, offset1>
(S1: level 0 – offset1: max = 0)

<S1, offset2>
(S1: level 0 – offset2: max = 3)

Kernel routine A

Kernel routine B

Admitted cross-segment jumps

Non-admitted cross-segment jump

User routine

S2 (level 2)
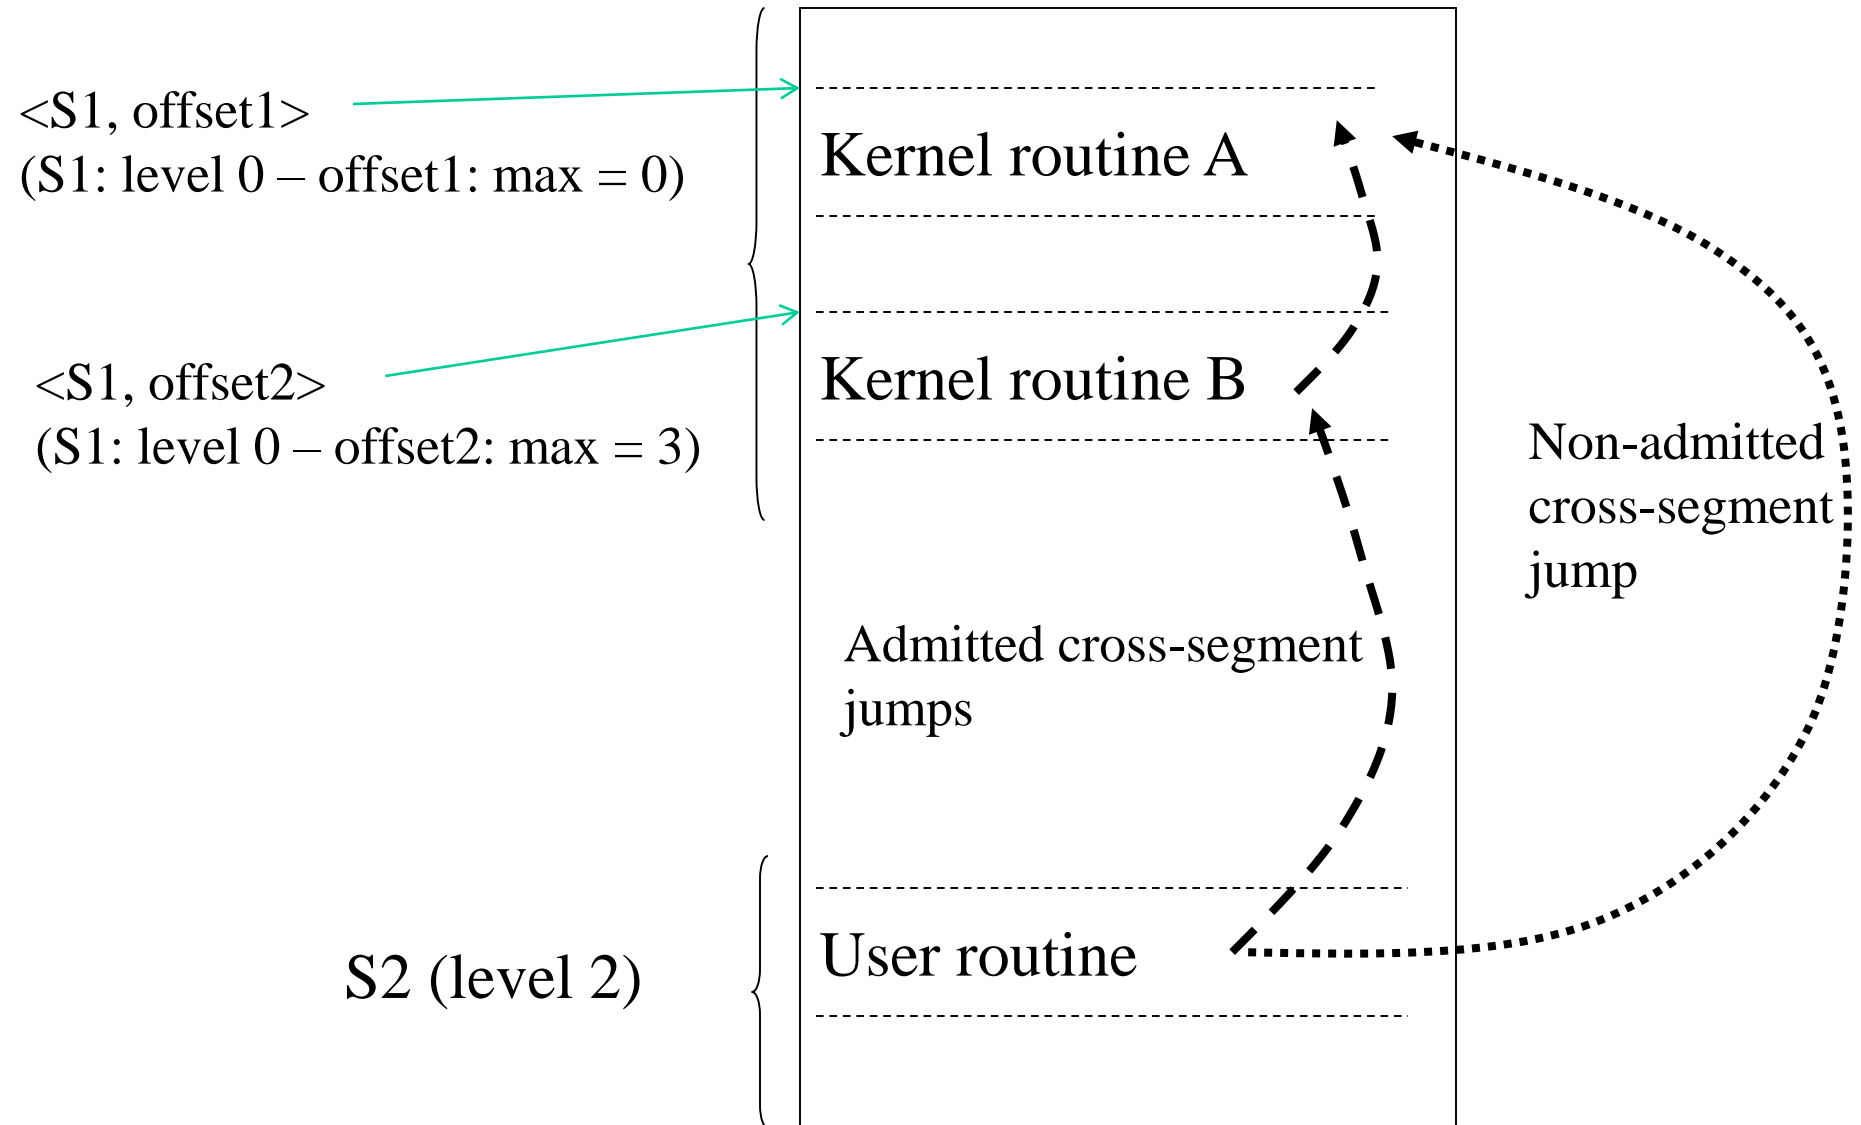
# Objectives of protection levels

- Denial of uncontrolled access to kernel level modules

- Kernel level access is controlled via specific "entry points" (the GATEs), which are explicitly used as destinations for jumps (more generally control flow variations) originated while running at worse protection levels

- In conventional operating systems, the entry points are typically associated with:

  ➢ **interrupt handlers** (asynchronous invocations)

  ➢ **software traps** (synchronous invocations)

# Ring scheme for x86 machines

**Protection Rings**

Operating System Kernel → Level 0

Operating System Services (Device Drivers, Etc.) → Level 1

Applications → Level 2

Level 3

Highest        Lowest

0    1    2    3

Privilege Levels
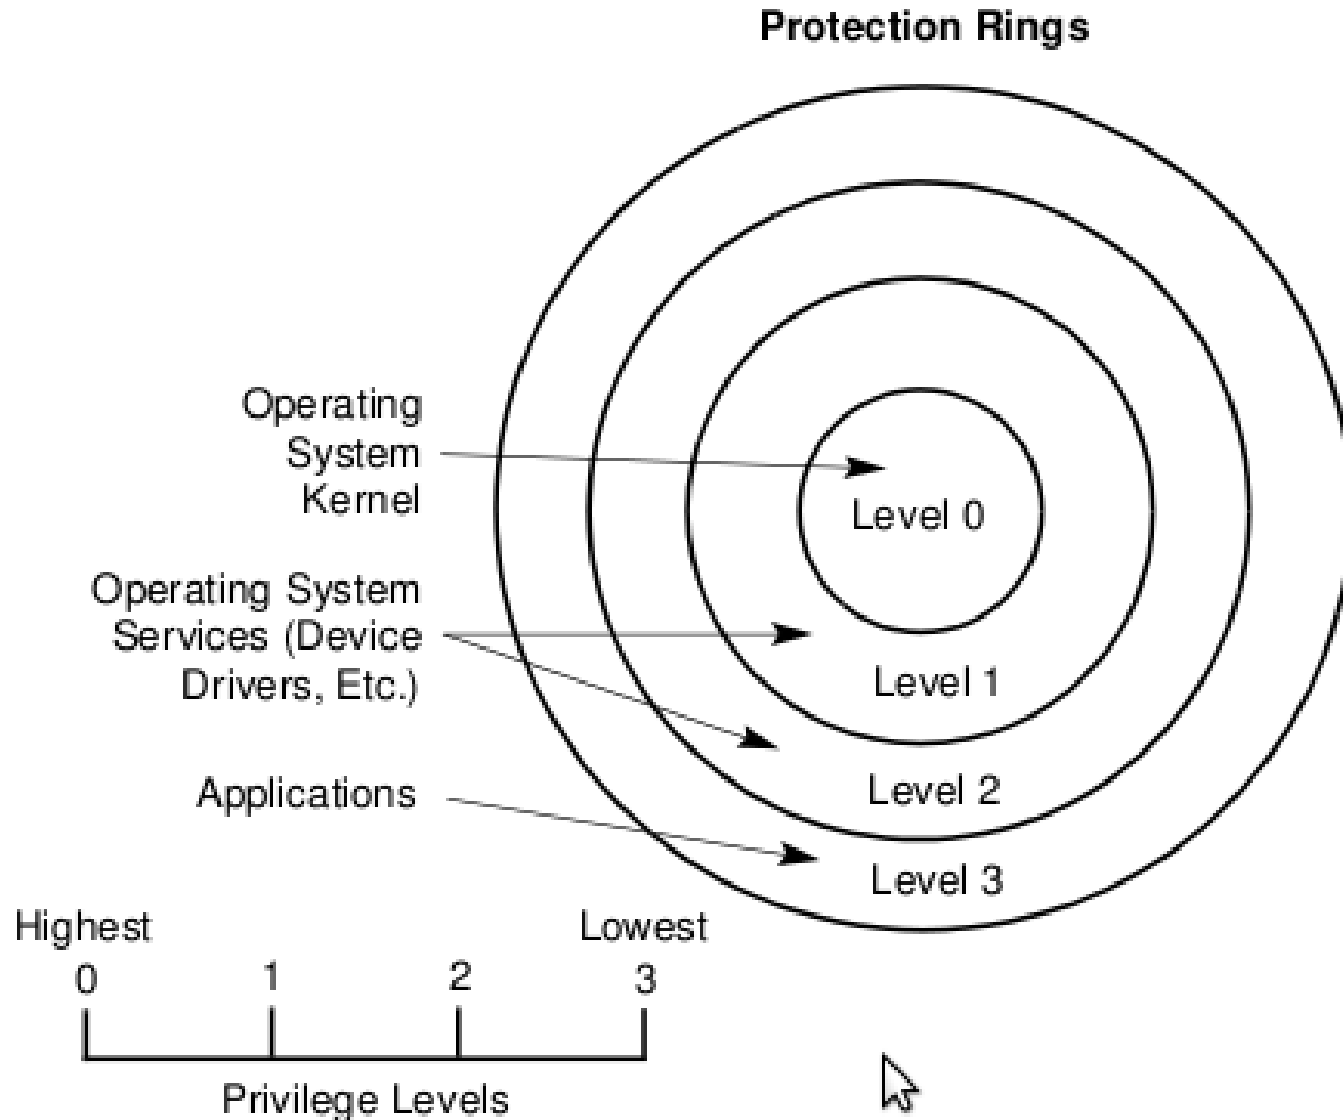
# x86 address composition with segmentation

- An address does not specify the segment ID directly

- It can specify a segment-selector register

- This register keeps information on the actual segment to which we are accessing

- An example:

```
<segment-selector-register,displacement>
```

# x86 details on the segmentation support



Segment Selector — 15 ... 3 2 1 0 — index | TI | RPL

TI = Table Indicator
RPL = Requestor Privilege Level

CS: code segment register
SS: stack segment register
DS: data segment register
ES: data segment register
FS: data segment register
GS: data segment register

added in 80386

For CS RPL is called **CPL**
This register is only writable by control flow variation instructions

- **CS (Code Segment Register)** points to the current segment. The 2 lsb identify the CPL (Current Privilege Level) for the CPU (from 0 to 3).

- **SS (Stack Segment Register)** points to the segment for the current stack.

- **DS (Data Segment Register)** points to the segment containing static and global data.

# Back to the very early x86 example

```
mov (%rax), %rbx

push %rbx
```

Here we are seamlessly (say implicitly) using <u>CS, and DS</u> for the first instruction and <u>CS and SS</u> for the second instruction

ES is an additional (to DS) implicit segment for specific classes of machine instructions, e.g. string-targeted ones like `stos` and `movs`

# x86 GDT entries (segment descriptors)

| 31 | 16 | 15 | 0 |
|----|----|----|---|
| Base 0:15 | | Limit 0:15 | |

| 63 | 56 | 55 | 52 | 51 | 48 | 47 | 40 | 39 | 32 |
|----|----|----|----|----|----|----|----|----|----|
| Base 24:31 | | Flags | | Limit 16:19 | | Access Byte | | Base 16:23 | |

This directly supports protected mode

**Access byte content:**
**Pr -** Present bit. This must be **1** for all valid selectors.
**Privl -** Privilege, 2 bits. Contains the ring level (0 to 3)
**Ex -** Executable bit (**1** if code in this segment can be executed)
…….

**Flags:**
**Gr -** Granularity bit. If **0** the limit is in 1 B blocks (byte granularity), if **1** the limit is in 4 KB blocks (page granularity)
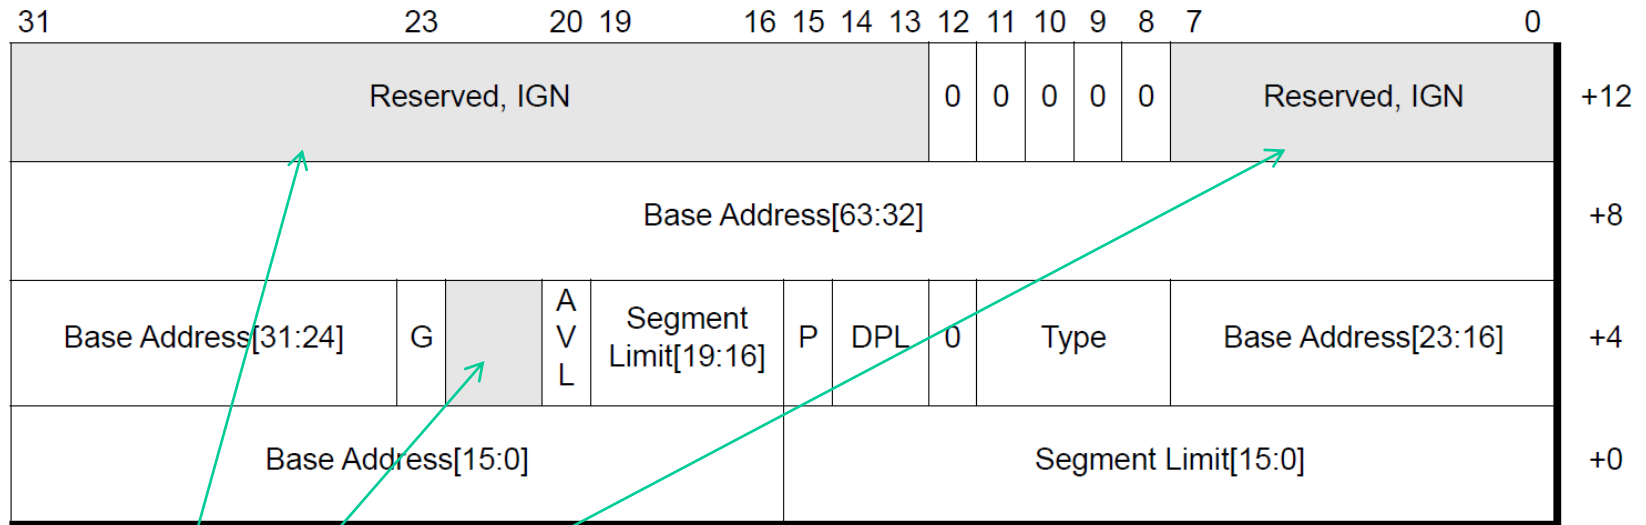
….

# Accessing GDT entries

- Given that a *segment descriptor* is 8 bytes in size, its relative address within GDT is computed by multiplying the 13 bits of the *index* field of *segment selector* by 8

- E.g, in case GDT is located at address 0x00020000 (value that is kept by the **gdtr register**) and the *index* value within *segment selector* is set to the value 2, the address associated with the *segment descriptor* is 0x00020000 + (2*8), namely 0x00020010

This is not only a pointer but actually a packed struct describing positioning and size of the GDT

# Long mode descriptors

| 31 | 23 | 20 19 | 16 15 14 13 | 12 | 11 | 10 | 9 | 8 | 7 | 0 | |
|---|---|---|---|---|---|---|---|---|---|---|---|
| Reserved, IGN | | | | 0 | 0 | 0 | 0 | 0 | Reserved, IGN | | +12 |
| Base Address[63:32] | | | | | | | | | | | +8 |
| Base Address[31:24] | G | | AVL | Segment Limit[19:16] | P | DPL | 0 | Type | Base Address[23:16] | | +4 |
| Base Address[15:0] | | | | Segment Limit[15:0] | | | | | | | +0 |

ignored bits

# Store Global Descriptor Table Register

| Opcode | Mnemonic | Description |
|--------|----------|-------------|
| 0F 01 /0 | SGDT m | Store GDTR to m. |

### Description

Stores the content of the global descriptor table register (GDTR) in the destination operand. The destination operand specifies a 6-byte memory location. If the operand-size attribute is 32 bits, the 16-bit limit field of the register is stored in the low 2 bytes of the memory location and the 32-bit base address is stored in the high 4 bytes. If the operand-size attribute is 16 bits, the limit is stored in the low 2 bytes and the 24-bit base address is stored in the third, fourth, and fifth byte, with the sixth byte filled with 0s.

SGDT is only useful in operating-system software; however, it can be used in application programs without causing an exception to be generated.

See "LGDT/LIDT-Load Global/Interrupt Descriptor Table Register" in Chapter 3 for information on loading the GDTR and IDTR.
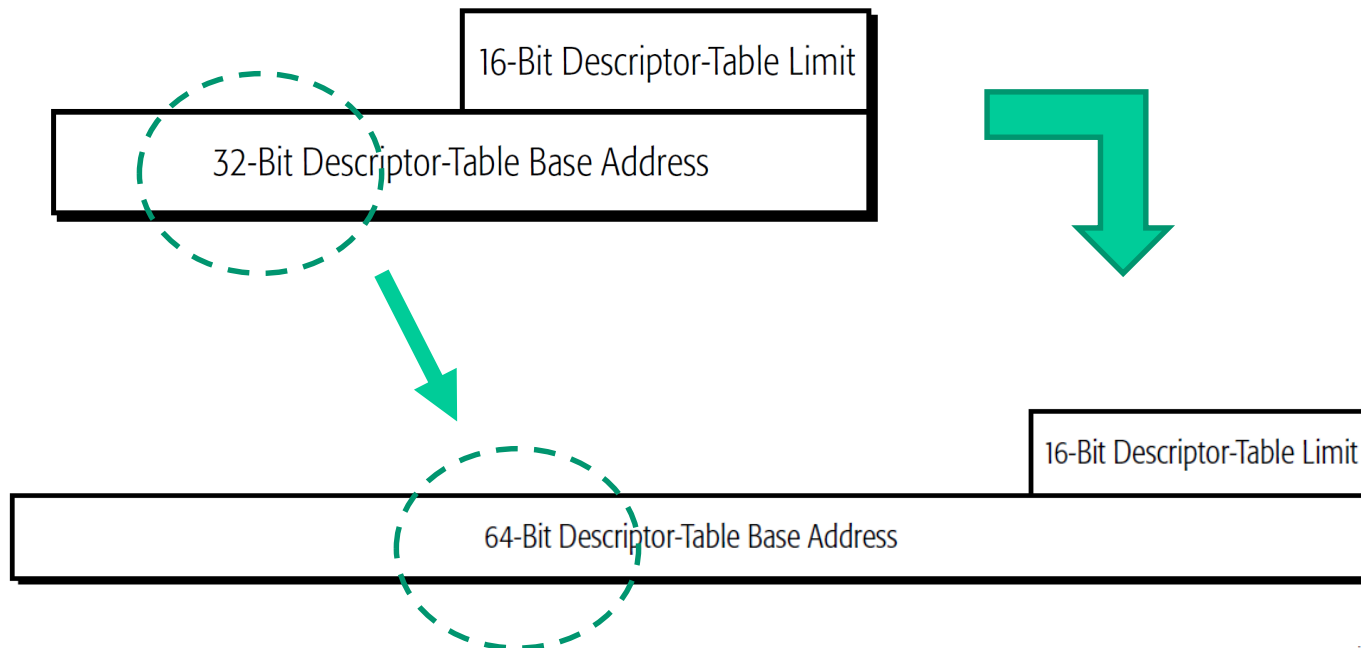
### Operation

```
if(OperandSize == 16) {
        Destination[0..15] = GDTR.Limit;
        Destination[16..39] = GDTR.Base; //24 bits of base address loaded
        Destination[40..47] = 0;
}
else { //32-bit Operand Size
        Destination[0..15] = GDTR.Limit;
        Destination[16..47] = GDTR.Base; //full 32-bit base address loaded
}
```

### IA-32 Architecture Compatibility

The 16-bit form of the SGDT is compatible with the Intel 286 processor if the upper 8 bits are not referenced. The Intel 286 processor fills these bits with 1s; the Pentium 4, Intel Xeon, P6 family, Pentium, Intel486, and Intel386 processors fill these bits with 0s.

x86 long mode provides 2 (the table size) + 8 (the table address) bytes

# Long mode GDTR extensions

16-Bit Descriptor-Table Limit

32-Bit Descriptor-Table Base Address

16-Bit Descriptor-Table Limit

64-Bit Descriptor-Table Base Address

# Example code

```c
#include <stdio.h>

struct desc_ptr {
        unsigned short size;
        unsigned long address;
} __attribute__((packed)) ;

#define store_gdt(ptr) asm volatile("sgdt %0":"=m"(*ptr))

int main (int argc, char**argv){
 struct desc_ptr gdtptr;
 char v[10];//another way to see 10 bytes packed in memory

 store_gdt(&gdtptr);
 store_gdt(v);

 printf("comparison is %d\n",memcmp(v,&gdtptr,10));
 printf("GDTR is at %x - size is %d\n",gdtptr.address, gdtptr.size);
 printf("GDTR is at %x - size is %d\n",((struct desc_ptr*)v)->address,
        ((struct desc_ptr*)v)->size);

}
```
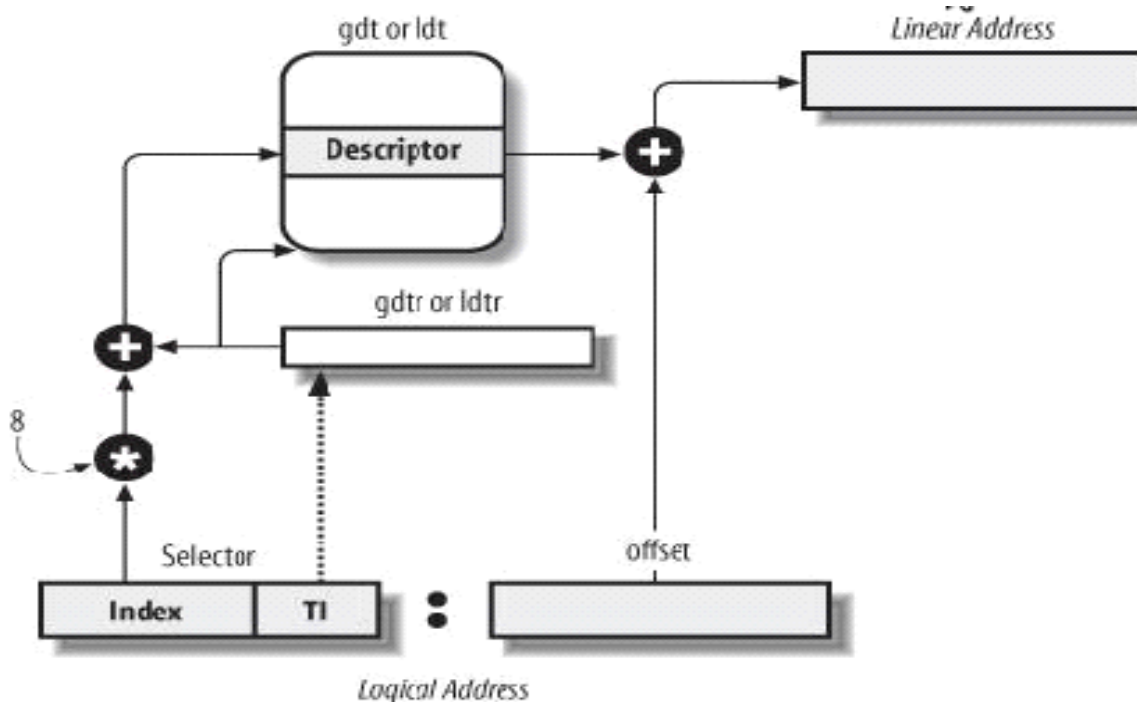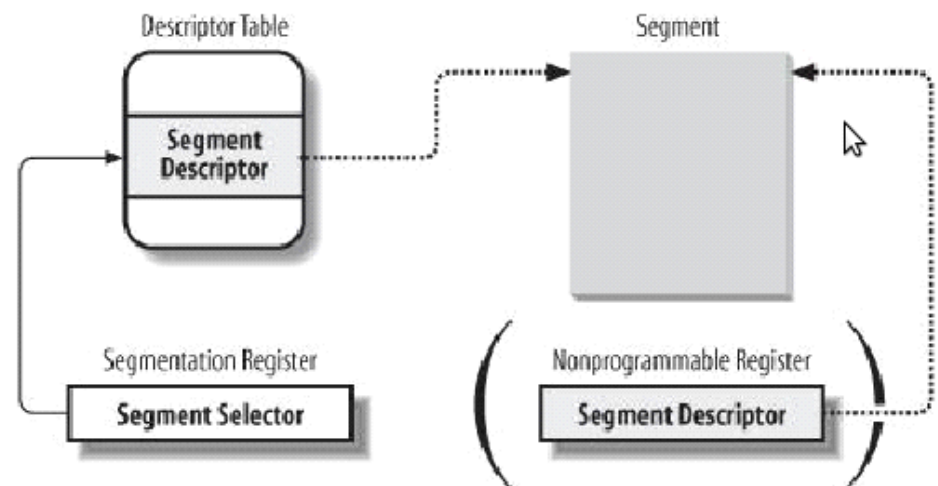
# Access scheme



Caching of descriptors
(1 cache register per segment
  selector – non-programmable)

Cache line filled upon  selector
update

# Making explicit usage of segments while coding

```c
#include <stdio.h>

#define load(ptr,var) asm volatile("mov %%ds:(%0), %%rax":"=a" (var):"a" (ptr))
#define store(val,ptr) asm volatile(" mov %0, %%ds:(%1)"\
                                     ::"a" (val), "b" (ptr):)

int main (int argc, char**argv){

        unsigned long x = 16;

        unsigned long y;

        load(&x,y);
        printf("variable y has value %u\n",y);

        store(y+1,&x);
        printf("variable x has value %u\n",x);

}
```
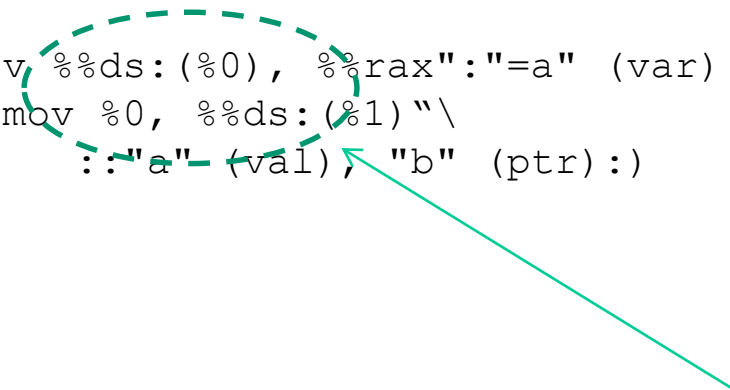
explicit reference to the data segment register (DS)

# Code/data segments for LINUX

| Segment | Base | G | Limit | S | Type | DPL | D/B | P |
|---|---|---|---|---|---|---|---|---|
| user code | 0x00000000 | 1 | 0xfffff | 1 | 10 | 3 | 1 | 1 |
| user data | 0x00000000 | 1 | 0xfffff | 1 | 2 | 3 | 1 | 1 |
| kernel code | 0x00000000 | 1 | 0xfffff | 1 | 10 | 0 | 1 | 1 |
| kernel data | 0x00000000 | 1 | 0xfffff | 1 | 2 | 0 | 1 | 1 |

Can we read/write/execute?

Is the segment present?

x86-64 directly forces base to 0x0 for
the corresponding segment registers

# x86-64 selector management details

CS
SS
DS
ES

$\longrightarrow$ Base = 0x0

Privilege level is still there and working

FS
GS

$\longrightarrow$ Arbitrary Base
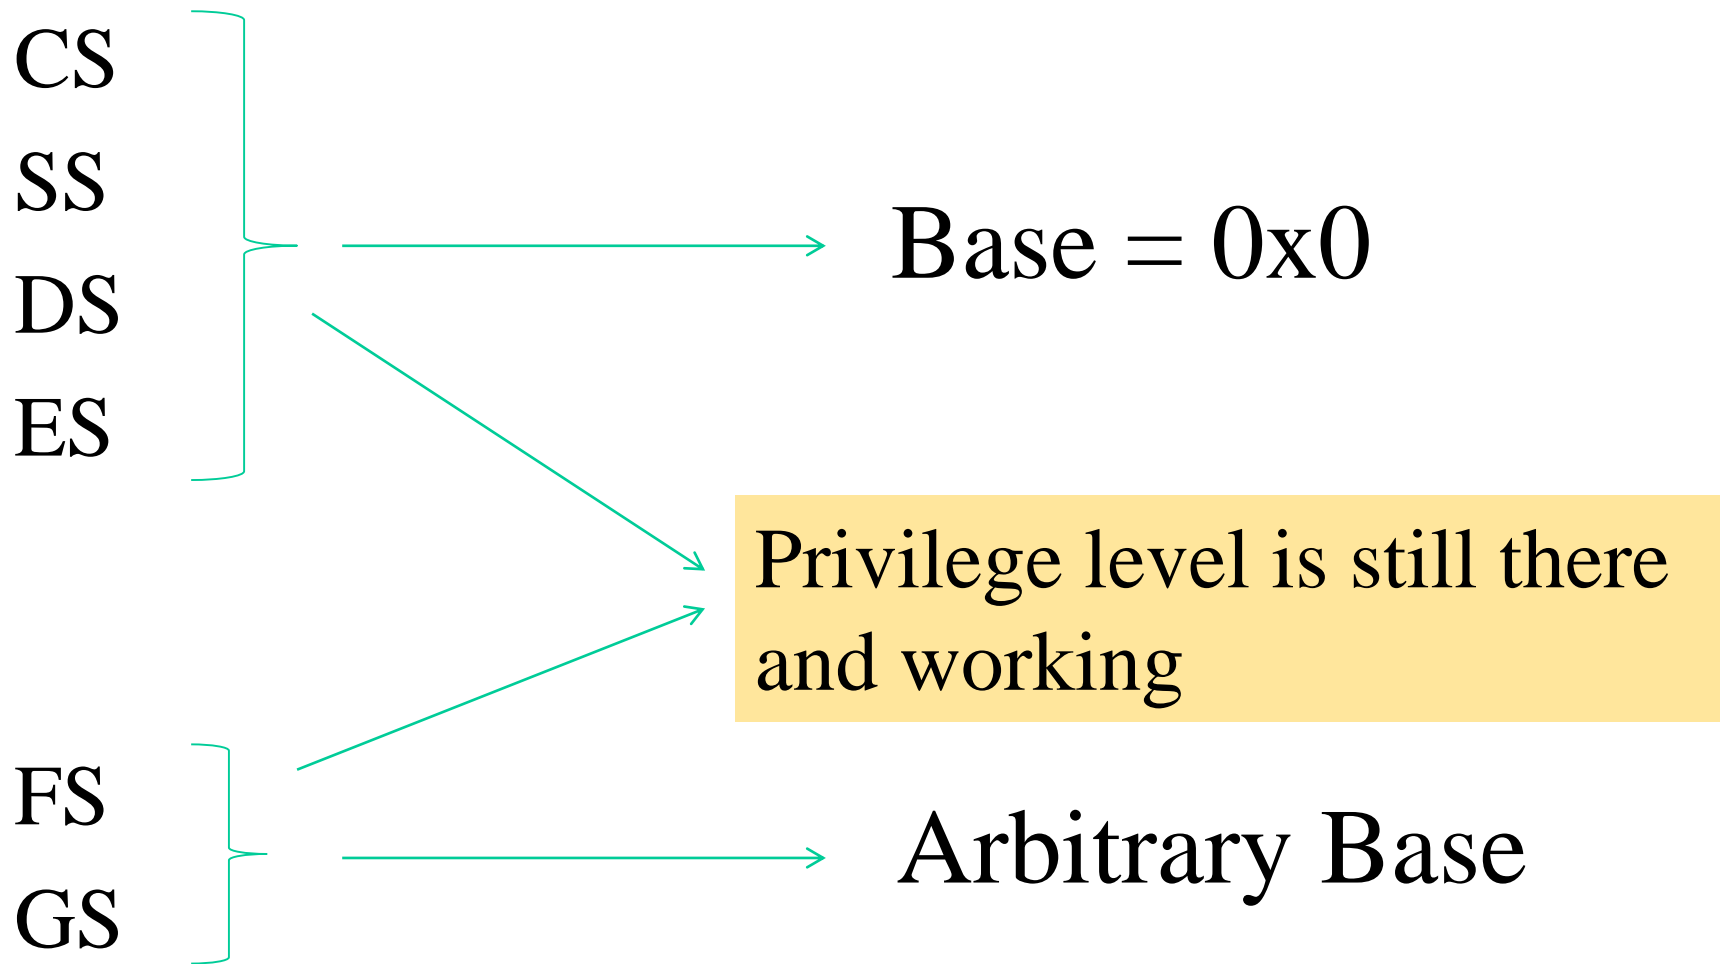
# Segment selectors update rules

- CS plays a central role, since it keeps the CPL (Current Privilege level)

- CS is only updated via control flow variations

- All the other segment registers can be updated if the segment descriptor they would point to after the update has DPL $\geq$ CPL

- Clearly, **with CPL = 0 we can update everything (ring 0 has no limit)**

# LINUX GDT on x86

| Linux's GDT | Segment Selectors | | Linux's GDT | Segment Selectors |
|---|---|---|---|---|
| null | 0x0 | | TSS | 0x80 |
| reserved | | | LDT | 0x88 |
| reserved | | | PNPBIOS 32-bit code | 0x90 |
| reserved | | | PNPBIOS 16-bit code | 0x98 |
| not used | | | PNPBIOS 16-bit data | 0xa0 |
| not used | | | PNPBIOS 16-bit data | 0xa8 |
| TLS #1 | 0x33 | | PNPBIOS 16-bit data | 0xb0 |
| TLS #2 | 0x3b | | APMBIOS 32-bit code | 0xb8 |
| TLS #3 | 0x43 | | APMBIOS 16-bit code | 0xc0 |
| reserved | | | APMBIOS data | 0xc8 |
| reserved | | | not used | |
| reserved | | | not used | |
| kernel code | 0x60 (__KERNEL_CS) | | not used | |
| kernel data | 0x68 (__KERNEL_DS) | | not used | |
| user code | 0x73 (__USER_CS) | | not used | |
| user data | 0x7b (__USER_DS) | | double fault TSS | 0xf8 |

Beware these

# TSS

- TSS (Task State Segment): the set of linear addresses associated with TSS is a subset of the linear address space destined to kernel data segment

- each TSS (one per CPU-core) is kept within the **int_tss** array

- the *Base* field within the $n$-th core TSS register points to the $n$-th entry of the **int_tss** array (transparently via the TSS segment)

- Gr=0 while *limit is 104 * 4 bytes*

- *DPL*=0, since the TSS segment cannot be accessed in user mode

# x86 TSS structure

| 31 | 15 | 0 | |
|---|---|---|---|
| I/O Map Base Address | | T | 100 |
| | LDT Segment Selector | | 96 |
| | GS | | 92 |
| | FS | | 88 |
| | DS | | 84 |
| | SS | | 80 |
| | CS | | 76 |
| | ES | | 72 |
| EDI | | | 68 |
| ESI | | | 64 |
| EBP | | | 60 |
| ESP | | | 56 |
| EBX | | | 52 |
| EDX | | | 48 |
| ECX | | | 44 |
| EAX | | | 40 |
| EFLAGS | | | 36 |
| EIP | | | 32 |
| CR3 (PDBR) | | | 28 |
| | SS2 | | 24 |
| ESP2 | | | 20 |
| | SS1 | | 16 |
| ESP1 | | | 12 |
| | SS0 | | 8 |
| ESP0 | | | 4 |
| | Previous Task Link | | 0 |

▢ Reserved bits. Set to 0.

Although it could be ideally used for hardware based context switches, it is not in Linux/x86

It is essentially used for privilege level switches (e.g. access to kernel mode), based on stack differentiation

# x86-64 variant

| offset | 31-16 | 15-0 |
|--------|-------|------|
| 0x00 | reserved | |
| 0x04 | RSP0 (low) | |
| 0x08 | RSP0 (high) | |
| 0x0C | RSP1 (low) | |
| 0x10 | RSP1 (high) | |
| 0x14 | RSP2 (low) | |
| 0x18 | RSP2 (high) | |

room for 64-bit
stack pointers has been created
sacrificing general registers
snapshots

# Loading the TSS register

- x86 ISA (Instruction Set Architecture) offers the instruction `LTR`

- This is privileged and must be executed at CPL = 0

- The TSS descriptor must be filled with a source operand

- The source can be a general-purpose register or a memory location

- Its value (16 bits) keeps the index of the TSS descriptor into the GDT

# LTR — Load Task Register

| Opcode | Instruction | Op/En | 64-Bit Mode | Compat/Leg Mode | Description |
|--------|-------------|-------|-------------|-----------------|-------------|
| 0F 00 /3 | LTR r/m16 | M | Valid | Valid | Load r/m16 into task register. |

## Instruction Operand Encoding ¶

| Op/En | Operand 1 | Operand 2 | Operand 3 | Operand 4 |
|-------|-----------|-----------|-----------|-----------|
| M | ModRM:r/m (r) | NA | NA | NA |

## Description ¶

Loads the source operand into the segment selector field of the task register. The source operand (a general-purpose register or a memory location) contains a segment selector that points to a task state segment (TSS). After the segment selector is loaded in the task register, the processor uses the segment selector to locate the segment descriptor for the TSS in the global descriptor table (GDT). It then loads the segment limit and base address for the TSS from the segment descriptor into the task register. The task pointed to by the task register is marked busy, but a switch to the task does not occur.

The LTR instruction is provided for use in operating-system software; it should not be used in application programs. It can only be executed in protected mode when the CPL is 0. It is commonly used in initialization code to establish the first task to be executed.

The operand-size attribute has no effect on this instruction.

# GDT replication

- By the discussion on TSS we might have already observed that different CPU-cores in a multi-core/multi-processor system may need to fill a given entry of the GDT with different values

- To achieve this goal the GDT is actually replicated in common operating systems, with one copy for each CPU-core

- Then each copy slightly diverges in a few entries

- The main (combined) motivations are
    - ✓ performance
    - ✓ transparency of data access separation

# Actual architectural scheme

RAM memory

CPU-core 0

gdtr

gdtr

CPU-core 1

The two tables may differ in a few entries!!

# Replication benefits: per-CPU seamless memory accesses



RAM memory

`gdtr`

`gdtr`

CPU-core 0

CPU-core 1

GS segment = X

GS segment = X

Base is B

Base is B'

Same displacement within segment X seamlessly leads the two CPU-cores to access different linear addresses

# Per-CPU memory

- No need for a CPU-core to call, e.g. CPUID (… devastating for the speculative pipeline …) to determine what memory portion is explicitly dedicated to it

- Fast access via GS segment displacing for per-CPU common operations such as

  ✓ Statistics update (non need for LOCKED CMPXCHG)

  ✓ Fast control operations

# Per-CPU memory setup in Linux

- Based on some per-CPU reserved zone in the linear addressing scheme

- The reserved zone is displaced by relying on the GS segment register

- Based on macros that select a displacement in the GS segment

- Based on macros that implement memory access relying on the selected displacement

# An example

```
DEFINE_PER_CPU(int, x);

int z;

z = this_cpu_read(x);
```

The above  statement results in a single instruction:

```
        mov ax, gs:[x]
```

To operate with no special define we can also get the actual address of the per-CPU data and work normally:

```
        y = this_cpu_ptr(&x)
```

# TLS – Thread Local Storage

- It is based on setting up different segments associated with FS and GS selectors

- Each time a thread is CPU-dispatched, kernel software restores its corresponding segment descriptors into TLS#1, TLS#2 and TLS#3 within the GDT

- We have system calls allowing us to change the segment descriptors to be posted on TLS entries

# Segment management system calls (i)

arch_prctl - set architecture-specific thread state

```
#include <asm/prctl.h>
#include <sys/prctl.h>

int arch_prctl(int code, unsigned long addr);
int arch_prctl(int code, unsigned long *addr);
```

**arch_prctl**() sets architecture-specific process or thread state.
*code* selects a subfunction and passes argument *addr* to it; *addr* is
interpreted as either an *unsigned long* for the "set" operations, or
as an *unsigned long* *, for the "get" operations.

Subfunctions for x86-64 are:

# Segment management system calls (ii)

Subfunctions for x86-64 are:

**ARCH_SET_FS**
Set the 64-bit base for the *FS* register to *addr*.

**ARCH_GET_FS**
Return the 64-bit base value for the *FS* register of the
current thread in the *unsigned long* pointed to by *addr*.

**ARCH_SET_GS**
Set the 64-bit base for the *GS* register to *addr*.

**ARCH_GET_GS**
Return the 64-bit base value for the *GS* register of the
current thread in the *unsigned long* pointed to by *addr*.

## RETURN VALUE   top

On success, **arch_prctl**() returns 0; on error, -1 is returned, and
*errno* is set to indicate the error.

# x86-64 control registers

- CR0-CR3 or CR0-CR4 (on more modern x86 CPUs)
- CR0: is the baseline one
- CR1: is reserved
- CR2: keeps the linear address in case of a fault
- CR3: is the page-table pointer

# CR0 structure vs long mode

Long mode uses a combination of this and the EFER (Extended Feature Enable Register) MSR (model specific register)

| Bit | Name | Full Name | Description |
|-----|------|-----------|-------------|
| 0 | PE | Protected Mode Enable | If 1, system is in protected mode, else system is in real mode |
| 1 | MP | Monitor co-processor | Controls interaction of WAIT/FWAIT instructions with TS flag in CR0 |
| 2 | EM | Emulation | If set, no x87 FPU is present, if clear, x87 FPU is present |
| 3 | TS | Task switched | Allows saving x87 task context upon a task switch only after x87 instruction used |
| 4 | ET | Extension type | On the 386, it allowed to specify whether the external math coprocessor was an 80287 or 80387 |
| 5 | NE | Numeric error | Enable internal x87 floating point error reporting when set, else enables PC style x87 error detection |
| 16 | WP | Write protect | When set, the CPU can't write to read-only pages when privilege level is 0 |
| 18 | AM | Alignment mask | Alignment check enabled if AM set, AC flag (in EFLAGS register) set, and privilege level is 3 |
| 29 | NW | Not-write through | Globally enables/disable write-through caching |
| 30 | CD | Cache disable | Globally enables/disable the memory cache |
| 31 | PG | Paging | If 1, enable paging and use the CR3 register, else disable paging |

# Interrupts/traps vs kernel access

- Interrupts are **asynchronous events** that are not correlated with the current CPU-core execution flow

- Interrupts are generated by external devices, and can be masked (vs non-masked)

- Traps, also known as **exceptions**, are **synchronous events**, strictly coupled with the current CPU-core execution (e.g. division by zero)

- Multiple executions of the same program, under the same input, may (but not necessarily do) give rise to the same exceptions

- Traps are (actually have been historically) used as the mechanism for on demand access to kernel mode (via system calls)

# Management of trap/interrupt events

- The kernel keeps a **trap/interrupt table**

- Each table entry keeps a **GATE descriptor**, which provides information on the address associated with the GATE (e.g. <seg.id,offset>) and the GATE protection level

- The content of the trap/interrupt table is exploited to determine whether the access to the GATE can be enabled

- The check relies on the current content of CPU registers, the segment registers, which specify the current privilege level (CPL)

- In principle, it may occur that a given GATE **is described within multiple entries** of the trap/interrupt table (aliasing), possibly with different protection specifications

# Summary on x86 control flow variations

- **intra-segment**: standard jump instruction (e.g. JMP <displacement> on x86 architectures)
  - ➢ firmware only verifies whether the displacement is within the current segment boundary
- **cross-segment**: long jump instructions (e.g. LJMP <seg.id>, <displacement> on x86 architectures)
  - ➢ Firmware verifies whether jump is enabled on the basis of privilege levels (no CPL improvement is admitted)
  - ➢ Then, firmware checks whether the displacement is within the segment boundaries
- **cross-segment via GATEs**: trap instructions (e.g. INT <table displacement> on x86 architectures)
  - ➢ Firmware checks whether jumping is admitted depending on the privilege level associated with the target GATE as specified within the **trap/interrupt table**

# An overview

Always admitted
(requires anyway consulting
 the segment Tables)

Not always admitted
(requires consulting the
 Trap/interrupt table
 +
Segment Tables)

Move across
segments

Seg 0 – level = 0

Seg 1 – level ≥ 0

Seg i – level ≤ n

# GATE details for the x86 architecture (i)

- The trap/interrupt table is called **Interrupt Descriptor Table (IDT)**

- Any entry keeps
  - ➢ The ID of the target segment and the segment displacement
  - ➢ the *max* level starting from which the access to the GATE is granted

- IDT is accessible via the `idtr` register which is a packed structure keeping the linear address of the IDT and the size (number of entries, each made up by 8 or 16 bytes, depending on whether extended 64-bit mode is active)

- The register is loadable via the LIDT machine instruction

# GATE details for the x86 architecture (ii)

- We know the **<u>current privilege level</u>** is kept within CS

- If protection information enables jumping, the segment ID within IDT is used to access GDT in order to check whether jumping is within the segment boundaries

- If check succeeds the current privilege level gets updated

- The new value is taken from the **<u>corresponding entry</u>** of GDT (this value corresponds to the privilege level of the target segment)

- The GATE description also tells whether the activated code is interruptible or not

# Conventional operating systems

- For LINUX/Windows systems, the GATE for on-demand access (via software traps) to the kernel **is unique**

- For i386 machines the corresponding software traps are

  ➢ INT 0x80 for LINUX (with backward compatibility in x86-64)

  ➢ INT 0x2E for Windows

- Any other GATE is reserved for the management of run-time errors (e.g. divide by zero exceptions) and interrupts

- They are not usable for on-demand access via software (clearly except if you hack the kernel)

- The software module associated with the on-demand access GATE implements **a dispatcher that is able to trigger the activation of the specific system call** targeted by the application

# Data structures for system call dispatching

- There exists a **"system call table"** that keeps, in any entry, the address of a specific system call

- Such an address becomes the target for a subroutine activation by the dispatcher

- To access the correct entry, the dispatcher gets in input the **number (the numerical code – the index) of the target system call** (typically this input is provided within a CPU register)

- The code is used to identify the target entry within the system call table

- Then the dispatcher invokes the system call routine (as a "jump sub-routine" – CALL instruction on x86)

- The actual system call, once executed, provides its output (return) value within a CPU register

# The trap-based dispatching scheme

**User level**

define input and
access GATE (trap)

retrieve system call
return value

**Kernel level**

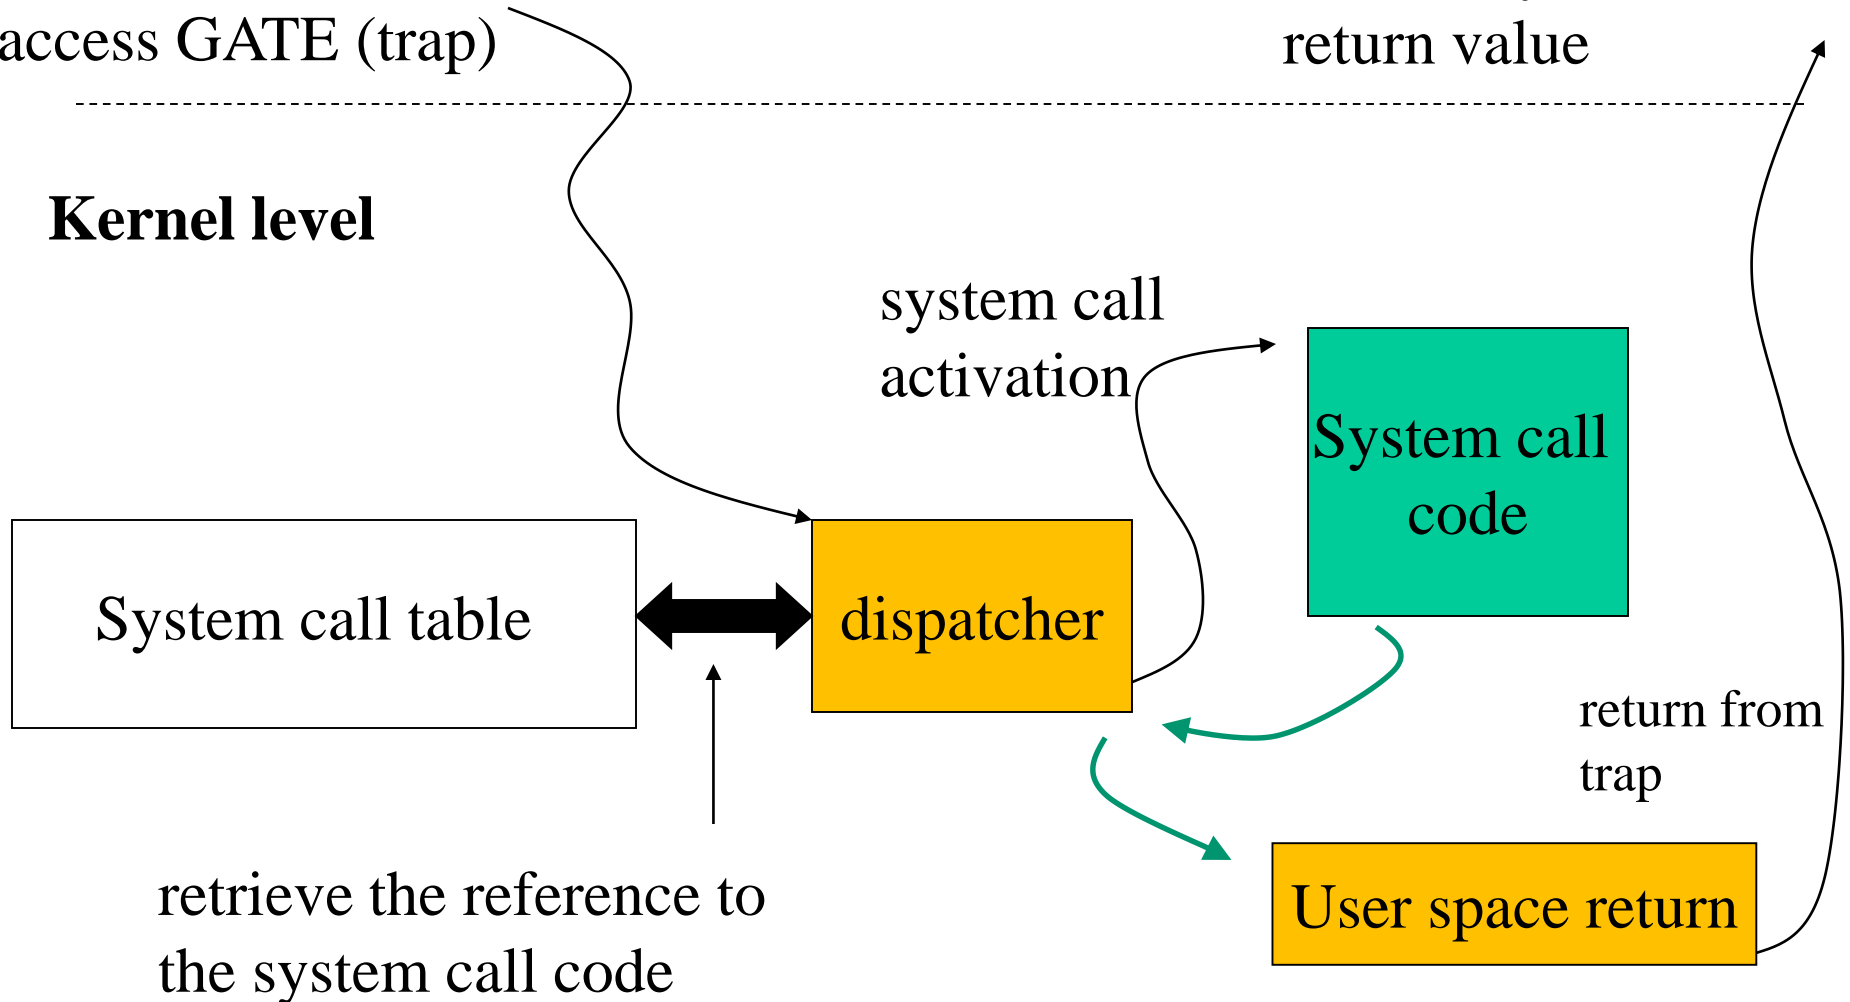system call
activation

System call
code

System call table ⬌ dispatcher

return from
trap

User space return

retrieve the reference to
the system call code

# Trap vs interruptible execution

- Differently from interrupts, **<u>trap</u>** management is typically configured so as not to entail/enable automatically resetting the interruptible-state for the CPU-core

- Any critical code portion associated with the management of the trap within the kernel requires explicit set of the interruptible-state bit, and the reset after job is complete (e.g. via CLI e STI instructions in x86 processors)

- For SMP/multi-core machines this **<u>may not suffice</u>** for guaranteeing correctness (e.g. atomicity) while handling the trap

- To address this issue, spinlock mechanisms are adopted, which are base on atomic **test-end-set code portions** (e.g., generated via the x86 LOCK prefix on standard compilation tool chains)

# Test-and-set support

- Modern instruction sets offer a single instruction to atomically test-and-set memory, this is the CAS (Compare And Swap) intruction

- On x86 machines the actual CAS is called CMPXCHG (Compare And Exchange)

- ... but we already discussed of this while dealing with memory consistency!!

# System call software components

- User side: software module (a) providing the input parameters to the GATE (and to the actual system call) (b) activating the GATE and (c) recovering the system call return value

- kernel side:
  - ➢ dispatcher
  - ➢ system call table
  - ➢ actual system call code
- Addition of a new **system call** means working on both sides

- Typically, this happens with no intervention on the dispatcher **in all the cases where the system call format is compliant with those predefined for the target operating system**

# Linux along our path

- Kernel 2.4 : highly oriented to expansibility modifiability
- Kernel 2.6: more scalable
- Kernel 3.x: more structured and secure
- Kernel 4.x, 5.x: even more secure

# LINUX system calls support: the whole (long) history up to kernel 5.x

# System call indexing

- This has changes along time

- We originally had the so called UNISTD_32 indexing scheme

- This is still supported in modern kernel versions (e.g. 4.x and 5.x)

- Now we have the UNISTD_64 indexing

- Given that the system call indexes are used/needed at user space, we can exploit them for user code programming via the `/usr/include/asm` directory (or `/usr/include/x86_64-linux-gnu/asm`)

- The two indexing schemes are stated in
    - `unistd_32.h`
    - `unistd_64.h`

- Two indexing schemes imply two different system call tables at kernel level, which coexist with each other (and of course two dispatchers)

# UNISTD_32 listing

```
#ifndef _ASM_X86_UNISTD_32_H
#define _ASM_X86_UNISTD_32_H 1

#define __NR_restart_syscall 0
#define __NR_exit 1
#define __NR_fork 2
#define __NR_read 3
#define __NR_write 4
#define __NR_open 5
#define __NR_close 6
#define __NR_waitpid 7
#define __NR_creat 8
#define __NR_link 9
#define __NR_unlink 10
#define __NR_execve 11
#define __NR_chdir 12
#define __NR_time 13
#define __NR_mknod 14
#define __NR_chmod 15
#define __NR_lchown 16
......
```

# UNISTD_64 listing

```
#ifndef _ASM_X86_UNISTD_64_H
#define _ASM_X86_UNISTD_64_H 1

#define __NR_read 0
#define __NR_write 1
#define __NR_open 2
#define __NR_close 3
#define __NR_stat 4
#define __NR_fstat 5
#define __NR_lstat 6
#define __NR_poll 7
#define __NR_lseek 8
#define __NR_mmap 9
#define __NR_mprotect 10
#define __NR_munmap 11
#define __NR_brk 12
#define __NR_rt_sigaction 13
#define __NR_rt_sigprocmask 14
......
```

# User level tasks for accessing the gate GATE

1.  **Specification of the input parameters via CPU registers** (note that these include the actual system call parameters and the dispatcher ones)

2.  ASM instructions triggering the GATE (e.g. traps)

3.  Recovery of the return value of the systems call (upon returning from the trap associated with GATE activation)

# Predefined system call formats

- These ae specified in header files that enable using GATE access functions in C

- **Thes eheadre files define the standard formats for the user level module triggering access to the system GATE** (namely the module that activates the system call dispatcher), each for a different value of the number of system call parameters (from 0 to 6)

- Essentially these header files contain **ASM vs C directives** and architecture specific compilation directives

- They represent a meeting point between standard C programming and machine specific ASM language (in relation to the GATE access functionality)

# Code block for a standard system call with no parameter (e.g. fork()) – classical define style

```
#define _syscall0(type,name) \
type name(void) \
{ \
long __res; \
__asm__ volatile ("int $0x80" \
        : "=a" (__res) \
        : "0" (__NR_##name)); \
__syscall_return(type,__res); \
}
```
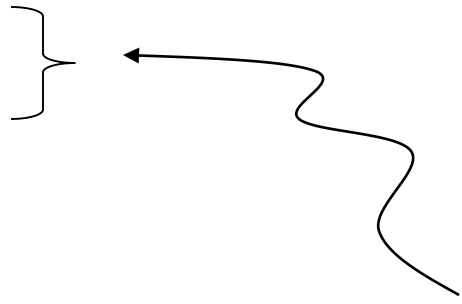
Assembler instructions

Tasks to be done after the execution of the assembler code block

Tasks preceding the assembler code block

# Managing the return value and `errno`

```
/* user-visible error numbers are in the range -1 - -124:
  see <asm-i386/errno.h> */


#define __syscall_return(type, res) \
do { \
        if ((unsigned long)(res) >= (unsigned long)(-125)) { \
                errno = -(res); \
                res = -1; \
        } \
        return (type) (res); \
} while (0)
```

Case of `res` within the interval [–1, -124]

# Note: why the do/while(0) construct?

It is a C construct that allows to
- #define a multi-statement operation
- put a semicolon after and
- still use within an **if** statement

# Code block for a standard system call with one parameter (e.g. close()) – classical define style

```
#define _syscall1(type,name,type1,arg1) \
type name(type1 arg1) \
{ \
long __res; \
__asm__ volatile ("int $0x80" \
    : "=a" (__res) \
    : "0" (__NR_##name),"b" ((long)(arg1))); \
__syscall_return(type,__res); \
}
```

2 registers used for the input

# Code block for a system call with six parameters – classical define style (tailored to UNISTD_32)

```
#define _syscall6(type,name,type1,arg1,type2,arg2,type3,arg3,type4,arg4, \
        type5,arg5,type6,arg6) \
type name (type1 arg1,type2 arg2,type3 arg3,type4 arg4,type5 arg5,type6 arg6) \
{ \
long __res; \
__asm__ volatile ("push %%ebp ; movl %%eax,%%ebp ; movl %1,%%eax ; int $0x80 ; pop %%ebp" \
        : "=a" (__res) \
        : "i" (__NR_##name),"b" ((long)(arg1)),"c" ((long)(arg2)), \
          "d" ((long)(arg3)),"S" ((long)(arg4)),"D" ((long)(arg5)), \
          "0" ((long)(arg6))); \
__syscall_return(type,__res); \
}
```

We use 4 general purpose registers (eax,ebx,ecx,edx) plus the additional registers ESI e EDI, and the ebp register (**base pointer** for the current stack frame, which is saved before overwriting) and a local integer variable "i"
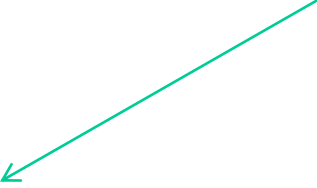
# UNISTD_32 calling conventions for system calls

```
/*
 *      0(%esp) - %ebx          ARGS
 *      4(%esp) - %ecx
 *      8(%esp) - %edx
 *      C(%esp) - %esi
 *     10(%esp) - %edi
 *     14(%esp) - %ebp          END ARGS
 *     18(%esp) - %eax
 *     1C(%esp) - %ds
 *     20(%esp) - %es
 *     24(%esp) - orig_eax
 *     28(%esp) - %eip
 *     2C(%esp) - %cs
 *     30(%esp) - %eflags
 *     34(%esp) - %oldesp
 *     38(%esp) - %oldss
 */
```

The stack layout representation complies with the traditional stack based passage of parameters

Ring and baseline CPU state information (firmware saved onto the system stack)

# UNISTD_64 calling conventions for system calls

```
/*
 * Register setup:
 * rax   system call number
 * rdi   arg0
 * rcx   return address for syscall/sysret, C arg3
 * rsi   arg1
 * rdx   arg2
 * r10   arg3 (--> moved to rcx for C)
 * r8    arg4
 * r9    arg5
 * r11   eflags for syscall/sysret, temporary for C
 * r12-r15,rbp,rbx saved by C code, not touched.
 *
 * Interrupts are off on entry.
 * Only called from user space.
 */
```

# System V AMD64 ABI additional details

- If the callee wishes to use registers RBX, RBP, and R12–R15, it must restore their original values before returning control to the caller

- All other registers must be saved by the caller if it wishes to preserve their values

# Details on passing parameters

- Once gained control, the dispatcher will take a complete snapshot of CPU registers

- The snapshot is taken within the **system level stack**

- **Then the dispatcher will invoke the system call as a subroutine call** (e.g. via a CALL instruction in x86 architectures)

- The actual system call will retrieve the parameters according to the proper ABI

- The taken snapshot can be modified upon the system call return (e.g. for delivering the return value)

# An example

## PHASE 1

registers

Sys call NR

System stack
upon triggering
dispatcher

## PHASE 2

Stack pointer

Stack pointer

Base pointer

Sys call NR

**Dispatcher execution**

## PHASE 3

Stack pointer

Base pointer

PC

Sys call NR

**system call execution**

# UNISTD_32 stack alignment

```
struct pt_regs {
        unsigned long bx;
        unsigned long cx;
        unsigned long dx;
        unsigned long si;
        unsigned long di;
        unsigned long bp;
        unsigned long ax;
        unsigned short ds;
        unsigned short __dsh;
        unsigned short es;
        unsigned short __esh;
        unsigned short fs;
        unsigned short __fsh;
        unsigned short gs;
        unsigned short __gsh;
        unsigned long orig_ax;
        unsigned long ip;
        unsigned short cs;
        unsigned short __csh;
        unsigned long flags;
        unsigned long sp;
        unsigned short ss;
        unsigned short __ssh;
}
```

Software saved
(no distinction between
caller/callee save)

Firmware saved

# UNISTD_64 stack alignment

```c
struct pt_regs {
        /* * C ABI says these regs are callee-preserved. They aren't saved on
kernel entry * unless syscall needs a complete, fully filled "struct pt_regs". */
        unsigned long r15;      unsigned long r14; unsigned long r13;
        unsigned long r12;      unsigned long bp;  unsigned long bx;
/* These regs are callee-clobbered. Always saved on kernel entry. */
        unsigned long r11;
        unsigned long r10;
        unsigned long r9;
        unsigned long r8;
        unsigned long ax;
        unsigned long cx;
        unsigned long dx;
        unsigned long si; unsigned long di;
/* * On syscall entry, this is syscall#. On CPU exception, this is error code. * On
hw interrupt, it's IRQ number: */
        unsigned long orig_ax;
/* Return frame for iretq */
        unsigned long ip;
        unsigned long cs;
        unsigned long flags;
        unsigned long sp;
        unsigned long ss;
        /* top of stack page */
};
```

Firmware
managed

# Simple examples for adding system calls to the <u>user API</u>

**Provide a C file which:**
- contains the definition of the numerical codes for the new system calls
- contains (or includes) the macro-definition for creating the actual standard  module associated with the new system calls (e.g. `_syscall0()` for UNISTD_32)

---

```
#include <unistd.h>
#define _NR_my_first_sys_call  254
#define _NR_my_second_sys_call  255

_syscall0(int,my_first_sys_call);
_syscall1(int,my_second_sys_call,int,arg);
```

# Simple overriding of the
# `fork()` UNISTD_32 system call

```c
#include <unistd.h>

#define __NR_my_fork 2 //same numerical code as the original
#define _new_syscall0(name) \
int name(void) \
{ \
 asm("int $0x80" : : "a" (__NR_##my_fork) ); \
 return 0; \
} \


_new_syscall0(my_fork)

int main(int a, char** b){
        my_fork();
        pause();  // there will be two processes pausing !!
}
```

# "int 0x80" system call path performance implications

- One memory access to the IDT

- One memory access to the GDT to retrieve the kernel CS segment

- One memory access to the GDT (namely the TSS) to retrieve the kernel level stack pointer

- A lot of clock cycles waiting for data coming from memory (just to control the execution flow)

- Asymmetric delays in asymmetric hardware (e.g. NUMA)

- Unreliable outcome for time-interval measures using system calls, see `gettimeofday()` (and *rdtsc*)

# The x86 revolution (starting with Pentium3)

- CS value for kernel code cached into an apposite MSR (Model Specific Register)

- Kernel entry point offset (the target EIP/RIP) kept into an apposite MSR

- Kernel level stack/data base kept into an apposite MSR

- Entering kernel code is as easy as flushing the MSRs values onto the corresponding original registers (e.g. CS, DS, SS …. recall that the corresponding bases are defaulted to 0x0)

- No memory access for activating the system call dispatcher

- **This is the fast system call path!!**

# Fast system call path additional details

**SYSENTER instruction for 32 bits - SYSCALL instruction for 64 bits**
    based on (pseudo) register manipulation

- CS register set to the value of (SYSENTER_CS_MSR)
- EIP register set to the value of (SYSENTER_EIP_MSR) for 32 bits
- SS register set to the sum of (8 plus the value in SYSENTER_CS_MSR for 32 bits – another bitmask for 64 bits)
- ESP register set to the value of (SYSENTER_ESP_MSR)

# SYSCALL — Fast System Call

| Opcode | Instruction | Op/En | 64-Bit Mode | Compat/Leg Mode | Description |
|--------|-------------|-------|-------------|-----------------|-------------|
| 0F 05 | SYSCALL | ZO | Valid | Invalid | Fast call to privilege level 0 system procedures. |

## Instruction Operand Encoding ¶

| Op/En | Operand 1 | Operand 2 | Operand 3 | Operand 4 |
|-------|-----------|-----------|-----------|-----------|
| ZO | NA | NA | NA | NA |

## Description ¶

SYSCALL invokes an OS system-call handler at privilege level 0. It does so by loading RIP from the IA32_LSTAR MSR (after saving the address of the instruction following SYSCALL into RCX). (The WRMSR instruction ensures that the IA32_LSTAR MSR always contain a canonical address.)

SYSCALL also saves RFLAGS into R11 and then masks RFLAGS using the IA32_FMASK MSR (MSR address C0000084H); specifically, the processor clears in RFLAGS every bit corresponding to a bit that is set in the IA32_FMASK MSR.

SYSCALL loads the CS and SS selectors with values derived from bits 47:32 of the IA32_STAR MSR. However, the CS and SS descriptor caches are **not** loaded from the descriptors (in GDT or LDT) referenced by those selectors. Instead, the descriptor caches are loaded with fixed values. See the Operation section for details. It is the responsibility of OS software to ensure that the descriptors (in GDT or LDT) referenced by those selector values correspond to the fixed values loaded into the descriptor caches; the SYSCALL instruction does not ensure this correspondence.

The SYSCALL instruction does not save the stack pointer (RSP). If the OS system-call handler will change the stack pointer, it is the responsibility of software to save the previous value of the stack pointer. This might be done prior to executing SYSCALL, with software restoring the stack pointer with the instruction following SYSCALL (which will be executed after SYSRET). Alternatively, the OS system-call handler may save the stack pointer and restore it before executing SYSRET.

**Instruction ordering.** Instructions following a SYSCALL may be fetched from memory before earlier instructions complete execution, but they will not execute (even speculatively) until all instructions prior to the SYSCALL have completed execution (the later instructions may execute before data stored by the earlier instructions have become globally visible).

# Fast system call path additional details

**SYSEXIT instruction for 32 bits - SYSRET instruction for 64 bits**

> based on pseudo register manipulation

- CS register set to the sum of (16 plus the value in SYSENTER_CS_MSR for 32 bits – another bitmask for 64 bits)
- EIP register set to the value contained in the EDX register for 32 bits – RCX for 64 bits
- SS register set to the sum of (24 plus the value in SYSENTER_CS_MSR for 32 bits – another bitmask for 64 bits)
- ESP register set to the value contained in the ECX register for 32 bits

# MSR and their setup

/usr/src/linux/include/asm/msr.h:
101 #define MSR_IA32_SYSENTER_CS   0x174
102 #define MSR_IA32_SYSENTER_ESP 0x175
103 #define MSR_IA32_SYSENTER_EIP  0x176


/usr/src/linux/arch/i386/kernel/sysenter.c:
36 wrmsr(MSR_IA32_SYSENTER_CS, __KERNEL_CS, 0);
37 wrmsr(MSR_IA32_SYSENTER_ESP, tss->esp1, 0);
38 wrmsr(MSR_IA32_SYSENTER_EIP,
                              (unsigned long) sysenter_entry, 0);


**rdmsr and wrmsr** are the actual machine instructions for
reading/writing the registers

# The `syscall()` construct
# (since Pentium3/kernel 2.6)

- `syscall()` is implemented within `glibc` (in `stdlib.h`)

- It allows triggering a trap to the kernel for the execution of a generic system call

- The first argument is the system call number

- The other parameters are the input for the system call code

- The actual ASM code implementation of `syscall()` is targeted and optimized for the specific architecture

- Specifically, the implementation (including the kernel level counterpart) relies on ASM instructions such as `sysenter/sysexit` or `syscall/sysret`, which have been made available starting from Pentium3 processors

# An example

```c
#include <stdlib.h>

#define __NR_my_first_sys_call  333
#define __NR_my_second_sys_call 334

int my_first_sys_call(){
     return syscall(__NR_my_first_sys_call);
}

int my_second_sys_call(int arg1){
     return syscall(__NR_my_second_sys_call, arg1);
}

int main(){
      int x;

     my_first_sys_call();
     my_second_sys_call(x);
}
```

# The system call table

- It is an array of function pointers

- However, we cannot easily resize the array and recompile the kernel

- This is because that table (like may other kernel level data structures) is positioned at compile time in specific zones of virtual addresses

- Simple enlarging on the table with no other modification of the kernel compilation layout will lead to data structures' overlap

- Such strict compilation rules depend on the fact that hardware setup for running the kernel may require CPU registers to be populated with compile time defined values

- The before described fast system call path is a clear example!!

# System call table hacking: entry reusage

- In older versions of the kernel the system call table was oversized, just for kernel modifiability purposes

- Hence the addition of new system calls in the kernel software could be based on the free entries

- In current (or more recent) kernel versions no oversize is put in place

- This is because the less "free" zones of data structures exist, the less the likelihood that they can be exploited against security

- But we are lucky because a few entries, although reserved, are not actually used to point to actual kernel level functions

- In the essence this is the scenario of kernel services that were planned (with given indexing) but not actually implemented

- All these entries point to the so called "`sys_ni_syscall`" kernel module, which simply returns upon its invocation

# x86 system call table details

- For kernel 2.4 and i386 machines the system call table is defined in `arch/i386/kernel/entry.S`

- For kernel 2.6.xx the table is posted on the file `arch/x86/kernel/syscall_table32.S`

- For kernel 4.15.xx and UNISTD_64 the table pointer is defined in `/arch/x86/entry/syscall_64.c`

- The `.S` files contain pre-processor ASM directives

- Any table entry keeps a symbolic reference to the kernel level name of a system call (typically, the kernel level name resembles the one used at application level)

- The above files (or other .S) also contains the code block for the dispatcher associated with the kernel access GATE

# Table structure – i386/UNISTD_32 style

```
ENTRY(sys_call_table)
     .long SYMBOL_NAME(sys_ni_syscall) /* 0  -  old "setup()"
system call*/
     .long SYMBOL_NAME(sys_exit)
     .long SYMBOL_NAME(sys_fork)
     .long SYMBOL_NAME(sys_read)
     .long SYMBOL_NAME(sys_write)
     .long SYMBOL_NAME(sys_open)                   /* 5 */
     .long SYMBOL_NAME(sys_close)

     ……
     .long SYMBOL_NAME(sys_sendfile64)
     .long SYMBOL_NAME(sys_ni_syscall) /* 240 reserved for futex
*/

     ………
     .long SYMBOL_NAME(sys_ni_syscall) /* 252
sys_set_tid_address */
```

**New symbols need to be inserted here**

```
     .rept NR_syscalls-(.-sys_call_table)/4
          .long SYMBOL_NAME(sys_ni_syscall)
     .endr
```

# Definition of system call symbols

- For the previous example, the actual system call specification will be

```
.long SYMBOL_NAME(sys_my_first_sys_call)
.long SYMBOL_NAME(sys_my_second_sys_call)
```

- The actual code for the system calls (generally based exclusively on C with compilation directives for the specific architecture) can be included within new modules added to the kernel or within already exiting modules

- The actual code can rely on the kernel global data structures and on functions already available within the kernel, except for the case where they are explicitly masked (e.g. masking with `static` declarations external to the file containing the system call)

# Definition of the system call table – UNISTD_64 style

- The kernel level source file that defines the system call table is
  [arch/x86/entry/syscall_64.c](arch/x86/entry/syscall_64.c)

```
asmlinkage const sys_call_ptr_t sys_call_table[__NR_syscall_max+1] = {
        [0 ... __NR_syscall_max] = &sys_ni_syscall,
        #include <asm/syscalls_64.h>
};
```

After the "include" expansion

```
asmlinkage const sys_call_ptr_t sys_call_table[__NR_syscall_max+1]  =  {
        [0 ... __NR_syscall_max] = &sys_ni_syscall,
        [0] = sys_read,
        [1] = sys_write,
        [2] = sys_open,
        ...
        ...
        ... };
```

# Compilation directives for kernel side systems calls

- Specific directives are used to make the system call code compliant with the dispatching rules

- **<u>Compliance is assessed on the basis of how the input parameters are passed/retrieved</u>**

- The input parameters passage by convention historically took place via the kernel stack

- The corresponding compilation directive is `asmlinkage`

- This directive is now mapped to the current ABI

- Hence for the previous examples we will have the following system call definitions

```
asmlinkage long sys_my_first_sys_call() { return 0;}
asmlinkage long sys_my_second_sys_call(int x) {
        return ((x>0)?x:-x);}
```

# The `ni_sys_call` module

```
asmlinkage long sys_ni_syscall(void) {

    return -ENOSYS;

}
```

# The actual dispatcher (trap driven activation – i386/UNISTD_32)

```
ENTRY(system_call)
        pushl %eax                              # save orig_eax
        SAVE_ALL
        GET_CURRENT(%ebx)
        testb $0x02,tsk_ptrace(%ebx)            # PT_TRACESYS
        jne tracesys
        cmpl $(NR_syscalls),%eax
        jae badsys
        call *SYMBOL_NAME(sys_call_table)(,%eax,4)
        movl %eax,EAX(%esp)                      # save the return value
ENTRY(ret_from_sys_call)
        cli                                     # need_resched and signals atomic test
        cmpl $0,need_resched(%ebx)
        jne reschedule
        cmpl $0,sigpending(%ebx)
        jne signal_return
restore_all:
        RESTORE_ALL
```

Manipulating the CPU snapshot in the stack

**Beware this!!!**

# The actual dispatcher (`syscall` driven activation – UNISTD_64/kernel 2.4)

```
ENTRY(system_call)
        swapgs
        movq         %rsp,PDAREF(pda_oldrsp)
        movq         PDAREF(pda_kernelstack),%rsp
        sti
        SAVE_ARGS 8,1
        movq  %rax,ORIG_RAX-ARGOFFSET(%rsp)
        movq  %rcx,RIP-ARGOFFSET(%rsp)
        GET_CURRENT(%rcx)
        testl $PT_TRACESYS,tsk_ptrace(%rcx)
        jne tracesys
        cmpq $__NR_syscall_max,%rax
        ja badsys
        movq %r10,%rcx
        call *sys_call_table(,%rax,8)  # XXX:              rip relative
        movq %rax,RAX-ARGOFFSET(%rsp)
        .globl ret_from_sys_call
ret_from_sys_call:
sysret_with_reschedule:
        GET_CURRENT(%rcx)
        cli
        cmpq $0,tsk_need_resched(%rcx)
        jne sysret_reschedule
        cmpl $0,tsk_sigpending(%rcx)
        jne sysret_signal
sysret_restore_args:
        ……….
```

#define PDAREF(field) %gs:field

Part of the stack switch work originally done via firmware is moved to software

**Beware this!!!**

# User vs kernel GS segment

---

# SWAPGS — Swap GS Base Register

| Opcode | Instruction | Op/En | 64-Bit Mode | Compat/Leg Mode | Description |
|--------|-------------|-------|-------------|-----------------|-------------|
| 0F 01 F8 | SWAPGS | ZO | Valid | Invalid | Exchanges the current GS base register value with the value contained in MSR address C0000102H. |

## Instruction Operand Encoding ¶

| Op/En | Operand 1 | Operand 2 | Operand 3 | Operand 4 |
|-------|-----------|-----------|-----------|-----------|
| ZO | NA | NA | NA | NA |

## Description ¶

SWAPGS exchanges the current GS base register value with the value contained in MSR address C0000102H (IA32_KERNEL_GS_BASE). The SWAPGS instruction is a privileged instruction intended for use by system software.

When using SYSCALL to implement system calls, there is no kernel stack at the OS entry point. Neither is there a straightforward method to obtain a pointer to kernel structures from which the kernel stack pointer could be read. Thus, the kernel cannot save general purpose registers or reference memory.

By design, SWAPGS does not require any general purpose registers or memory operands. No registers need to be saved before using the instruction. SWAPGS exchanges the CPL 0 data pointer from the IA32_KERNEL_GS_BASE MSR with the GS base register. The kernel can then use the GS prefix on normal memory references to access kernel data structures. Similarly, when the OS kernel is entered using an interrupt or exception (where the kernel stack is already set up), SWAPGS can be used to quickly get a pointer to the kernel data structures.

The IA32_KERNEL_GS_BASE MSR itself is only accessible using RDMSR/WRMSR instructions. Those instructions are only accessible at privilege level 0. The WRMSR instruction ensures that the IA32_KERNEL_GS_BASE MSR contains a canonical address.

# … moving to kernel 4.xx or later

Snippet taken from

https://github.com/torvalds/linux/blob/master/arch/x86/entry/entry_64.S

```asm
ENTRY(entry_SYSCALL_64)
        UNWIND_HINT_EMPTY
        /*
         * Interrupts are off on entry.
         * We do not frame this tiny irq-off block with TRACE_IRQS_OFF/ON,
         * it is too small to ever cause noticeable irq latency.
         */

        swapgs
        /*
         * This path is only taken when PAGE_TABLE_ISOLATION is disabled so it
         * is not required to switch CR3.
         */
        movq    %rsp, PER_CPU_VAR(rsp_scratch)
        movq    PER_CPU_VAR(cpu_current_top_of_stack), %rsp

        /* Construct struct pt_regs on stack */
        pushq   $__USER_DS                      /* pt_regs->ss */
        pushq   PER_CPU_VAR(rsp_scratch)        /* pt_regs->sp */
        pushq   %r11                            /* pt_regs->flags */
        pushq   $__USER_CS                      /* pt_regs->cs */
        pushq   %rcx                            /* pt_regs->ip */
GLOBAL(entry_SYSCALL_64_after_hwframe)
        pushq   %rax                            /* pt_regs->orig_ax */

        PUSH_AND_CLEAR_REGS rax=$-ENOSYS

        TRACE_IRQS_OFF

        /* IRQs are off. */
        movq    %rax, %rdi
        movq    %rsp, %rsi
        call    do_syscall_64                   /* returns with IRQs disabled */

        TRACE_IRQS_IRETQ                        /* we're about to change IF */

        /*
         * Try to use SYSRET instead of IRET if we're returning to
         * a completely clean 64-bit userspace context.  If we're not,
         * go to the slow exit path.
         */
```

Here we pass control to a C-stub, not to the actual system call

# Snippet taken from

https://github.com/torvalds/linux/blob/master/arch/x86/entry/common.c

```
271    #ifdef CONFIG_X86_64
272    __visible void do_syscall_64(unsigned long nr, struct pt_regs *regs)
273    {
274            struct thread_info *ti;
275
276            enter_from_user_mode();
277            local_irq_enable();
278            ti = current_thread_info();
279            if (READ_ONCE(ti->flags) & _TIF_WORK_SYSCALL_ENTRY)
280                    nr = syscall_trace_enter(regs);
281
282            /*
283             * NB: Native and x32 syscalls are dispatched from the same
284             * table.  The only functional difference is the x32 bit in
285             * regs->orig_ax, which changes the behavior of some syscalls.
286             */
287            nr &= __SYSCALL_MASK;
288            if (likely(nr < NR_syscalls)) {
289                    nr = array_index_nospec(nr, NR_syscalls);
290                    regs->ax = sys_call_table[nr](regs);
291            }
292
293            syscall_return_slowpath(regs);
294    }
295    #endif
```

Wrong-speculation cannot rely on arbitrary sys-call indexes!!!!

Also, from kernel 4.17 the system call table entry no longer points to the actual system call code, rather to another wrapper that masks from the stack non-useful values

# Overall

- For more security-oriented implementations we have

  - ✓ More strict checks and manipulation of the user provided information before any action is taken

  - ✓ A more layered architecture for better decoupling user/kernel information flows

- The latter point has reflection on programming aspects since for, e.g., Kernel 4.17 the kernel-side creation of a new system call should be based on kernel level macros for implementing a stub-based execution of the native system-call code

- These macros are `SYSCALL_DEFINE0, SYSCALL_DEFINE1, SYSCALL_DEFINE2, SYSCALL_DEFINE3` ……

# Actual usage/effect of kernel-side sys-call macros

- The `SYSCALL_DEFINE2` example (still representative of other macros)

```
SYSCALL_DEFINE2(name, param1type, param1name, param2type, param2name){

        actual body implementing the kernel side system call

}
```

The macro creates a function
`sys_name` (aliased by `SyS_name`) or
`__x86_sys_name` from kernel 4.17

In 4.17 this function passes only the requested values (i.e.
`param1name` and `param2name`) to the actual function
related to the above specified body - such a function has
now name `__se_sys_name`

# Finally … PTI (Page Table Isolation)

```
144
145    ENTRY(entry_SYSCALL_64)
146          UNWIND_HINT_EMPTY
147          /*
148           * Interrupts are off on entry.
149           * We do not frame this tiny irq-off block with TRACE_IRQS_OFF/ON,
150           * it is too small to ever cause noticeable irq latency.
151           */
152
153          swapgs
154          /* tss.sp2 is scratch space. */
155          movq    %rsp, PER_CPU_VAR(cpu_tss_rw + TSS_sp2)
156          SWITCH_TO_KERNEL_CR3 scratch_reg=%rsp
157          movq    PER_CPU_VAR(cpu_current_top_of_stack), %rsp
158
159          /* Construct struct pt_regs on stack */
160          pushq   $__USER_DS                      /* pt_regs->ss */
161          pushq   PER_CPU_VAR(cpu_tss_rw + TSS_sp2)    /* pt_regs->sp */
162          pushq   %r11                            /* pt_regs->flags */
163          pushq   $__USER_CS                      /* pt_regs->cs */
164          pushq   %rcx                            /* pt_regs->ip */
165    GLOBAL(entry_SYSCALL_64_after_hwframe)
166          pushq   %rax                            /* pt_regs->orig_ax */
167
168          PUSH_AND_CLEAR_REGS rax=$-ENOSYS
169
170          TRACE_IRQS_OFF
171
172          /* IRQs are off. */
173          movq    %rax, %rdi
174          movq    %rsp, %rsi
175          call    do_syscall_64           /* returns with IRQs disabled */
176
177          TRACE_IRQS_IRETQ                /* we're about to change IF */
178
179          /*
180           * Try to use SYSRET instead of IRET if we're returning to
181           * a completely clean 64-bit userspace context.  If we're not,
```
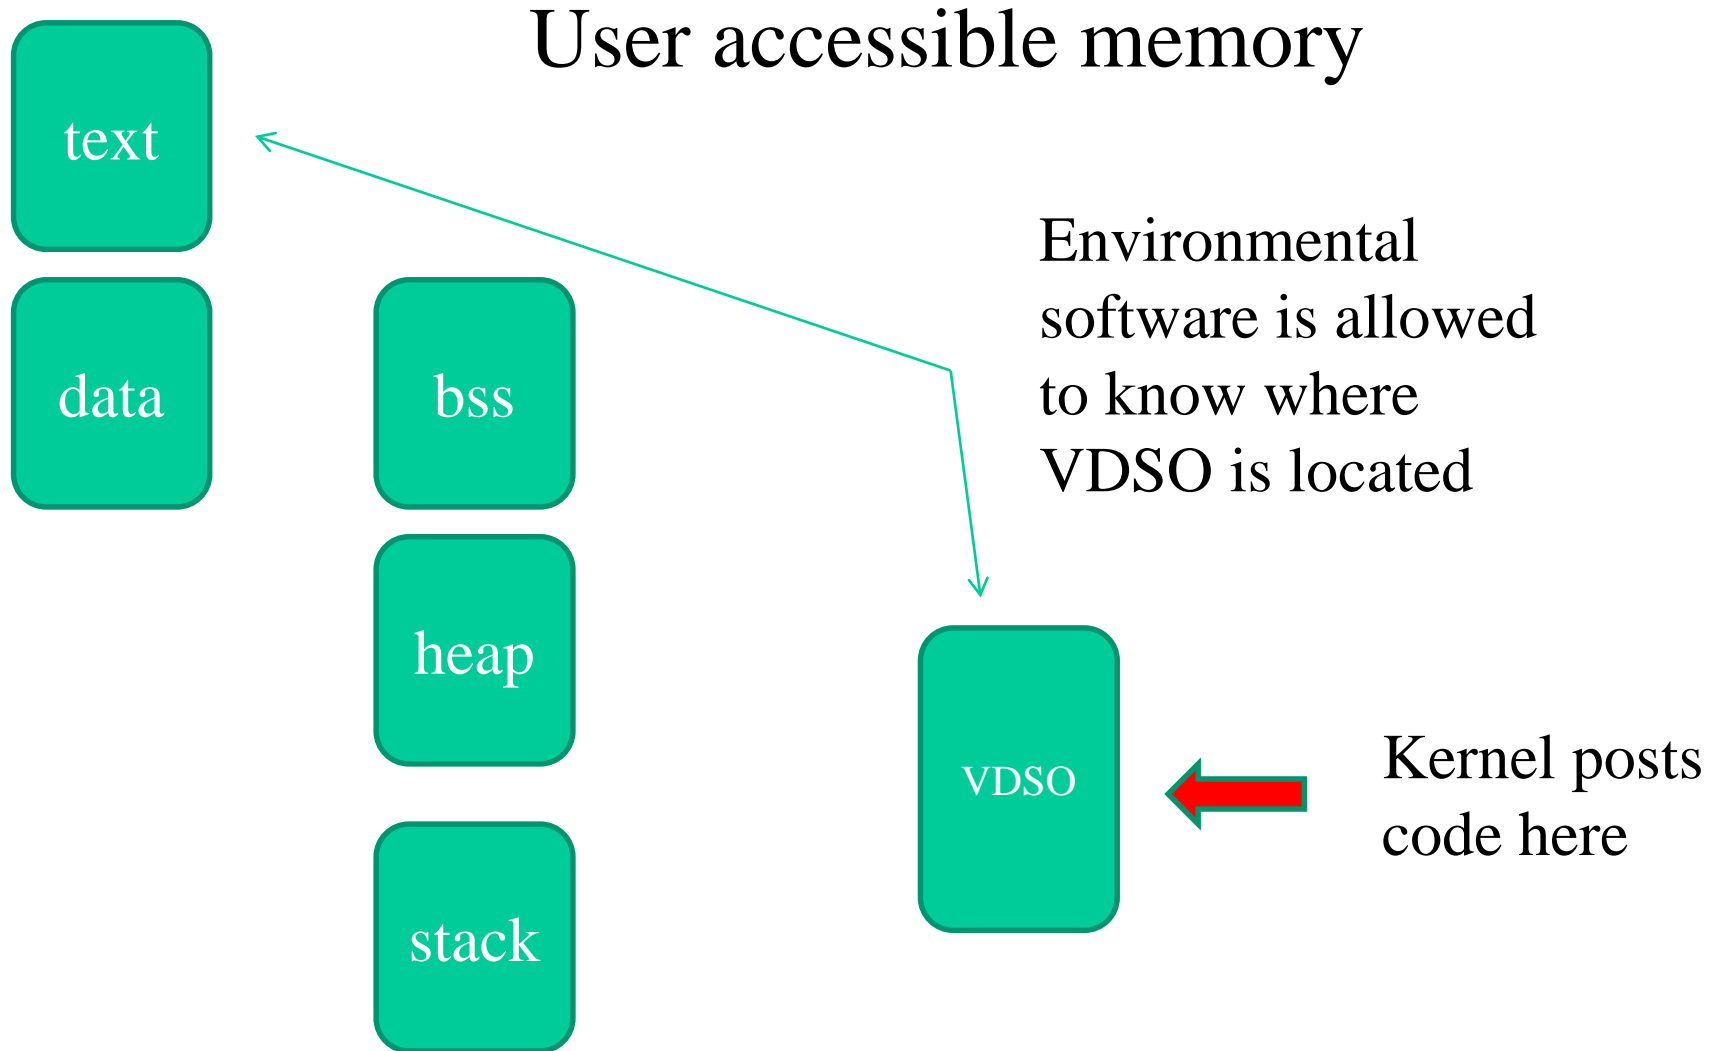
Switch to the kernel view of memory

# Virtual Dynamic Shared Object (VDSO)

- Kernel also setups system call entry/exit points for user processes

- Kernel creates a single page (or a few) in memory and attaches it to all processes' address space when they are loaded into memory.

- This page contains the actual implementation of the system call entry/exit mechanism

- Kernel calls this page **virtual dynamic shared object** (VDSO)

- Originally exploited for making the fast system call path available (in relation to a few services)

# VDSO and the address space

User accessible memory

text

data

bss

heap

stack

Environmental
software is allowed
to know where
VDSO is located

VDSO

Kernel posts
code here

# Application exposed facilities

**SYNOPSIS**
**#include <sys/auxv.h>**

**void *vdso = (uintptr_t) getauxval(AT_SYSINFO_EHDR);**

**DESCRIPTION**

The "vDSO" (virtual dynamic shared object) is a small shared library that the kernel automatically maps into the address space of all user-space applications. Applications usually do not need to concern themselves with these details as the vDSO is most commonly called by the C library. This way you can code in the normal way using standard functions and the C library will take care of using any functionality that is available via the vDSO.

# The actual VDSO (getcpu example)

# Performance effects

- The VDSO exploits flat (linear) addressing proper of operating system memory managers in order to bypass segmentation and the related operations

- It therefore reduces the number of accessed to memory in order to support the change to kernel mode

- Studies show that the reduction of clock cycles for system calls can be of the order of 75%

- This is in the end typical for any usage of the fast system call path

# The current picture

- VDSO is now used to replace the old facilities supported via the **vsyscall** section, say support for specific system calls (e.g. query system calls such as `gettimeofday()`)

- VDSO is randomized (in terms of positioning into the address space) so security gets increased

- The system call mechanism in the wide, which relies on `sysenter/syscall` and `sysexit/sysret`, is in charge of the dynamic linker (ld-linux.so)

# Back to the coexistence of slow and fast system call paths

- **Slow path**

  ✓ Still based on `int 0x80`

  ✓ Still accessing IDT/GDT (which is the reason why the target entry still requires to be populated)

  ✓ The kernel level system call dispatcher accesses the i386/INISTD_32 system call table

- **Fast path**

  ✓ Base on the `syscall` instruction (no IDT/GDT access)

  ✓ The kernel level dispatcher (different from the previous one) accesses the x86-64/UNISTD_64 system call table

# Kernel software organization

- About the 80-90% of the actual code for system calls is embedded within a few main portions of the kernel archive

- These are contained in the following directories

  - ➤ `kernel` (process and used management)

  - ➤ `mm` (basic memory management)

  - ➤ `ipc` (interprocess communication management)

  - ➤ `fs` (virtual file system management)

  - ➤ `net` (network management)

# Kernel compiling

- You can exploit **make**

- It executes a set of tasks (compilation, assembly and linking tasks) which are specified via a `Makefile`

- This file can specify differentiated actions to be done (possibly exhibiting dependencies) which are described within a field called **target**

- Each action can be specified by the following syntax:

```
action-name: [ dependency-name]*{new-line}
{tab} action-body
```

- Further, we can define variables via the syntax:

```
variable-name = value
```

- Any variable can be accessed via the syntax:

```
$(variable-name)
```

# Standard compilation steps (current tyle)

**make config (or menuconfig)**

**make**

**make modules**

**make modules_install (ROOT)**

**make install (ROOT)**

**mkinitrd (or mkinitramfs) –o initrd.img-<vers> <vers>**

**update-grub**

**OR**

**grub-mkconfig -o /boot/grub/grub.cfg (ROOT)**

# About 'config'

- The possibilities
  - **allyesconfig** (likelihood of conflicting modules)
  - **allnoconfig** (likelohood of non-sufficient services in the kernel image)
  - Answer to the individual questions you may be asked for
  - Retrieve a good configuration file (depending on you machine/settings) on the web
  - Reuse the configuration files(s) you find in the **/boot** directory of your root file system (likely works when recompiling the same kernel version you already have)

# Role of initrd

- It is a RAM disk
- It can be (temporary) mounted as the root file system and programs can be run from it
- A different root file system can be then mounted from a different device
- The previous root (from initrd) can then be moved to a directory and can be subsequently unmounted
- With initrd system startup can occur in two phases
  - the kernel initially comes up with a minimum set of compiled-in drivers
  - additional modules are loaded from initrd

# Step effects

**make config (or menuconfig)**

**make**

**make modules**

**make modules_install (ROOT) (writes into /lib/modules)**

**make install (ROOT) (writes into /boot: the kernel image, the system map and the config file)**

**update-grub**
**OR**
**grub-mkconfig -o /boot/grub/grub.cfg (ROOT)**

# "Extended" Kernel compilation (current style)

- Makefile updates
  1. setting of the `EXTRAVERSION` variable (non-mandatory)
  2. use obj- directive to add a file or a directory into the compilation tree
  3. the addition is within already available makefiles (or new ones)

# Kernel anatomy: the system map

- It contains the symbols and the corresponding virtual memory reference (as determined at compile/link time – beware randomization) for:
  - Kernel  functions (steady state ones)
  - Kernel data structures
- Each symbol is also associated with a tag that defines the 'storage class' as determined by the compiling process
- As an example, 'T' usually denotes a global (non-static but not necessarily exported) function, 't' a function local to the compilation unit (i.e. static), 'D'  global data, 'd' data local to the compilation unit. 'R' and 'r' same as 'D'/'d' but for read-only data

# System map applications

- Kernel debugging
- Kernel run-time hacking
- The system map is also (partially) reported by the (pseudo) file `/proc/kallsysm`
- The latter is exploited for run-time kernel 'hacking' via the modules' technology

# Just an example

2.6.5-7.282-smp #1 SMP ……. i686 i686 i386 GNU/Linux

c03a8a00 D sys_call_table

**Read/write data**

2.6.32-5-amd64 #1 SMP ……… x86_64 GNU/Linux

ffffffff81308240 R sys_call_table

**Read-only data**