MS degree in Computer Engineering University of Rome Tor Vergata Lecturer: Francesco Quaglia

# **Topics:**

- 1. The very base on boot vs memory management
- 2. Memory 'Nodes' (UMA vs NUMA)
- 3. Paging support in x86
- 4. Boot and steady state behavior of the memory system in LINUX
- 5. Kernel level memory allocation/deallocation services

## **Basic terminology**

- firmware: a program coded on a ROM device, which can be executed when powering a processor on
- **bootsector**: predefined device (e.g. disk) sector keeping executable code for system startup
- bootloader: the actual executable code loaded and launched right before giving control to the target operating system
  - ➤ this code is partially kept within the bootsector, and partially kept into other sectors
  - ➤It can be used to parameterize the actual operating system boot

### **Startup tasks**

- The firmware gets executed, which loads in memory and launches the bootsector content
- The loaded bootsector code gets launched, which may load other bootloader portions
- The bootloader ultimately loads the actual operating system kernel and gives it control
- The kernel performs its own startup actions, which may entail architecture setup, data structures and software setup, and process activations
- To emulate a steady state unique scenario, at least one process is derived from the boot thread (namely the IDLE PROCESS)

#### Firmware on x86

- It is called BIOS (Basic I/O System)
- Interactive mode can be activated via proper interrupts (e.g. the F1 key)
- Interactive mode can be used to parameterize firmware execution (the parameterization is typically kept via CMOS rewritable memory devices powered by apposite temporary power suppliers)
- The BIOS parameterization can determine the order for searching the boot sector on different devices
- A device boot sector will be searched for only if the device is registered in the BIOS list

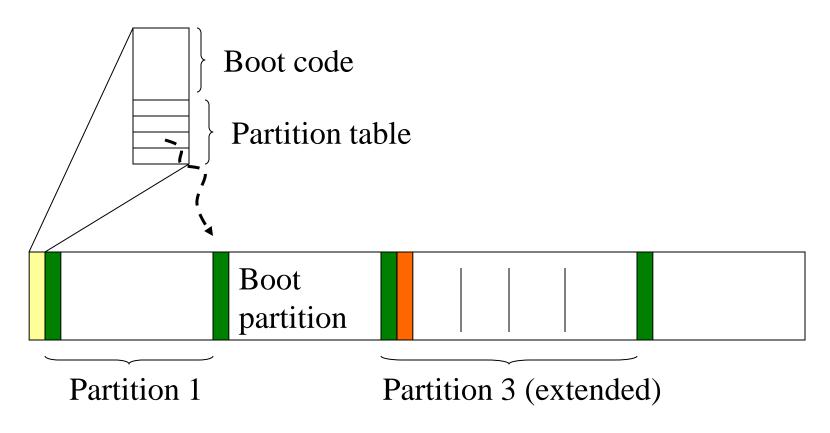
#### Bootsector on x86

- The first device sector keeps the so called <u>master boot</u> record (MBR)
- This sector keeps executable code and a 4/8-entry tables, each one identifying a different device partition (in terms of its positioning on the device)
- The first sector in each partition can operate as the partition boot sector (BS)
- In case the partition is extended, then it can additionally keep up to 4 sub-partitions (hence the partition boot sector can be structured to keep an additional partitioning table)
- Each sub-partition can keep its own boot sector

# RAM image of the x86 MBR

Offset	Size (bytes)	Description						
0	436 (to 446, if you need a little extra)	MBR Bootstrap (flat binary executable code)						
0x1b4	10	Optional "unique" disk ID <sup>1</sup>						
0x1be	64	MBR Partition Table, with 4 entries (below)						
0x1be	16	First partition table entry						
0x1ce	16	Second partition table entry						
0x1de	16	Third partition table entry						
0x1ee	16	Fourth partition table entry						
0x1fe	2	(0x55, 0xAA) "Valid bootsector" signature bytes						

# An example scheme



- Boot sector
- Extended partition boot record

# Historical LINUX bootsector organization for i386

- The historical bootsector code for LINUX (i386) is within the kernel file arch/i386/bootsect.S (it is no more used in case boot operates via LILO or GRUB bootloaders)
- This code loads arch/i386/bootsetup. S and the kernel image in memory
- The code within arch/i386/bootsetup. S gets launched for initializing the architecture (e.g. the processor initial state for the actual kernel boot)
- This code ultimately gives control to the initial kernel image

# Current LINUX bootsector organization on x86

- Similar steps are executed via, e.g., GRUB on generic x86 machines
- The machine setup code ultimately passes control to the initial kernel image
- This kernel image executes starting from the start\_kernel() in the init/main.c
- This kernel image is way different, both in size and structure, from the one that will operate at steady state
- Just to name one reason, boot is highly configurable!

### What about boot on multi-core machines

- The start\_kernel() function is executed along a single CPU-core (the master)
- All the other cores (the slaves) only keep waiting that the master has finished
- The kernel internal function smp\_processor\_id()
  can be used for retrieving the ID of the current core
- This function is based on ASM instructions implementing a hardware specific ID detection protocol
- This function operates correctly either at kernel boot or at steady state

# The actual support for CPU-core identification

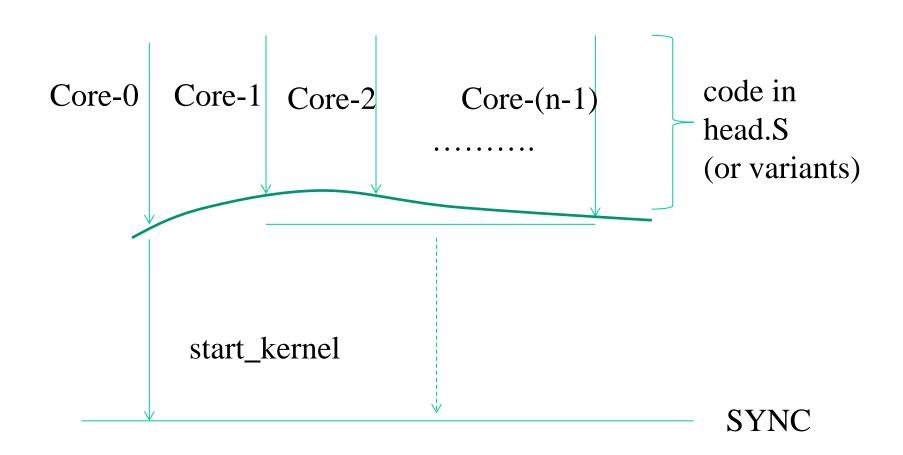
# x86 Instruction Set Reference

# **CPUID**

#### **CPU Identification**

	Opcode	Mnemonic	Description
6	F A2	CPUID	Returns processor identification and feature information to the EAX, EBX, ECX, and EDX registers, according to the input value
			entered initially in the EAX register.

# Actual kernel startup scheme



# An example head. S code snippet: triggering paging (IA32 case)

```
/* * Enable paging */ 3:
movl $swapper_pg_dir-__PAGE_OFFSET, %eax
movl %eax, %cr3 /* set the page table
pointer.. */
movl %cr0, %eax
orl $0x80000000, %eax
movl %eax, %cr0 /* ..and set paging (PG) bit
*/
```

# Hints on the signature of the start kernel function (as well as others)

..... \_\_init start\_kernel(void)

This only lives in memory during kernel boot (or startup)

The reason for this is to recover main memory storage which is relevant because of both:

- Reduced available RAM (older configurations)
- Increasingly complex (and hence large in size) startup code

Recall that the kernel image is not subject to swap out (kernel is resident )

# Management of init functions

- The kernel linking stage locates these functions on specific logical pages (recall what we told about the fixed positioning of specific kernel level stuff in the kernel layout!!)
- These logical pages are identified within a "bootmem" subsystem that is used for managing memory when the kernel is not yet at steady state of its memory management operations
- Essentially the bootmem subsystem keeps a bitmask with one bit indicating whether a given page has been used (at compile time) for specific stuff

# How is RAM memory organized on modern (large scale/parallel) machines?

- In modern chipsets, the CPU-core count continuously increases
- However, it is increasingly difficult to build architectures with a flat-latency memory access (historically referred to as UMA)
- Current machines are typically NUMA
- Each CPU-core has some RAM banks that are close and other that are far
- Generally speaking, each memory bank is associated with a so called NUMA-node
- Modern operating systems are designed to handle NUMA machines (hence UMA as a special case)

# Looking at the NUMA setup via Operating System facilities

- A very simple way is the numactl command
- It allows to discover
  - ✓ How many NUMA nodes are present
  - ✓ What are the nodes close/far to/from any CPU-core
  - ✓ What is the actual distance of the nodes (from the CPU-cores)

Let's see a few 'live' examples ......

# Actual kernel data structures for managing memory

- Kernel Page table
  - This is a kind of 'ancestral' page table (all the others are somehow derived from this one)
  - ➤It keeps the memory mapping for kernel level code and data (thread stack included)
- Core map
  - The map that keeps status information for any frame (page) of physical memory, and for any NUMA node
- Free list of physical memory frames, for any NUMA node

None of them is already finalized when we startup the kernel

# Setup of the kernel page table

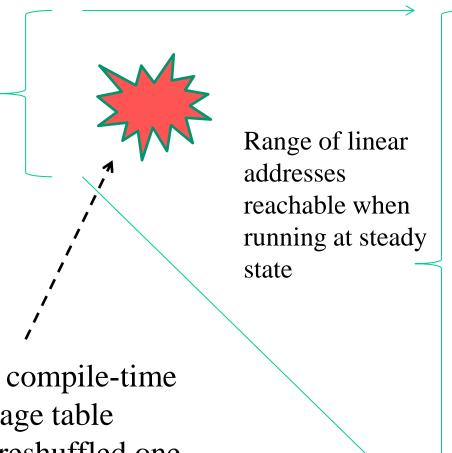
- This setup takes place via setup\_arch() which is present in arch/XX/kernel/setup.c
- The main function invoked here is a configuration specific version of paging init()
- As an example, for i386 processors this function is specified in arch/i386/mm/init.c
- Clearly, new generations of processors provide different (more powerful) paging support and clearly a different page table structure
- We will have a tour starting from i386 and then finally reaching x86-64 processors (via incremental differences)

# Objectives of the kernel page table setup

- These are basically two:
  - ✓ Allowing the kernel software to use virtual addressed while executing (either at startup or at steady state)
  - ✓ Allowing the kernel software (and consequently the application software) to reach (in read and/or write mode) the maximum admissible (for the specific machine) or available RAM storage
- The finalized shape of the kernel page table is therefore typically not setup into the original image of the kernel loaded in memory, e.g., given that the available RAM to drive can be parameterized

#### A scheme

Range of linear addresses reachable when switching protected mode plus paging



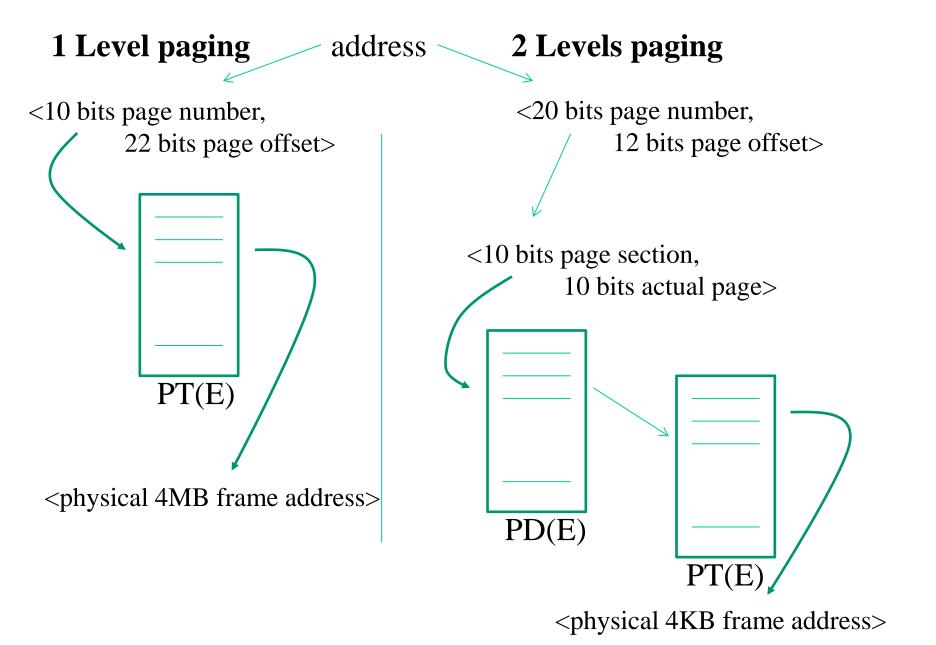
Increase
of the size
Of reachable
RAM locations

Passage from a compile-time defined kernel-age table
To a boot time reshuffled one

## Virtual memory vs boot sequence (2.4 example)

- Upon kernel startup addressing relies on a simple single level paging mechanism that only maps 2 pages (each of 4 MB) up to 8 MB physical addresses
- The actual paging rule (namely the page granularity and the number of paging levels up to 2 in i386) is identified via proper bits within the entries of the page table
- The physical address of the setup page table is kept within the CR3 register
- The steady state paging scheme used by LINUX will be activated during the kernel boot procedure
- The max size of the address space for LINUX processes on i386 machines is 4 GB
  - > 3 GB are within user level segments
  - ➤ 1 GB is within kernel level segments

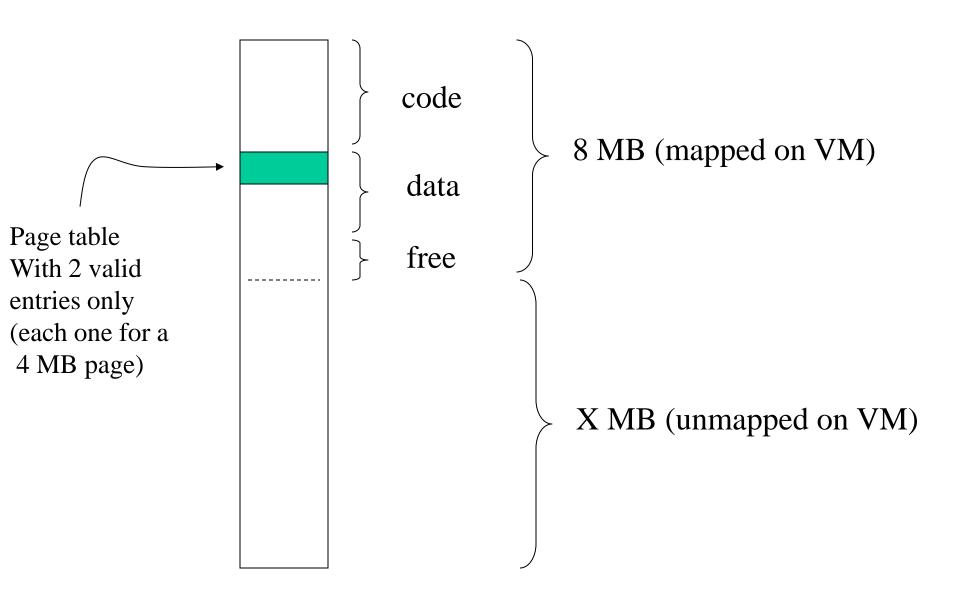
# Details on the page table structure in i386 (i)



## Details on the page table structure in i386 (ii)

- It is allocated in physical memory into 4KB blocks, which can be non-contiguous
- In typical LINUX configurations, once set-up it maps 4 GB addresses, of which 3 GB at null reference and (almost) 1 GB onto actual physical memory
- Such a mapped 1 GB corresponds to kernel level virtual addressing and allows the kernel to span over 1 GB of physical addresses
- To drive more physical memory, additional configuration mechanisms need to be activated, or more recent processors needs to be exploited as we shall see

# i386 memory layout at kernel startup for 2.4 kernels



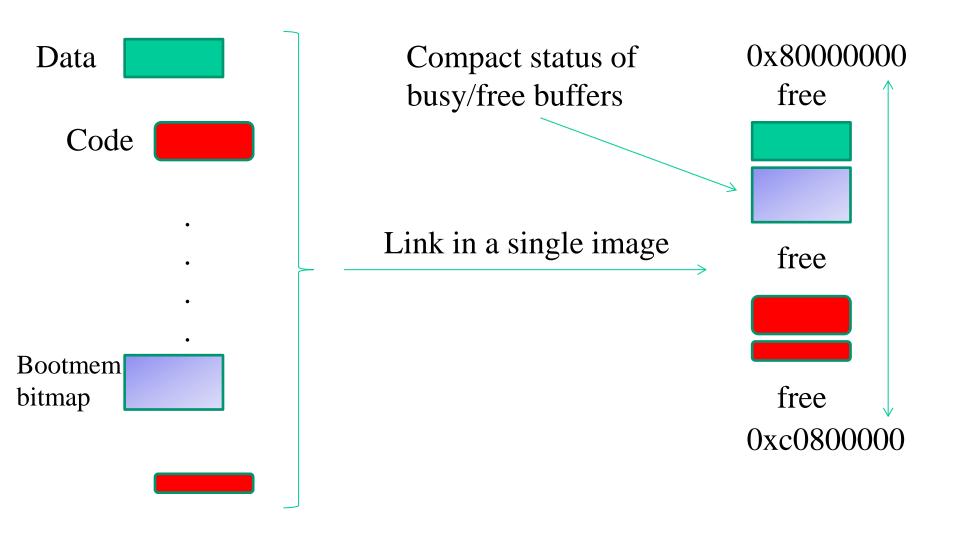
### Actual issues to be tackled

- 1. We need to reach the correct granularity for paging (4KB rather than 4MB)
- 2. We need to span logical to physical address across the whole 1GB of manageable physical memory
- 3. We need to re-organize the page table in two separate levels
- 4. So we need to determine 'free buffers' within the already reachable memory segment to initially expand the page table
- 5. We cannot use memory management facilities other than paging (since core maps and free lists are not at steady state)

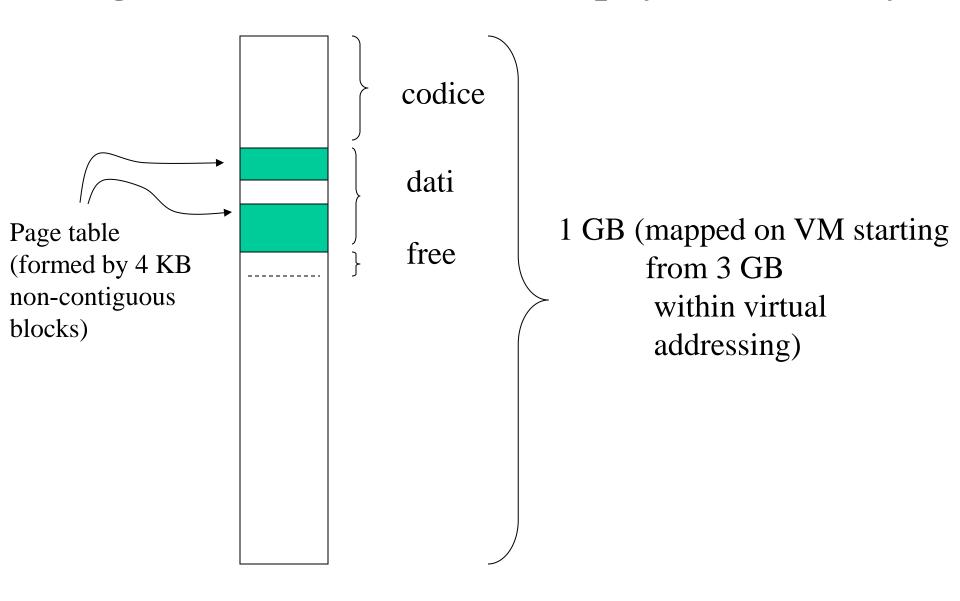
# Back to the concept of bootmem

- 1. Memory occupancy and location of the initial kernel image is determined by the compile/link process
- 2. A compile/link time memory manager is embedded into the kernel image, which is called bootmem manager
- 3. It relies on bitmaps telling if any 4KB page in the currently reachable memory image is busy or free
- 4. It also offers API (to be employed at boot time) in order to get free buffers
- 5. These buffers are sets of contiguous (or single) page alligned areas
- 6. This subsystem is in charge of handling \_init marked functions in terms of final release of the corresponding buffers

# An exemplified picture of bootmem (i386/kernel 2.4)

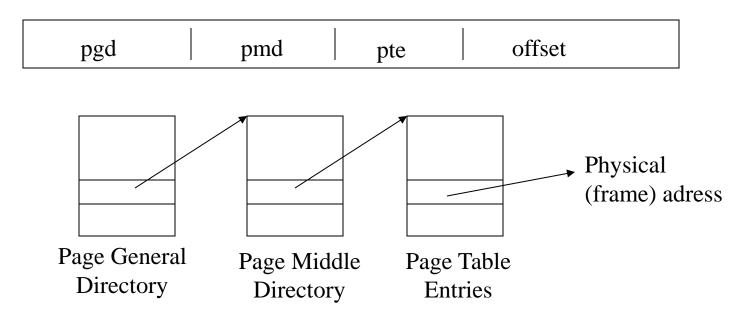


# Page table collocation within physical memory



## LINUX paging vs i386

• LINUX virtual addresses exhibit (at least) 3 indirection levels



- On i386 machines, paging is supported limitedly to 2 levels (pde, page directory entry pte, page table entry)
- Such a dicotomy is solved by setting null the pmd field, which is proper of LINUX, and mapping
  - ▶pgd LINUX on pde i386
  - >pte LINUX on pte i386

### i386 page table size

- Both levels entail 4 KB memory blocks
- Each block is an array of 4-byte entries
- Hence we can map 1 K x 1K pages
- Since each page is 4 KB in sixe, we get a 4 GB virtual addressing space
- The following macros define the size of the page tables blocks (they can be found in the file include/asm-i386/pgtable-2level.h)

```
▶#define PTRS_PER_PGD 1024
▶#define PTRS_PER_PMD 1
▶#define PTRS_PER_PTE 1024
```

• the value1 for PTRS\_PER\_PMD is used to simulate the existence of the intermediate level such in a way to keep the 3-level oriented software structure to be compliant with the 2-level architectural support

## Page table data structures

- A core structure is represented by the symbol swapper\_pg\_dir which is defined within the file arch/i386/kernel/head.S
- This symbol expresses the virtual memory address of the **PGD (PDE)** portion of the kernel page table
- This value is initialized at compile time, depending on the memory layout defined for the kernel bootable image
- Any entry within the PGD is accessed via displacement starting from the initial PGD address
- The C types for the definition of the content of the page table entries on i386 are defined in include/asm-i386/page.h
- They are

```
typedef struct { unsigned long pte_low; } pte_t;
typedef struct { unsigned long pmd; } pmd_t;
typedef struct { unsigned long pgd; } pgd_t;
```

## **Debugging**

- The redefinition of different structured types, which are identical in size and equal to an unsigned long, is done for debugging purposes
- Specifically, in C technology, <u>different aliases for the same type</u> are considered as identical types
- For instance, if we define typedef unsigned long pgd\_t;

typedef unsigned long pte\_t;

pgd\_t x; pte\_t y;

the compiler enables assignments such as x=y and y=x

• Hence, there is the need for defining different structured types which simulate the base types that would otherwise give rise to compiler equivalent aliases

### i386 PDE entries

#### Page-Directory Entry (4-KByte Page Table)

31		12	11	9	8	7	6	5	4	3	2	1	0
	Page-Table Base Address		Ava	iil	G	PS	0	Α	PCD	P W T	U / S	R / W	Р
	Available for system programmer's use Global page (Ignored) ————————————————————————————————————												

## i386 PTE entries

#### Page-Table Entry (4-KByte Page)

31		12	11	9	8	7	6	5	4	3	2	1	0
	Page Base Address		Avai	l	G	P A T	D	А	P C D	P W T	U / S	R / W	Р
	Available for system programmer's use Global Page Page Table Attribute Index Dirty Accessed Cache Disabled Write-Through User/Supervisor Read/Write Present												

#### **Field semantics**

- **Present**: indicates whether the page or the pointed page table is loaded in physical memory. This flag is not set by firmware (rather by the kernel)
- **Read/Write**: define the access privilege for a given page or a set of pages (as for PDE). Zero means read only access
- **User/Supervisor**: defines the privilege level for the page or for the group of pages (as for PDE). Zero means supervisor privilege
- Write Through: indicates the caching policy for the page or the set of pages (as for PDE). Zero means write-back, non-zero means write-through
- Cache Disabled: indicates whether caching is enabled or disabled for a page or a group of pages. Non-zero value means disabled caching (as for the case of memory mapped I/O)

- **Accessed**: indicates whether the page or the set of pages has been accessed. This is a sticky flag (no reset by firmware). Reset is controlled via software
- **Dirty**: indicates whether the page has been write-accessed. This is also a sticky flag
- Page Size (PDE only): if set indicates 4 MB paging otherwise 4 KB paging
- Page Table Attribute Index: ..... Do not care ......
- Page Global (PTE only): defines the caching policy for TLB entries. Non-zero means that the corresponding TLP entry does not require reset upon loading a new value into the page table pointer CR3

# Bit masking

- in include/asm-i386/pgtable.h there exist some macros defining the positioning of control bits within the entries of the PDE or PTE
- There also exist the following macros for masking and setting those bits

• These are all machine dependent macros

#### An example

```
pte t x;
x = ...;
if ((x.pte low) & PAGE PRESENT) {
    /* the page is loaded in a frame */
else{
     /* the page is not loaded in any
         frame */
```

## Relations the trap/interrupt events

- Upon a TLB miss, firmware accesses the page table
- The first checked bit is typically \_PAGE\_PRESENT
- If this bit is zero, a page fault occurs which gives rise to a trap (with a given displacement within the trap/interrupt table)
- Hence the instruction that gave rise to the trap can get finally re-executed
- Re-execution might give rise to additional traps, depending on firmware checks on the page table
- As an example, the attempt to access a read only page in write mode will give rise to a trap (which triggers the segmentation fault handler)

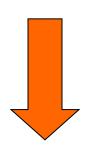
#### Run time detection of current page size on IA-32 processors

# Low level "pages"

Load undersized page table (kernel page size not finalized: 4MB)

- 4 KB (1K entry)

Finalize kernel handled page size (4KB)



Expand page table via boot mem low pages (not marked in the page table)

- compile time identification



Kernel boot

## Kernel page table initialization

- As said, the kernel PDE is accessible at the virtual address kept by swapper\_pg\_dir (now init\_level4\_pgt on x86-64/kernel 3 or init top pgt on x86-64/kernel4)
- The room for PTE tables gets reserved within the 8MB of RAM that are accessible via the initial paging scheme
- Reserving takes place via the macro alloc\_bootmem\_low\_pages() which is defined in include/linux/bootmem.h (this macro returns a virtual address)
- Particularly, it returns the pointer to a 4KB (or 4KB x N) buffer which is page aligned
- This function belongs to the (already hinted) basic memory management subsystem upon which the LINUX memory system boot lies on

# **Initialization algorithm**

- we start by the PGD entry which maps the address 3 GB, namely the entry numbered 768
- cyclically
  - 1. We determine the virtual address to be memory mapped (this is kept within the vaddr variable)
  - 2. One page for the PTE table gets allocated which is used for mapping 4 MB of virtual addresses
  - 3. The table entries are populated
  - 4. The virtual address to be mapped gets updated by adding 4 MB
  - 5. We jump to step 1 unless no more virtual addresses or no more physical memory needs to be dealt with (the ending condition is recorded by the variable end)

# Initialization function pagetable init()

```
for (; i < PTRS PER PGD; pgd++, i++) {</pre>
       vaddr = i*PGDIR SIZE; /* i is set to map from 3 GB */
       if (end && (vaddr >= end)) break;
       pmd = (pmd t *)pgd;/* pgd initialized to (swapper pg dir+i) */
       for (j = 0; j < PTRS PER PMD; pmd++, j++) {
       pte base = pte = (pte t *) alloc bootmem low pages(PAGE SIZE);
       for (k = 0; k < PTRS PER PTE; pte++, k++) {
               vaddr = i*PGDIR SIZE + j*PMD SIZE + k*PAGE SIZE;
               if (end && (vaddr >= end)) break;
                      *pte = mk pte phys( pa(vaddr), PAGE KERNEL);
       set pmd(pmd, pmd( KERNPG TABLE + pa(pte base)));
```

#### Note!!!

- The final PDE buffer coincides with the initial page table that maps 4 MB pages
- 4KB paging gets activated upon filling the entry of the PDE table (since the Page Size bit gets updated)
- For this reason the PDE entry is set only after having populated the corresponding PTE table to be pointed
- Otherwise memory mapping would be lost upon any TLB miss

# The set\_pmd macro

```
#define set pmd(pmdptr, pmdval) (*(pmdptr) = pmdval)
```

- Thia macro simply sets the value into one PMD entry
- Its input parameters are
  - > the pmdptr pointer to an entry of PMD (the type is pmd\_t)
  - The value to be loaded pmdval (of type pmd\_t, defined via casting)
- While setting up the kernel page table, this macro is used in combination with \_\_pa() (physical address) which returns an unsigned long
- The latter macro returns the physical address corresponding to a given virtual address within kernel space (except for some particular virtual address ranges)
- Such a mapping deals with [3,4] GB virtual addressing onto [0,1] GB physical addressing

# The mk\_pte\_phys() macro

```
mk_pte_phys(physpage, pgprot)
```

- The input parameters are
  - A frame physical address physpage, of type unsigned long
  - > A bit string paprot for a PTE, of type paprot t
- The macro builds a complete PTE entry, which includes the physical address of the target frame
- The result type is pte\_t
- The result value can be then assigned to one PTE entry

# PAE (Physical address extension)

- increase of the bits used for physical addressing
- offered by more recent x86 processors (e.g. Intel Pentium Pro) which provide up to 36 bits for physical addressing
- we can drive up to 64 GB of RAM memory
- paging gets operated at 3 levels (instead of 2)
- the traditional page tables get modified by extending the entries at 64-bits and reducing their number by a half (hence we can support 1/4 of the address space)
- an additional top level table gets included called "page directory pointer table" which entails 4 entries, pointed by CR3
- CR4 indicates whether PAE mode is activated or not (which is done via bit 5 PAE-bit)

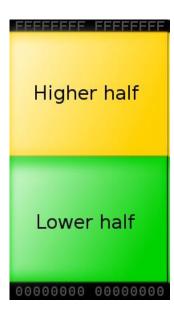
# x86-64 architectures

- They extend the PAE scheme via a so called "long addressing mode"
- Theoretically they allow addressing 2^64 bytes of logical memory
- In actual implementations we reach up to 2^48 canonical form addresses (lower/upper half within a total address space of 2^48)
- The total allows addressing spans over 256 TB
- Not all operating systems allow exploiting the whole range up to 256 TB of logical/physical memory
- LINUX currently allows for 128 TB for logical addressing of individual processes and 64 TB for physical addressing

# **Addressing scheme**

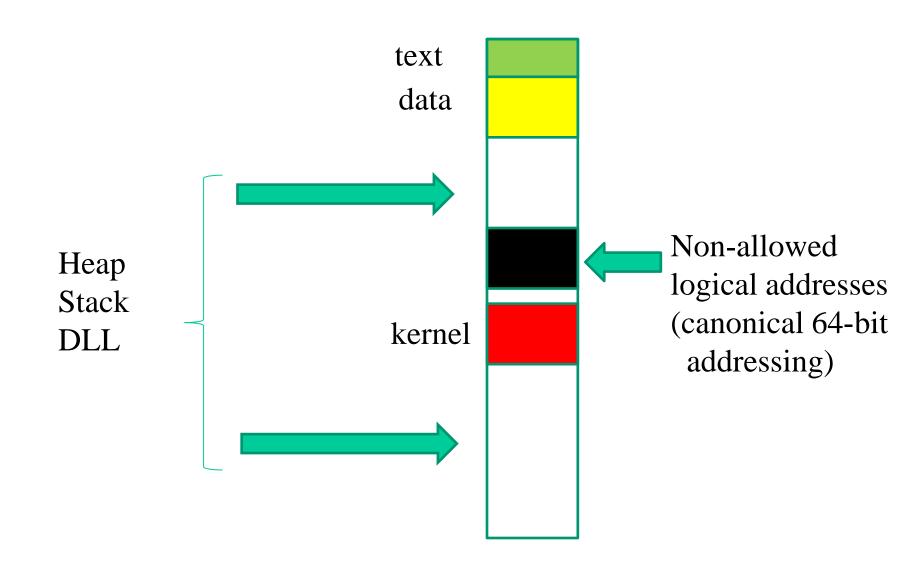
64-bit

48 out of 64-bit



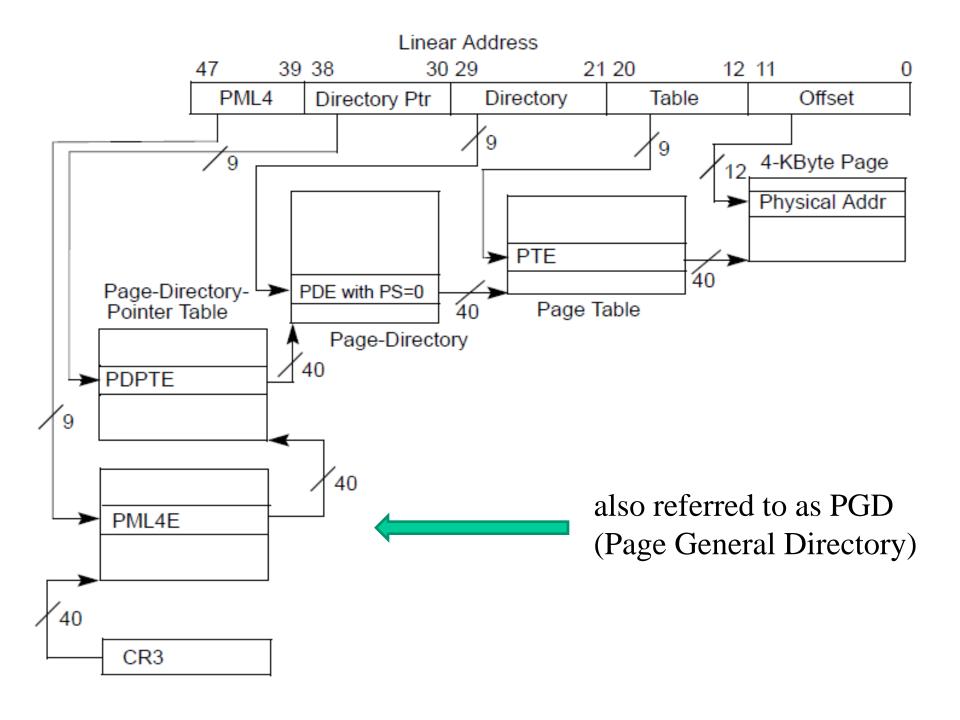


# Linux address space on x86-64 processors



# 48-bit addressing: page tables

- Page directory pointer has been expanded from 4 to 512 entries
- An additional paging level has been added thus reaching 4 levels, this is called "Page-Map level"
- Each Page-Map level table has 512 entries
- Hence we get 512<sup>4</sup> pages of size 4 KB that are addressable (namely, a total of 256 TB)



6 3	6 6 6 5 5 5 5 5 5 5 5 5 5 2 1 0 9 8 7 6 5 4 3 2	5 1 M <sup>1</sup>	M-1 3 3 3 2 1 0	2 2 2 2 2 2 2 2 2 2 9 8 7 6 5 4 3 2 1	2 1 1 1 1 1 1 1 1 0 9 8 7 6 5 4 3	1 1 1 2 1 0 9	8 7	6 5	4	3 2 1	0	
	Reserved <sup>2</sup>		Address of PML4 table				Ignored C			P W Ign T		CR3
3 D	Ignored	Rsvd.	Address of page-directory-pointer table Ign.			Rs vo	g A n	P C D	P U R W/S T	1	PML4E: present	
	Ignored										<u>0</u>	PML4E: not present
X D	Ignored	Rsvd.	Address of IGB page frame Reserved T		P A Ign. T	G <u>1</u>			P U R W/S W	- 1	PDPTE: 1GB page	
X D	Ignored	Rsvd.	Address of page directory Ign. D I PPUR CW/SW						1	PDPTE: page directory		
Ignored											<u>0</u>	PDTPE: not present
X D	Ignored	Rsvd.	Address of Reserved P A T		P A Ign. T	G <u>1</u>			P U R W/S T		PDE: 2MB page	
X D	Ignored	Rsvd.	Address of page table Ign. 0 I A P P U R I D T /S W						1	PDE: page table		
Ignored											<u>0</u>	PDE: not present
X D	Ignored	Rsvd.	Address of 4KB page frame						1	PTE: 4KB page		
				Ignored							<u>0</u>	PTE: not present

Figure 4-11. Formats of CR3 and Paging-Structure Entries with IA-32e Paging

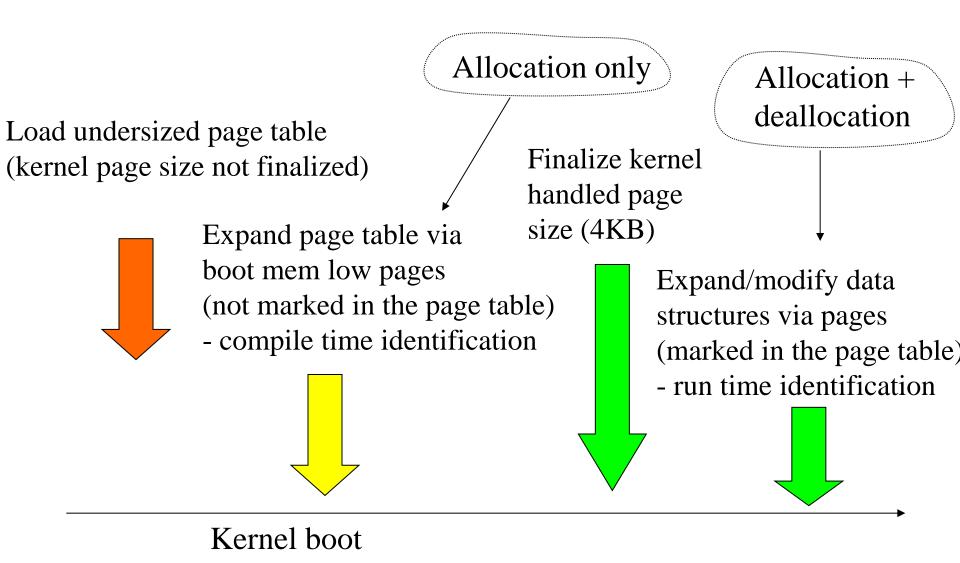
Table 4-19. Format of an IA-32e Page-Table Entry that Maps a 4-KByte Page

Bit Position(s)	Contents
0 (P)	Present; must be 1 to map a 4-KByte page
1 (R/W)	Read/write; if 0, writes may not be allowed to the 4-KByte page referenced by this entry (see Section 4.6)
2 (U/S)	User/supervisor; if 0, user-mode accesses are not allowed to the 4-KByte page referenced by this entry (see Section 4.6)
3 (PWT)	Page-level write-through; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
4 (PCD)	Page-level cache disable; indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
5 (A)	Accessed; indicates whether software has accessed the 4-KByte page referenced by this entry (see Section 4.8)
6 (D)	Dirty; indicates whether software has written to the 4-KByte page referenced by this entry (see Section 4.8)
7 (PAT)	Indirectly determines the memory type used to access the 4-KByte page referenced by this entry (see Section 4.9.2)
8 (G)	Global; if CR4.PGE = 1, determines whether the translation is global (see Section 4.10); ignored otherwise
11:9	Ignored
(M-1):12	Physical address of the 4-KByte page referenced by this entry
51:M	Reserved (must be 0)
62:52	Ignored
63 (XD)	If IA32_EFER.NXE = 1, execute-disable (if 1, instruction fetches are not allowed from the 4-KByte page controlled by this entry; see Section 4.6); otherwise, reserved (must be 0)

# **Huge pages**

- Ideally x86-64 processors support them starting from PDPT
- Linux typically offers the support for huge pages pointed to by the PDE (page size 512\*4KB)
- See: /proc/meminfo and /proc/sys/vm/nr\_hugepages
- These can be "mmaped" via file descriptors and/or mmap parameters (e.g. MAP\_HUGETLB flag)
- They can also be requested via the madvise (void\*, size\_t, int) system call (with MADV\_HUGEPAGE flag)

# Reaching vs allocating/deallocating memory



# Core map

- It is an array of mem\_map\_t (also known as struct page) structures defined in include/linux/mm.h
- The actual type definition is as follows:

```
typedef struct page {
    struct list head list;
                                      /* ->mapping has some page lists. */
                                       /* The inode (or ...) we belong to. */
    struct address space *mapping;
    unsigned long index;
                                       /* Our offset within mapping. */
    struct page *next hash;
                                       /* Next page sharing our hash bucket in
                                          the pagecache hash table. */
                                      /* Usage count, see below. */
    atomic t count;
                                      /* atomic flags, some possibly
    unsigned long flags;
                                         updated asynchronously */
                                       /* Pageout list, eg. active list;
    struct list head lru;
                                         protected by pagemap lru lock !! */
    struct page **pprev_hash; /* Complement to *next hash. */
    struct buffer head * buffers; /* Buffer maps us to a disk block. */
#if defined(CONFIG HIGHMEM) || defined(WANT PAGE VIRTUAL)
    void *virtual;
                                      /* Kernel virtual address (NULL if
                                          not kmapped, ie. highmem) */
#endif /* CONFIG HIGMEM || WANT_PAGE_VIRTUAL */
} mem map t;
```

#### **Fields**

- Most of the fields are used to keep track of the interactions between memory management and other kernel sub-systems (such as I/O)
- Memory management proper fields are
  - > struct list\_head list (whose type is defined in include/linux/lvm.h), which is used to organize the frames into free lists
  - ➤ atomic\_t count, which counts the virtual references mapped onto the frame (it is managed via atomic updates, such as with LOCK directives)
  - > unsigned long flags, this field keeps the status bits for the frame, such as:

```
#define PG_locked 0
#define PG_referenced 2
#define PG_uptodate 3
#define PG_dirty 4
#define PG_lru 6
#define PG_lru 6
```

# Core map initialization (i386/kernel 2.4)

- Initially we only have the core map pointer
- This is mem\_map and is declared in mm/memory.c
- Pointer initialization and corresponding memory allocation occur within free area init()
- After initializing, each entry will keep the value 0 within the count field and the value 1 into the PG\_reserved flag within the flags field
- Hence no virtual reference exists for that frame and the frame is reserved
- Frame un-reserving will take place later via the function mem\_init() in arch/i386/mm/init.c (by resetting the bit PG reserved)

# Free list organization: single NUMA zone – or NUMA unaware – protected mode case (e.g. kernel 2.4)

- we have 3 free lists of frames, depending on the frame positioning within the following zones: DMA (DMA ISA operations), NORMAL (room where the kernel can reside), HIGHMEM (room for user data)
- The corresponding defines are in include/linux/mmzone.h:

```
#define ZONE_DMA 0
#define ZONE_NORMAL 1
#define ZONE_HIGHMEM 2
#define MAX NR ZONES 3
```

• The corresponding sizes are usually defined as

#### Free list data structures

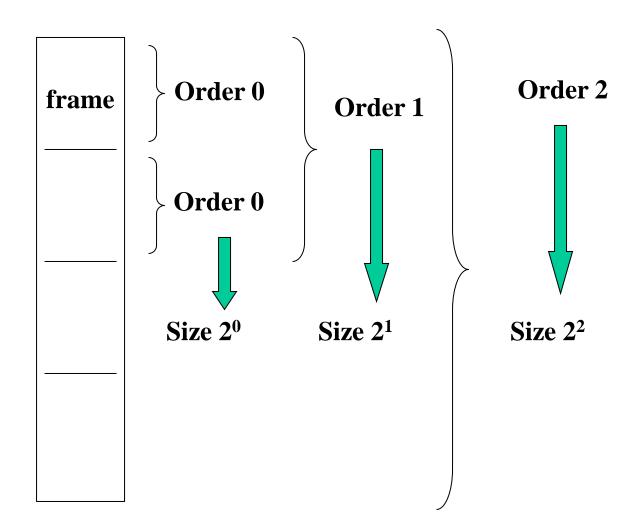
• Free lists information is kept within the pg\_data\_t data structure defined in include/linux/mmzone.h, and whose actual instance is contig\_page\_data, which is declared in mm/numa.c

```
typedef struct pglist data {
           zone t node zones[MAX NR ZONES];
           zonelist t node zonelists[GFP ZONEMASK+1];
           int nr zones;
Field of
           struct page *node mem map;
           unsigned long *valid addr bitmap;
interest
           struct bootmem data *bdata;
           unsigned long node start paddr;
           unsigned long node start mapnr;
           unsigned long node size;
           int node id;
           struct pglist data *node next;
       pg data t;
```

• the zone\_t type is defined in include/linux/mmzone.h as follows

```
typedef struct zone struct {
            spinlock t
                                 lock;
            unsigned
                                 long
                                        free pages;
            zone watermarks t watermarks [MAX NR ZONES];
                                need balance;
            unsigned long
            unsigned long nr active pages, nr_inactive_pages;
            unsigned long
                                 nr cache pages;
            free area t
                                 free area[MAX ORDER];
                                 * wait table;
            wait queue head t
            unsigned long
                                 wait table size;
                                wait table shift;
           unsigned long
                                                     Up to 11 in
Fields of
            struct pglist data
                                 *zone pgdat;
                                                     recent
interest
                                 *zone_mem_map;
           struct page
                                                      kernel versions
            unsigned long
                                 zone start paddr;
                                                      (it was typically
            unsigned long
                                 zone start mapnr;
                                                      5 before)
            char
                                 *name;
            unsigned long
                                 size;
            unsigned long
                                 realsize;
       zone t;
```

# **Buddy system features**



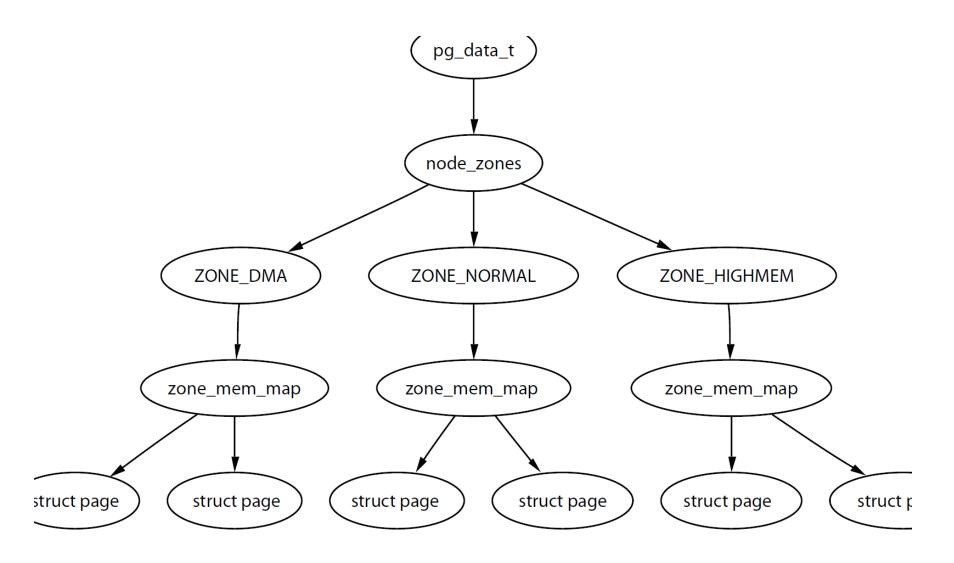
• free area t is defined in the same file as

where

```
struct list_head {
     struct list_head *next, *prev;
}
```

- overall, any element of the free\_area[] array keeps
  - A pointer to the first free frame associated with blocks of a given order
  - A pointer to a bitmap that keeps fragmentation information according to the 'buddy system' organization

# A scheme (picture from: Understanding the Linux Virtual Memory Manager – Mel Gorman)



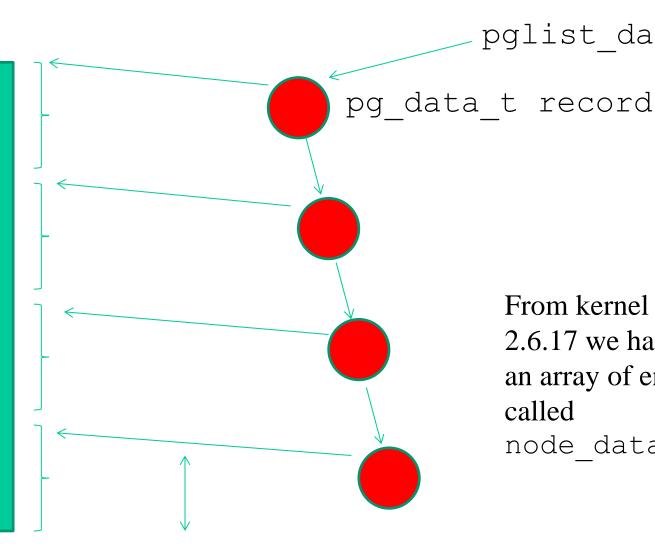
# Jumping to more recent kernels for actually NUMA machines – kernel 2.6

- The concept of multiple NUMA zones is represented by a struct pglist\_data even if the architecture is Uniform Memory Access (UMA)
- This struct is always referenced by its typedef pg\_data\_t
- Every node in the system is kept on a NULL terminated list called pgdat\_list, and each node is linked to the next with the field pg data t-node next
- For UMA architectures like PC desktops, only one static pg\_data\_t structure, as already seen called contig\_page\_data is used

mem map

# A scheme

One buddy allocator per each node



From kernel 2.6.17 we have an array of entries called node data[]

pglist data

struct page \*node mem map

## **Summarizing (still for 2.4 kernel)**

- Architecture setup occurs via setup\_arch() which will give rise to
  - ➤ a Core Map with all reseved frames
  - Free lists that looks like if no frame is available (in fact they are all reserved at this stage)
- Most of the work is done by free\_area\_init()
- This relies on bitmaps allocation services based on alloc\_bootmem()

# Releasing boot used pages to the free lists

```
static unsigned long init
free all bootmem core(pg data t *pgdat)
      for (i = 0; i < idx; i++, page++) {
            if (!test bit(i, bdata->node bootmem map)) {
      // il frame non deve restare riservato
                  count++;
                  ClearPageReserved(page);
                  set page count (page, 1);
                   free page(page);
      total += count;
      return total;
```

# Allocation contexts (more generally, kernel level execution contexts)

#### Process context

- Allocation is caused by a system call
  - Not satisfiable → wait is experienced along the current execution trace
  - Priority based schemes

#### • Interrupt

- Allocation requested by an interrupt handler
  - Not satisfiable → no-wait is experienced along the current execution trace
  - Priority independent schemes

# **Buddy-system API**

- After booting, the memory management system can be accessed via proper APIs, which drive operations on the aforementioned data structures
- The prototypes are in #include linux/malloc.h>
- The very base allocation APIs are (bare minimal page aligned allocation)
  - unsigned long get\_zeroed\_page(int flags)
    removes a frame from the free list, sets the content to zero and returns
    the virtual address
  - unsigned long \_\_get\_free\_page(int flags)
    removes a frame from the free list and returns the virtual address
  - unsigned long \_\_get\_free\_pages(int flags, unsigned long order)
    - removes a block of contiguous frames with given order from the free list and returns the virtual address of the first frame

- void free\_page (unsigned long addr)
  puts a frame into the free list again, having a given initial virtual
  address
- ➤ void free\_pages (unsigned long addr, unsigned long order) puts a block of frames of given order into the free list again Note!!!!!! Wrong order gives rise to kernel corruption

### flags: used contexts

GFP\_ATOMIC the call cannot lead to sleep (this is for interrupt contexts)

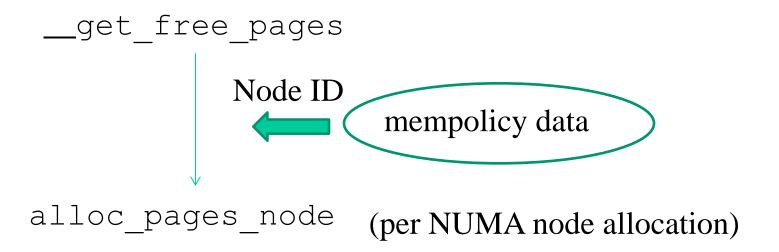
GFP\_USER - GFP\_BUFFER - GFP\_KERNEL the call can lead to sleep

# Binding actual allocation to NUMA nodes

The real core of the Linux page allocator is the function

```
struct page *alloc_pages_node(int nid, unsigned
  int flags, unsigned int order);
```

Hence the actual allocation chain is:



### Mempolicy details

- Generally speaking, mempolicies determine what NUMA node needs to be involved in a specific allocation operation which is thread specific
- Starting from kernel 2.6.18, the determination of mempolicies can be configured by the application code via system calls

### **Synopsis**

```
#include <numaif.h>
int set_mempolicy(int mode, unsigned long
    *nodemask, unsigned long maxnode);
```

sets the NUMA memory policy of the calling process, which consists of a policy mode and zero or more nodes, to the values specified by the *mode*, *nodemask* and *maxnode* arguments

The *mode* argument must specify one of MPOL\_DEFAULT, MPOL\_BIND, MPOL\_INTERLEAVE or MPOL\_PREFERRED

### ... another example

### **Synopsis**

```
#include <numaif.h>
int mbind(void *addr, unsigned long len, int
mode, unsigned long *nodemask, unsigned long
maxnode, unsigned flags);
```

sets the NUMA memory policy, which consists of a policy mode and zero or more nodes, for the memory range starting with *addr* and continuing for *len* bytes. The memory policy defines from which node memory is allocated.

# ... finally you can also move pages around

### **Synopsis**

```
#include <numaif.h>
long move_pages(int pid, unsigned long count,
void **pages, const int *nodes, int *status,
int flags);
```

moves the specified *pages* of the process *pid* to the memory nodes specified by *nodes*. The result of the move is reflected in *status*. The *flags* indicate constraints on the pages to be moved.

# The case of frequent allocation/deallocation of a target-specific data structures

- Here we are talking about allocation/deallocation operations of data structures
  - 1.that are used for a target-specific objective (e.g. in terms of data structures to be hosted)
  - 2.which are requested/released frequently
- The problem is that getting the actual buffers (pages) from the buddy system will lead to contention and consequent synchronization costs (does not scale)
- In fact the (per NUMA node) buddy system operates with spinlock synchronized critical sections
- Kernel design copes with this issue by using pre-reserved buffers with lightweight allocation/release logic

### ... a classical example

- The allocation and deletion of page tables, at any level, is a very frequent operation, so it is important the operation is as quick as possible
- Hence the pages used for the page tables are cached in a number of different lists called *quicklists*
- For 3 levels paging, PGDs, PMDs and PTEs have two sets of functions each for the allocation and freeing of page tables.
- The allocation functions are pgd\_alloc(), pmd\_alloc() and pte\_alloc(), respectively the free functions are, predictably enough, called pgd\_free(), pmd\_free() and pte\_free()
- Broadly speaking, these APIs implement caching

### **Actual quicklists**

- Defined in include/linux/quicklist.h
- They are implemented as a list of per-core page lists
- There is no need for synchronization
- If allocation fails, they revert to

```
__get_free_page()
```

### **Quicklist API**

```
static inline void *quicklist alloc(int nr, gfp_t flags, ...) {
      struct quicklist *q;
      void **p = NULL;
      q = &get cpu var(quicklist)[nr];
      p = q->page;
      if (likely(p)) {
             q->page = p[0];
             p[0] = NULL;
             q->nr pages--;
      put_cpu_var(quicklist);
      if (likely(p))
             return p;
      p = (void *) get free page(flags | GFP ZERO);
      return p;
```

### SLAB (or SLUB) allocator: a cache of 'small' size buffers

- The prototypes are in #include ux/malloc.h>
- The main APIs are
  - void \*kmalloc(size\_t size, int flags)
    allocation of contiguous memory of a given size it returns the virtual address
  - void kfree(void \*obj)
    frees memory allocated via kmalloc()
- Main features:
  - ➤ Cache aligned delivery of memory chunks (performance optimal access of related data within the same chunk)
  - ➤ Fast allocation/deallocation support
- Clearly, we also can perform node-specific requests via

### What about large size allocations

- Classically employed while adding large size data structures to the kernel in a stable way
- This is the case when, e.g., mounting external modules
- The main APIs are:
  - ➤ void \* vmalloc (unsigned long size); allocates memory of a given size, which can be non-contiguous, <u>and</u> <u>returns the virtual address</u> (the corresponding frames are anyhow reserved)
  - F void vfree (void \* addr)
    frees the above mentioned memory

Be careful of performance effects due to global TLB invalidations!!

# Logical/Physical address translation for kernel directly mapped memory (not vmalloc one)

### We can exploit the macros

\_\_pa() \_\_va()

In generic kernel versions

# kmalloc vs vmalloc: an overall picture

- Allocation size:
  - > 128 KB for kmalloc (cache aligned)
  - ➤ 64/128 MB for vmalloc
- Physical contiguousness
  - > Yes for kmalloc
  - ➤ No for vmalloc
- Effects on TLB
  - ➤ None for kmalloc
  - ➤ Global for vmalloc (transparent to vmalloc users)